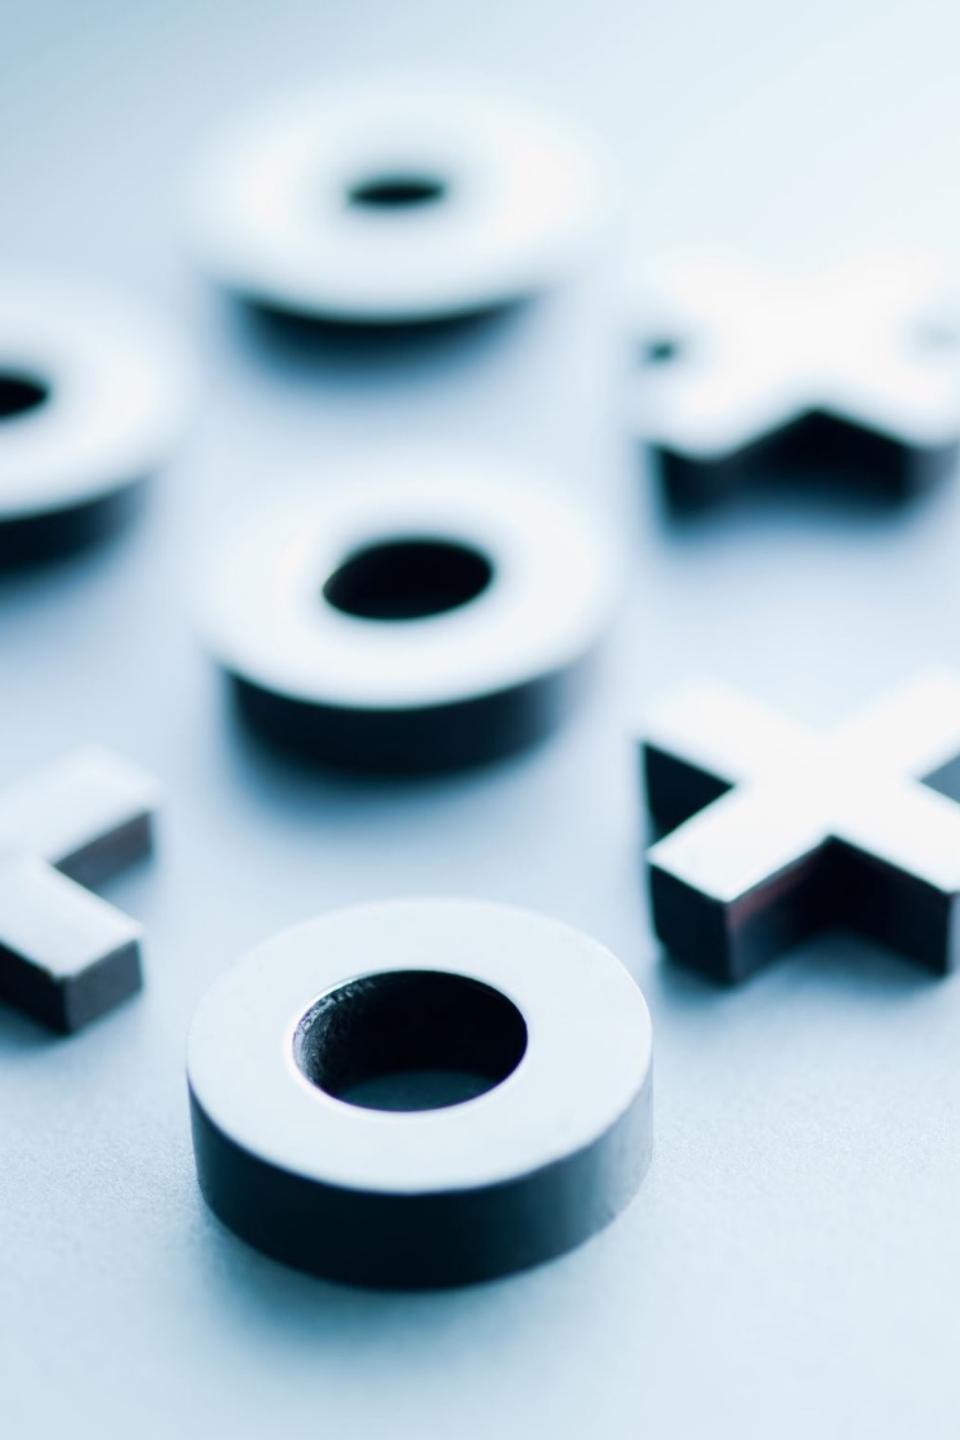


Redes Neuronales y Deep learning

M.C. JUAN EQUIHUA



Agenda

1:20

Calentamieto

Multi-Layer Perceptron y FNN

Break (15min)

NLP basics

Demo 1 – Sentiment Analysis

Computer vision y deep learning

1:00

Demo 2 – Clasificando imágenes

Break (15min)

1:20

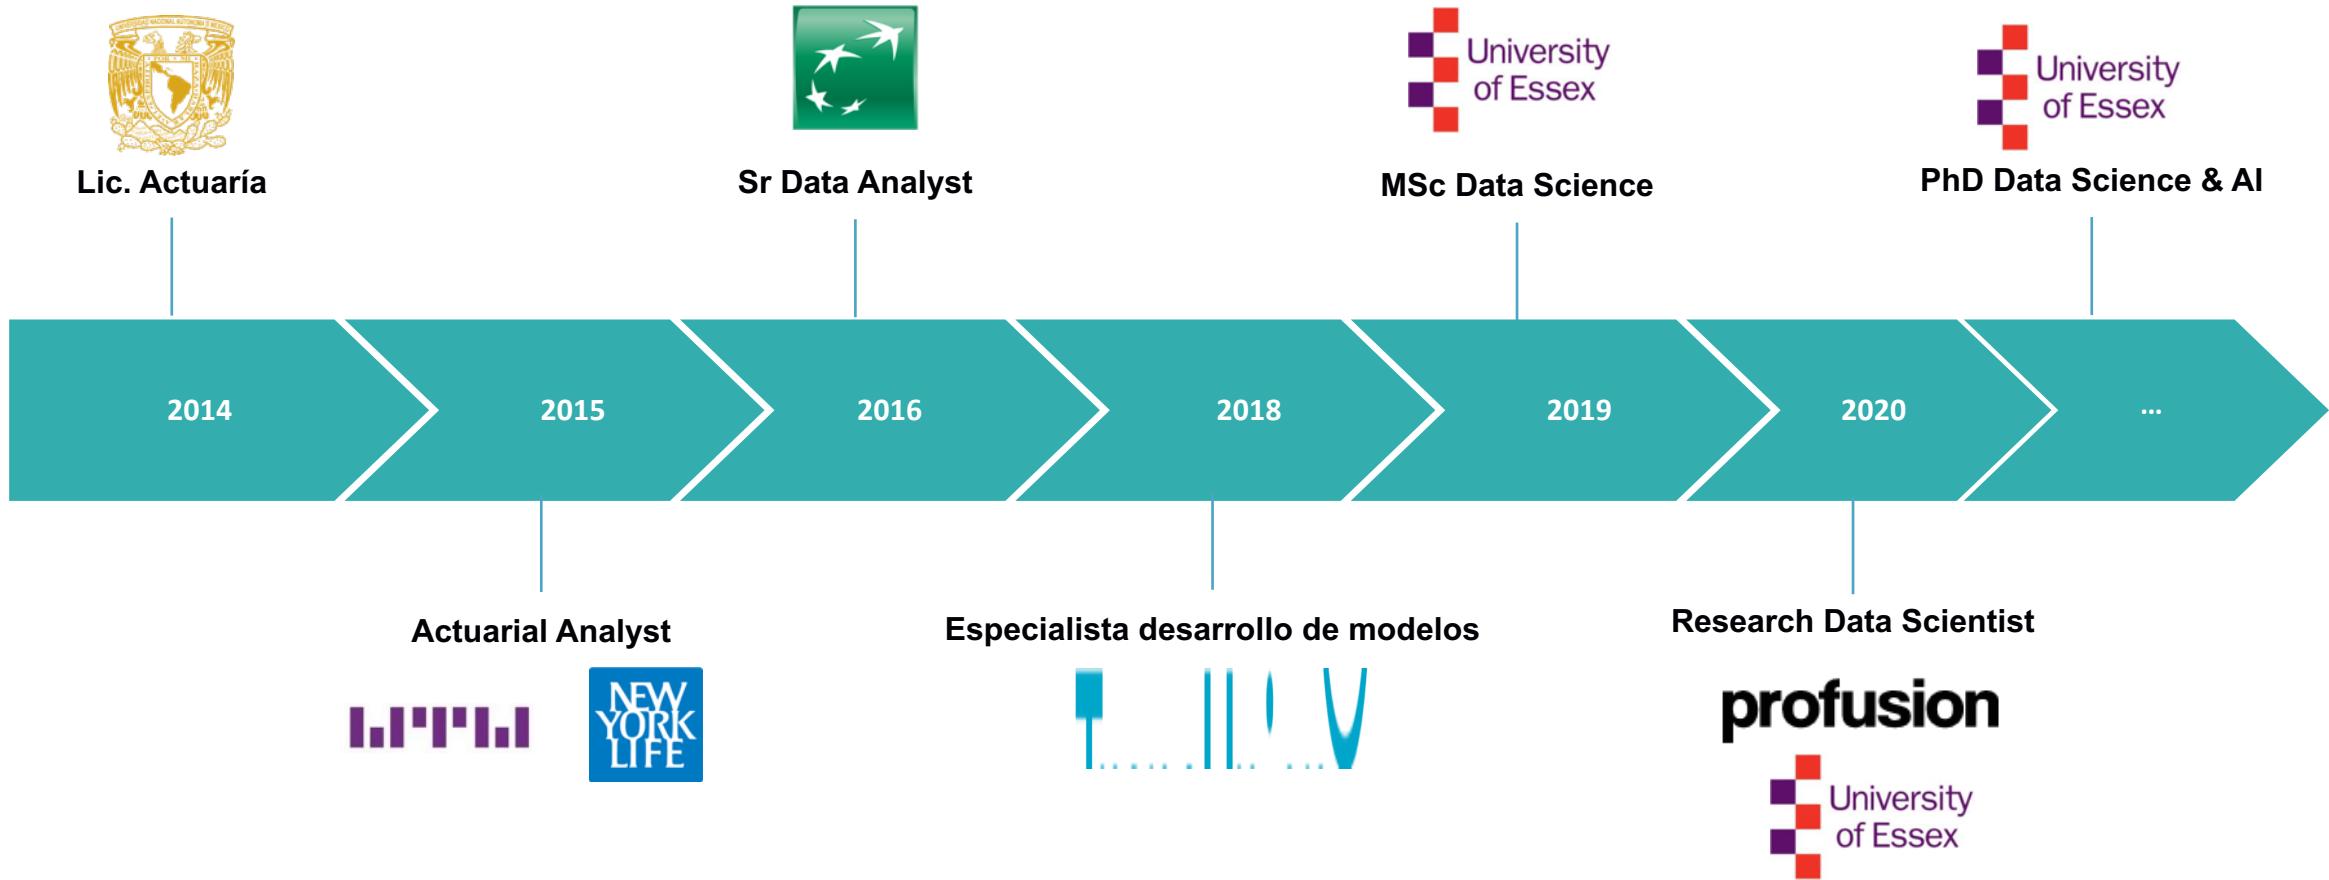
RNN

Arquitecturas avanzadas en NN

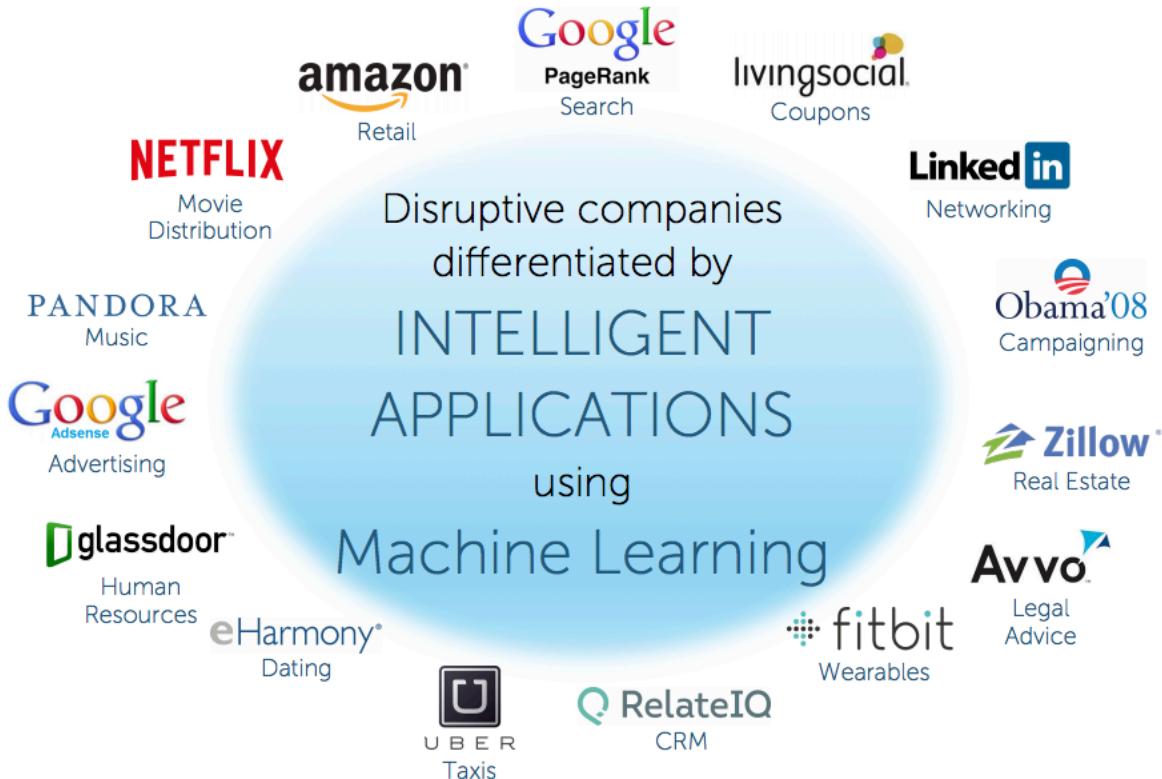
Demo 3 – AE & LSTM

Resumen y Recursos adicionales

Juan Equihua



Machine Learning está en todos los lados



Redes Neuronales en la vida real

Grammarly at correcting sentences

Type your title

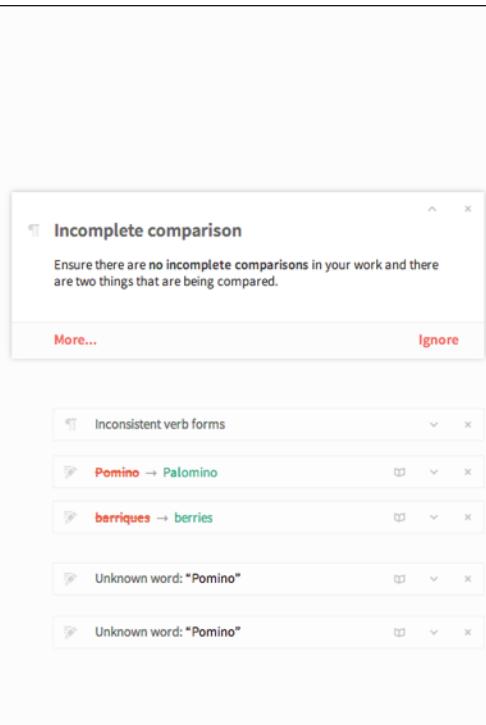
The venerable house Marchesi de Frescobaldi has been in the wine business in Tuscany for seven centuries. I've talked about the family history before—trading Michelangelo wine for paintings, ledgers that still list sales to Henry VIII, etc. There are not that many industries out there that are growing, have new companies coming into the market every year, and yet still have established brands going back so long. It also makes me wonder which American producers will still be going strong in the 28th century.

The Pomino DOC is much younger by comparison, established in the 1970s and permitting the use of traditionally French grapes in both red and white blends. This is my first exposure to this particular DOC, and I was pleased with the style. This particular wine is aged in stainless steel with a small amount held in French barriques for three months.

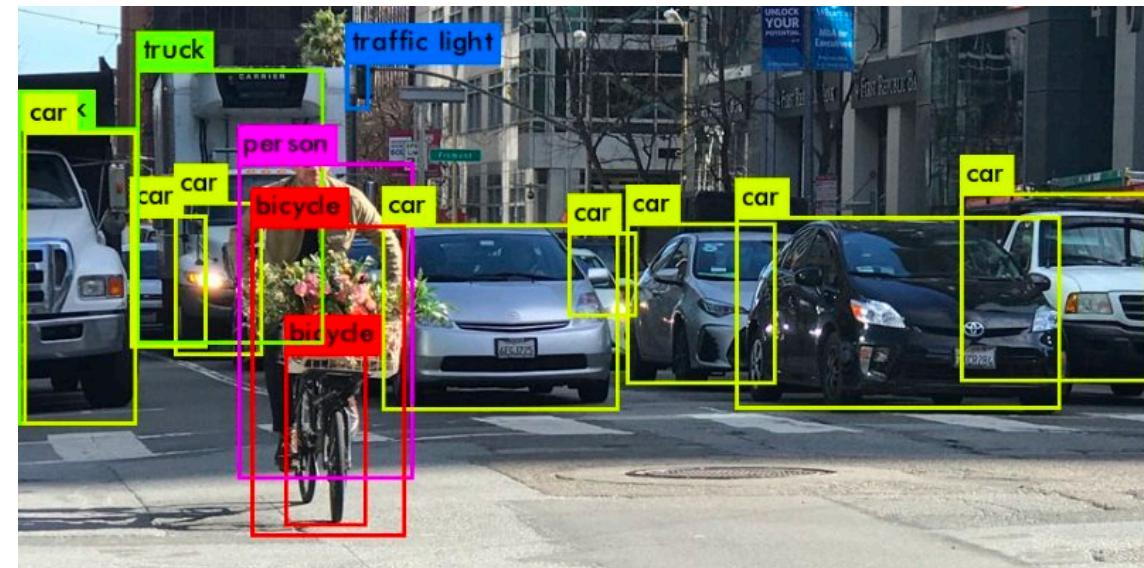
2012 Castello di Pomino Bianco

Pomino DOC

Chardonnay and Pinot Bianco with other unspecified white grapes
\$15, 12.5% abv.

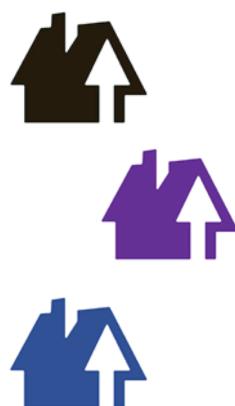


Object recognition for self-driving cars

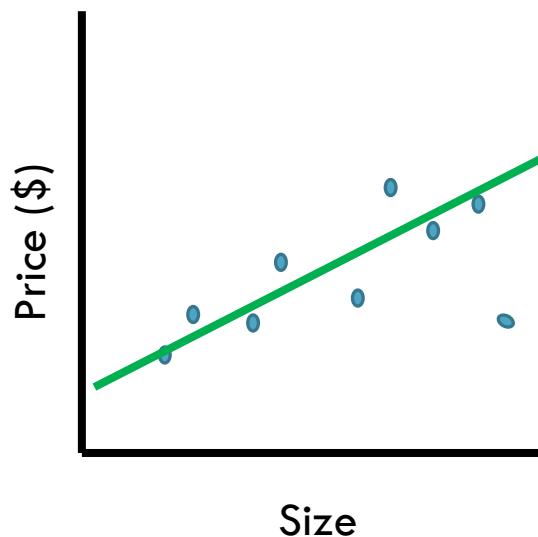


Ejemplo 1. Prediciendo precios de casas

Datos



Entrenamiento



Predicción



Predecir para
nuevas
observaciones

Precio?

Objetivo

Encontrar los pesos w para la regression lineal. Como?
minimizando el una función de Perdida.

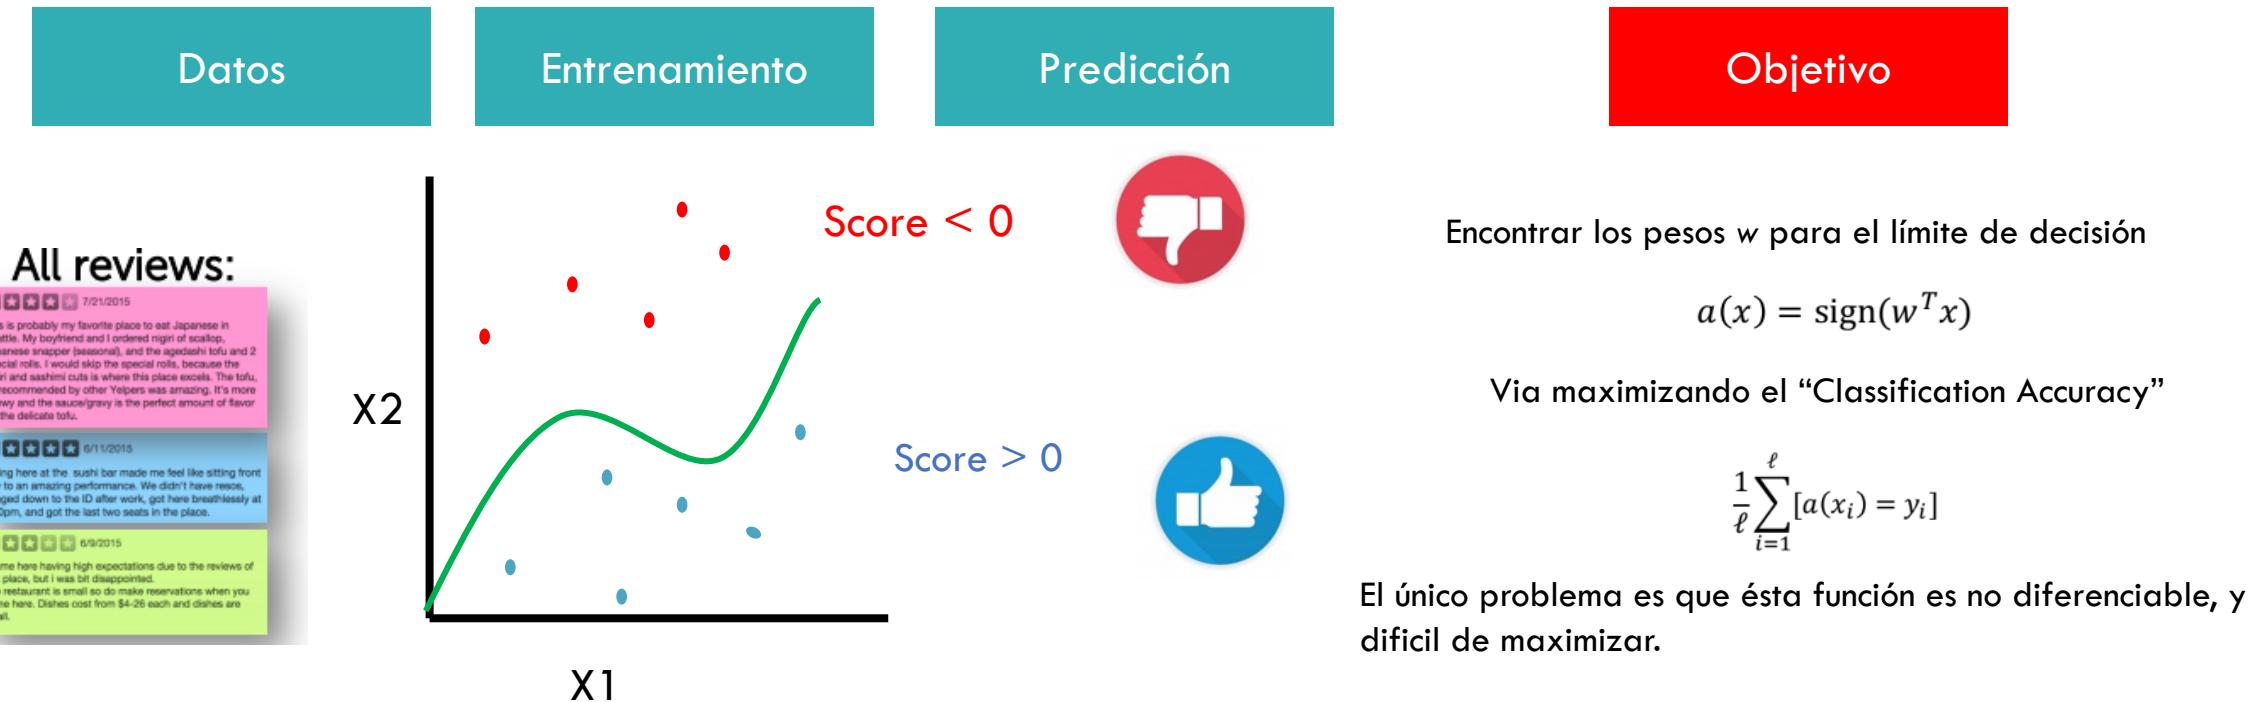
$$d(x) = w_0 + w_1 x_1 + w_2 x_2$$

$$L(w) = \frac{1}{\ell} \|Xw - y\|^2$$

Metricas comunes:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Ejemplo 2. Prediciendo sentimientos en reviews

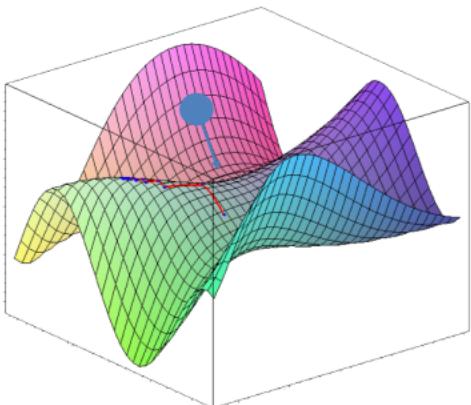


Como encontrar w ?

Stochastic Gradient Decent

SGD es un algoritmo para minimizar cualquier función de pérdida diferenciable.

Optimization problem $L(w) = \min(w)$



w^0 — initialization

$\nabla L(w^0) = \left(\frac{\partial L(w^0)}{\partial w_1}, \dots, \frac{\partial L(w^0)}{\partial w_n} \right)$ — gradient vector

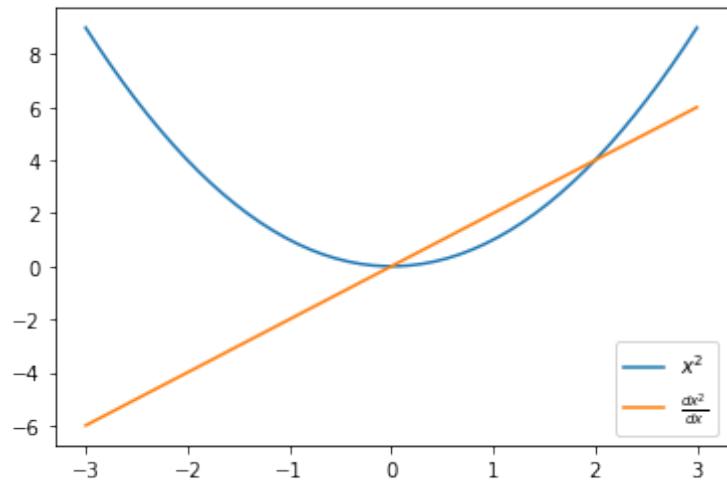
$w^1 = w^0 - \eta_1 \nabla L(w^0)$ — gradient step

while True:

$$w^t = w^{t-1} - \eta_t \nabla L(w^{t-1})$$

if $\|w^t - w^{t-1}\| < \epsilon$ then break

1-Dimension SGD



$$\nabla L(w^0) = \left(\frac{\partial L(w^0)}{\partial w_1}, \dots, \frac{\partial L(w^0)}{\partial w_n} \right) \text{ — gradient vector}$$

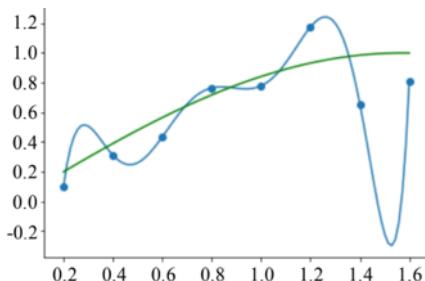
$$w^1 = w^0 - \eta_1 \nabla L(w^0) \text{ — gradient step}$$

L1 y L2 Losses

L1 Penalty

$$L_{reg}(w) = L(w) + \lambda \|w\|_1 \rightarrow \min_w$$

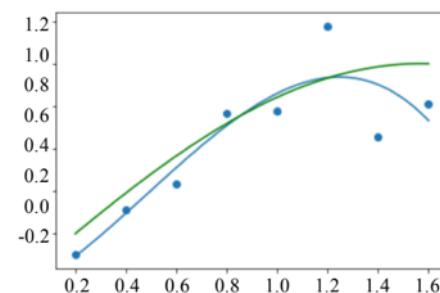
- $\|w\|_1 = \sum_{j=1}^d |w_j|$
- Drives some weights **exactly** to zero
- Learns sparse models
- Cannot be optimized with simple gradient methods



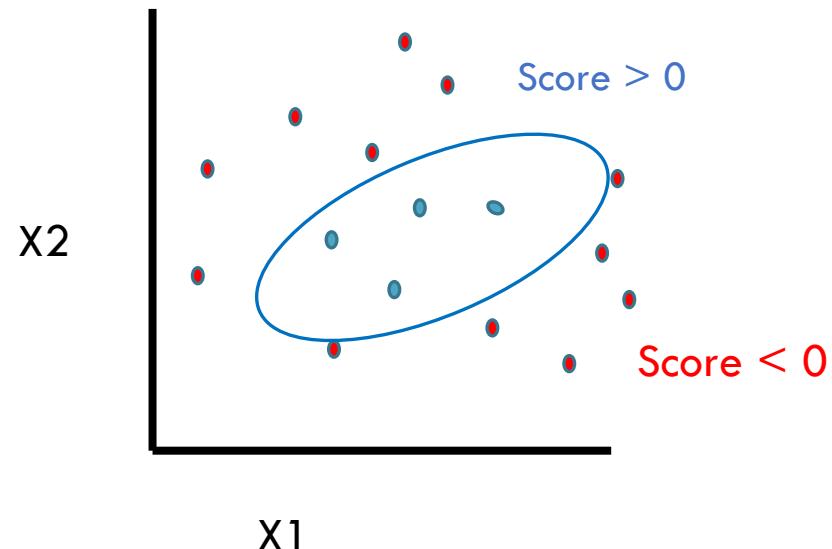
L2 Penalty

$$L_{reg}(w) = L(w) + \lambda \|w\|^2 \rightarrow \min_w$$

- $\|w\|^2 = \sum_{j=1}^d w_j^2$
- Drives all weights **closer** to zero
- Can be optimized with gradient methods



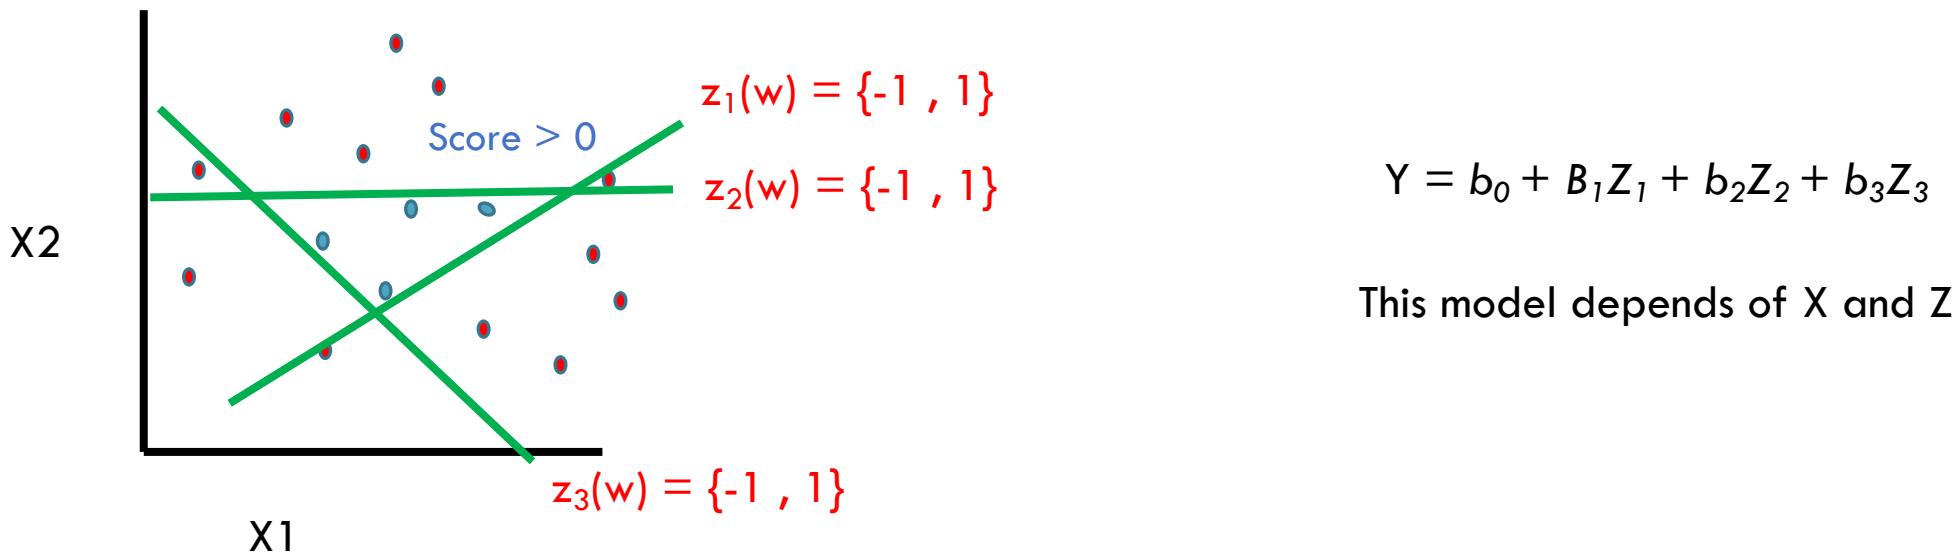
Como resolver un problema no linear?



No existe un único límite de decisión (decision boundary) que separe las dos diferentes clases

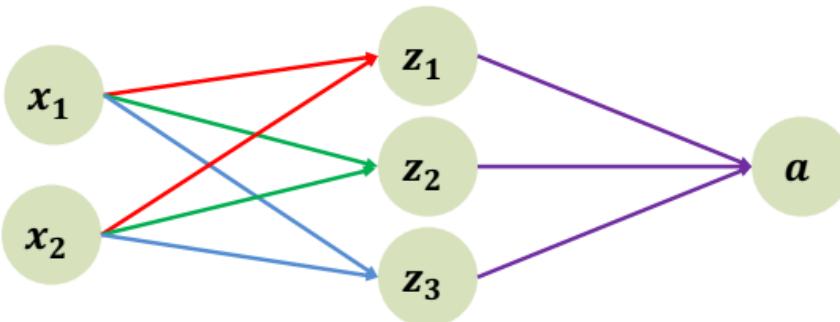
Regressões agregadas

Podemos resolver este problema mediante diferentes 3 regresiones lineales y después una regresión logística.



La primer red neuronal: MLP

- Sabemos la forma de nuestro clasificador
 - $z_i = \sigma(\mathbf{w}_{0,i} + \mathbf{w}_{1,i}x_1 + \mathbf{w}_{2,i}x_2)$
 - $a(x) = \sigma(\mathbf{w}_0 + \mathbf{w}_1z_1(x) + \mathbf{w}_2z_2(x) + \mathbf{w}_3z_3(x))$
- Ahora podemos reescribir nuestro clasificador como gráfica computacional



Nodes: computed variables ($x_1, x_2, z_1, z_2, z_3, a$)

Edges: dependencies (we need x_1 and x_2 to compute z_1)

Esta gráfica tiene el nombre de Multi-Layer Perceptron y es la red neuronal más simple:

- Input Layer (X_i)
- Hidden layer (Z_i)
- Output Layer (a)

Cada nodo es una regression (o neurona) con su propia function de activación

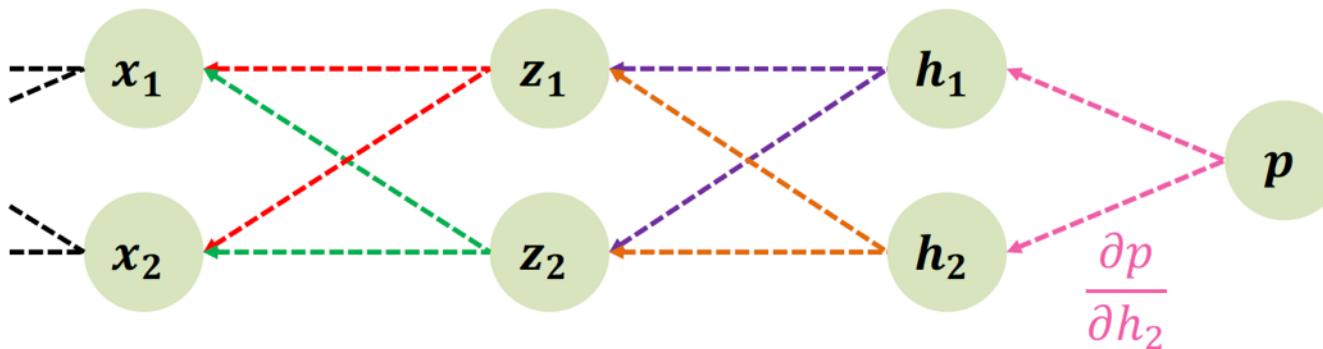
Back - Propagation

Como Encontrar w en un MLP?

$$h_2 = \sigma(w_0 + w_1 z_1 + w_2 z_2)$$

Gradient
Descent:

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial p} \frac{\partial p}{\partial w_1} = \frac{\partial L}{\partial p} \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial w_1}$$



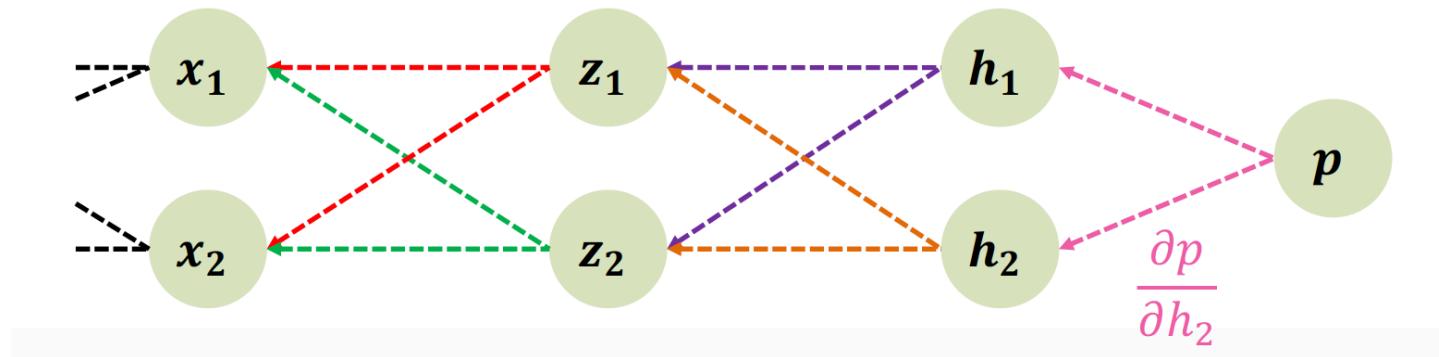
$$3: \frac{\partial p}{\partial h_1} \quad \frac{\partial p}{\partial h_2}$$

We will need these for GD

$$2: \frac{\partial p}{\partial z_1} = \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_1} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_1} \quad \frac{\partial p}{\partial z_2} = \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2}$$

$$1: \frac{\partial p}{\partial x_1} = \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_1} \frac{\partial z_1}{\partial x_1} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_1} \frac{\partial z_1}{\partial x_1} + \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} \frac{\partial z_2}{\partial x_1} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2} \frac{\partial z_2}{\partial x_1}$$

$$\frac{\partial p}{\partial x_2} = \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_1} \frac{\partial z_1}{\partial x_2} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_1} \frac{\partial z_1}{\partial x_2} + \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} \frac{\partial z_2}{\partial x_2} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2} \frac{\partial z_2}{\partial x_2}$$



$$3: \frac{\partial p}{\partial h_1} \quad \frac{\partial p}{\partial h_2}$$

$$2: \frac{\partial p}{\partial z_1} = \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_1} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_1}$$

$$1: \frac{\partial p}{\partial x_1} = \boxed{\frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_1} \frac{\partial z_1}{\partial x_1} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_1} \frac{\partial z_1}{\partial x_1}} + \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} \frac{\partial z_2}{\partial x_1} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2} \frac{\partial z_2}{\partial x_1}$$

$$1: \frac{\partial p}{\partial x_2} = \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_1} \frac{\partial z_1}{\partial x_2} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_1} \frac{\partial z_1}{\partial x_2} + \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} \frac{\partial z_2}{\partial x_2} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2} \frac{\partial z_2}{\partial x_2}$$

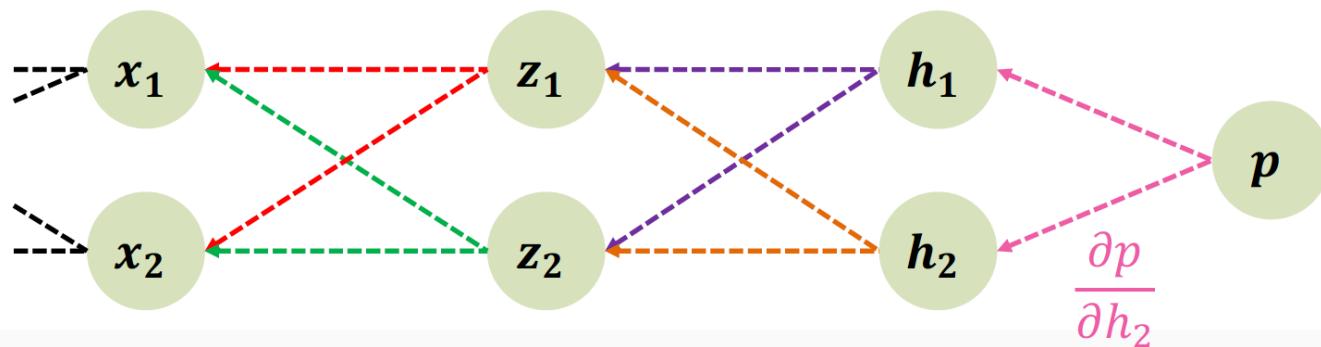
We will need these for GD

$$\frac{\partial p}{\partial z_2} = \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2}$$

Podemos re-utilizar derivadas anteriores

$$1. \frac{\partial p}{\partial x_1} = \boxed{\left(\frac{\partial p}{\partial z_1} \right) \frac{\partial z_1}{\partial x_1}} + \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} \frac{\partial z_2}{\partial x_1} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2} \frac{\partial z_2}{\partial x_1}$$

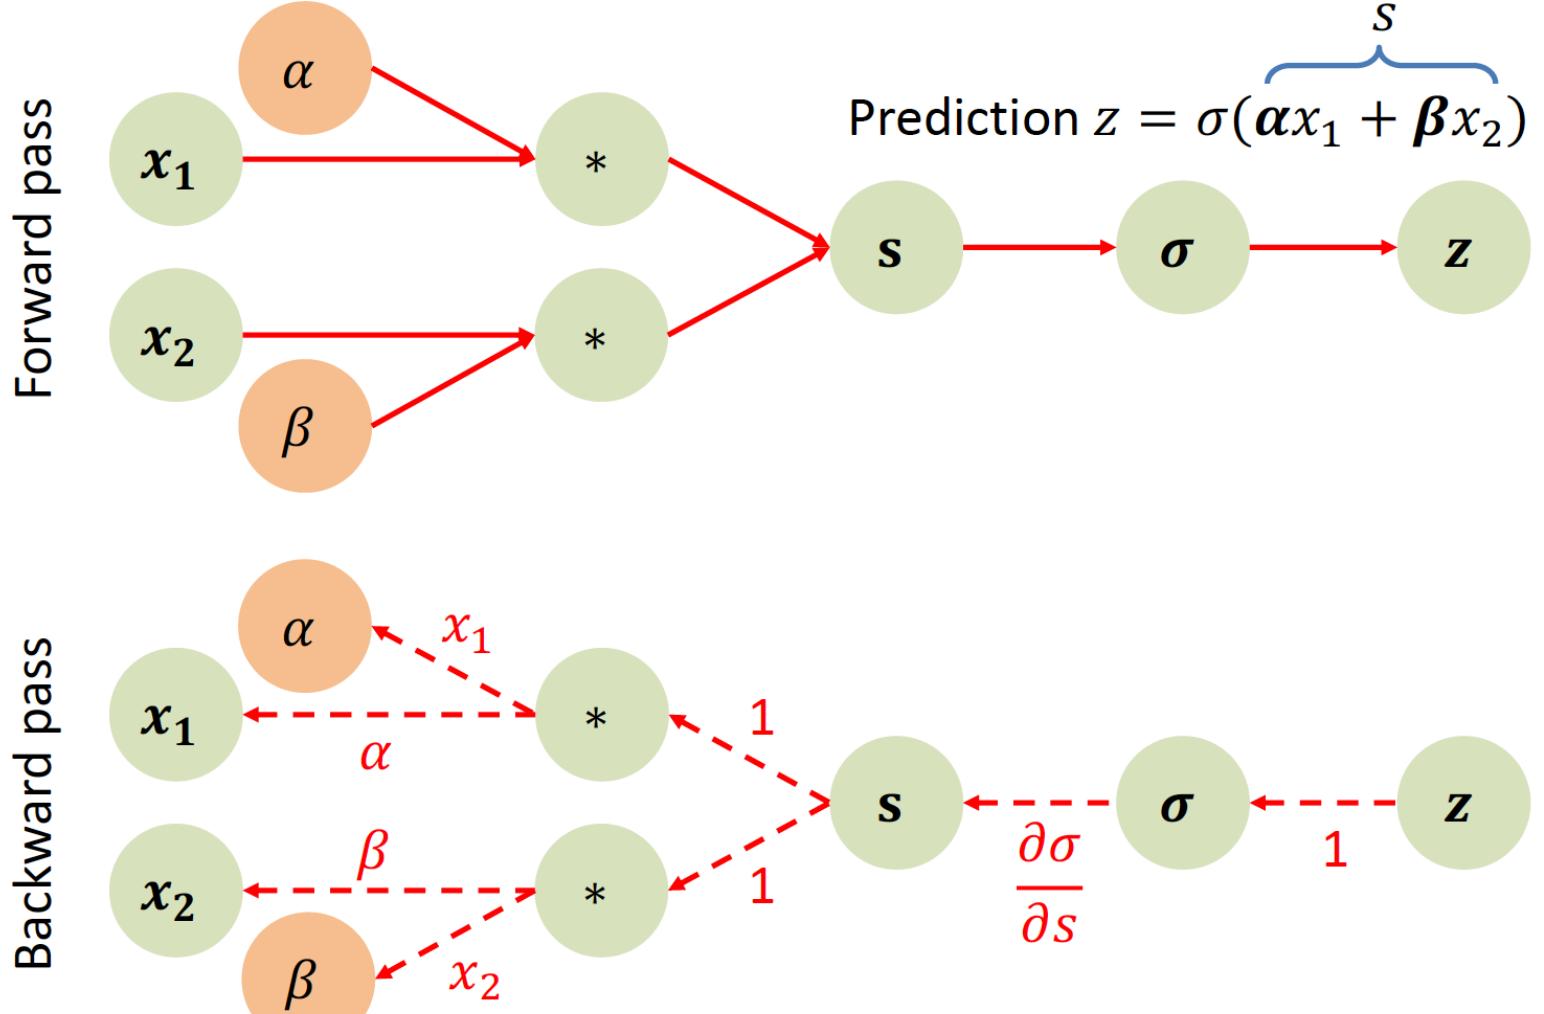
$$\frac{\partial p}{\partial x_2} = \boxed{\left(\frac{\partial p}{\partial z_1} \right) \frac{\partial z_1}{\partial x_2}} + \frac{\partial p}{\partial h_1} \frac{\partial h_1}{\partial z_2} \frac{\partial z_2}{\partial x_2} + \frac{\partial p}{\partial h_2} \frac{\partial h_2}{\partial z_2} \frac{\partial z_2}{\partial x_2}$$



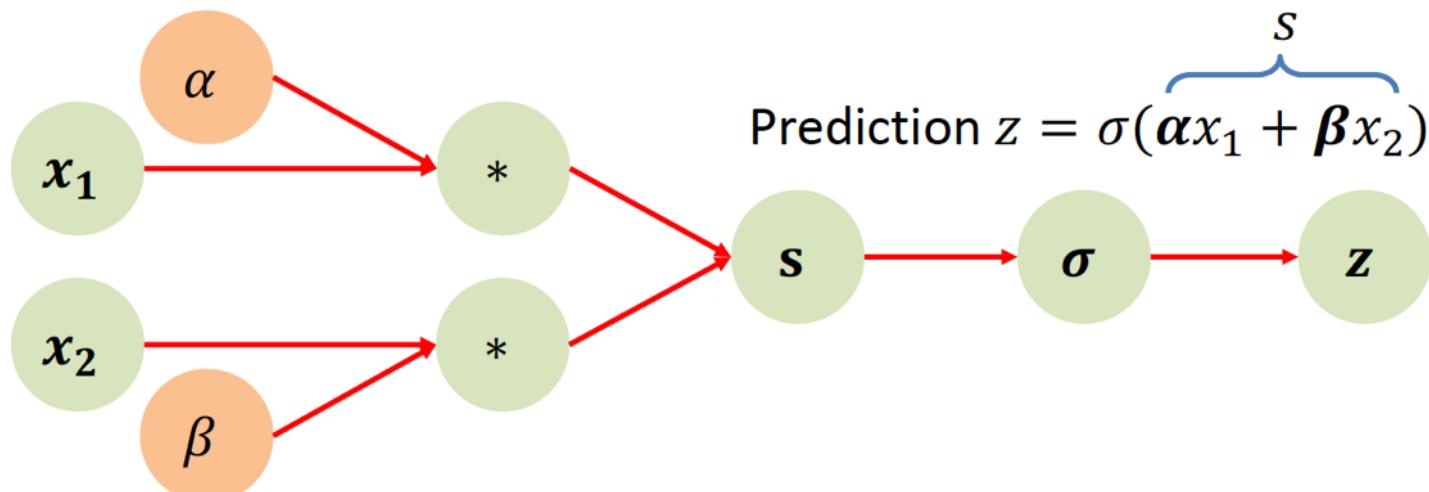
Back-propagation es un modo de diferenciacion automático

Es utilizado en todos los frameworks de AI hoy en dia.

Es muy rápido



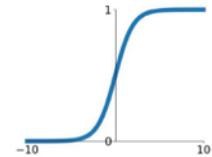
Funciones de Activación



$$\text{Prediction } z = \sigma(\alpha x_1 + \beta x_2 + s)$$

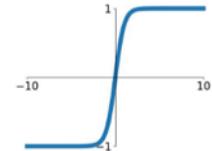
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



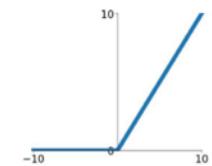
tanh

$$\tanh(x)$$



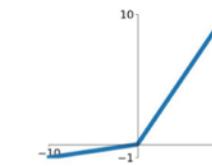
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

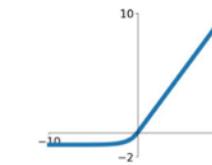


Maxout

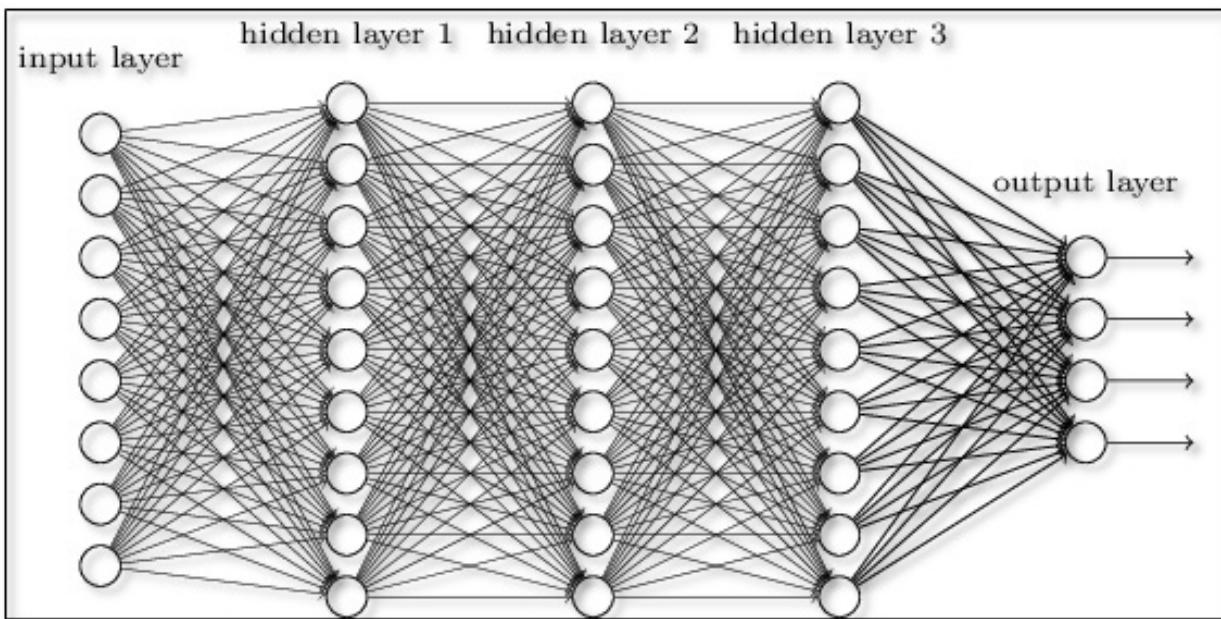
$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Deep Neural Networks



```
tf.keras.backend.clear_session()
```

```
model = Sequential()
```

```
model.add(Dense(hidden_nodes1, input_dim=input_nodes,  
activation='relu'))
```

```
model.add(Dense(10, activation='relu'))
```

```
model.add(Dense(10, activation='relu'))
```

```
model.add(Dense(10, activation='relu'))
```

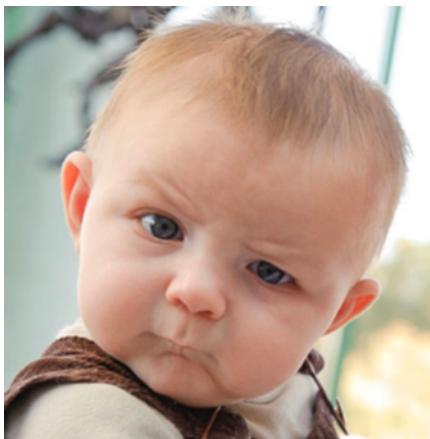
```
model.add(Dense(1, activation='sigmoid'))
```

```
model.compile(loss='binary_crossentropy',  
optimizer='adam',  
metrics=['accuracy'])
```

Por qué las NN son tan poderosas?

1. A MLP can approximate any function.

Cybenko, G. (1989) - Approximation by Superpositions of a Sigmoidal Function



Lo que se resume en que no importa que tan complejo es el decision boundary, un MLP siempre puede aproximar lo

(Cuando tienes datos y suficiente poder computacional)



Questions

[Home](#) > [Catalog](#) >

FREE COURSE

Intro to TensorFlow for Deep Learning

by  TensorFlow

This course is a practical approach to deep learning for software developers

[START FREE COURSE](#)

About this Course

Learn how to build deep learning applications with TensorFlow. This



RELATED NANODEGREE PROGRAM

Machine Learning Engineer

by  AWS

Get a Nanodegree certificate that accelerates your career! [»»»](#)

COURSE COST

Free

TIMELINE

Approx. 2 Months

SKILL LEVEL

 Intermediate

Primeros pasos en NLP



Bag of words and tokens



Lowercase



Stemming



Word embeddings

Filtros y Tokens

Cómo podemos representar una cadena de texto?

Evolution of the hyaluronan synthase (has) operon in
Streptococcus zooepidermidicus and other pathogenic
streptococci

Filtering

Evolution of the hyaluronan synthase **has** operon in
Streptococcus zooepidermidicus and othet pathogenic
streptococci

Tokenization

Evolution of the hyaluronan synthase has

Eliminar todos los caracteres innecesarios
(, . \$ % & / ! “ · < >) _ ^ * Ç .. ` + ‘

Separar las cadenas en palabras individuales

Lowercase & Stemming

Lowercase

Raw	Lowercased
Canada	canada
Canada	canada
CANADA	canada
TOMCAT	tomcat
Tomcat	tomcat
toMcat	tomcat

Stemming

	words	stemmed words
0	connect	connect
1	connected	connect
2	connection	connect
3	connections	connect
4	connects	connect

Bag of Words

Journal of Artificial Intelligence Research

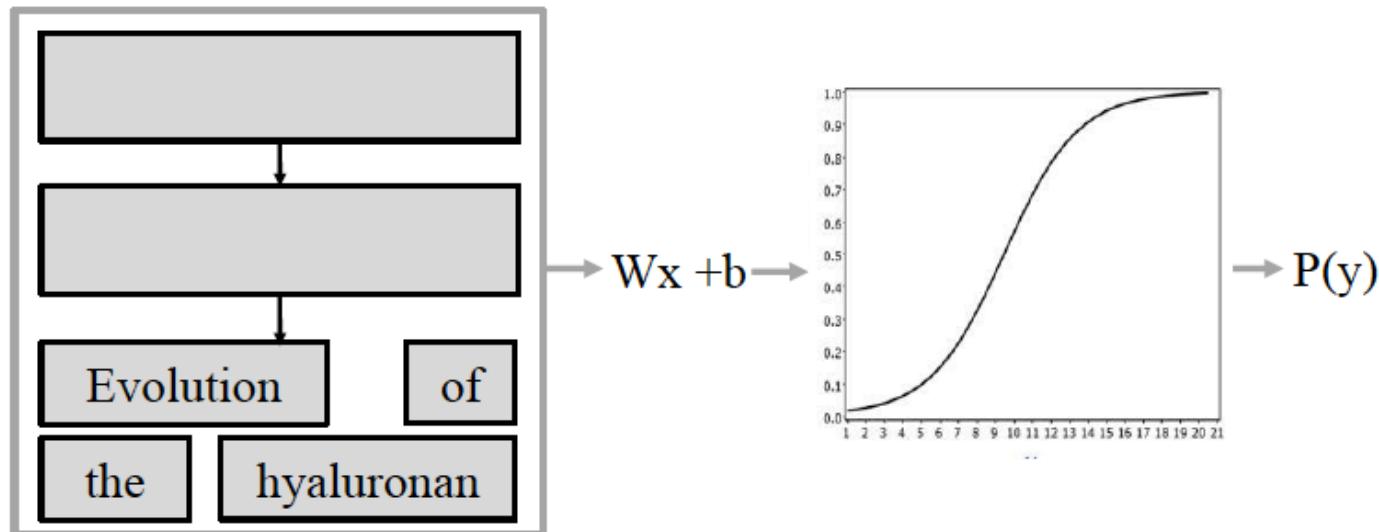
JAIR is a refereed journal, covering the areas of Artificial intelligence, which is distributed free of charge over the Internet. Each volume of the journal is also published by Morgan Kaufmann

0	learning
3	journal
2	intelligence
0	text
1	Internet
...	...



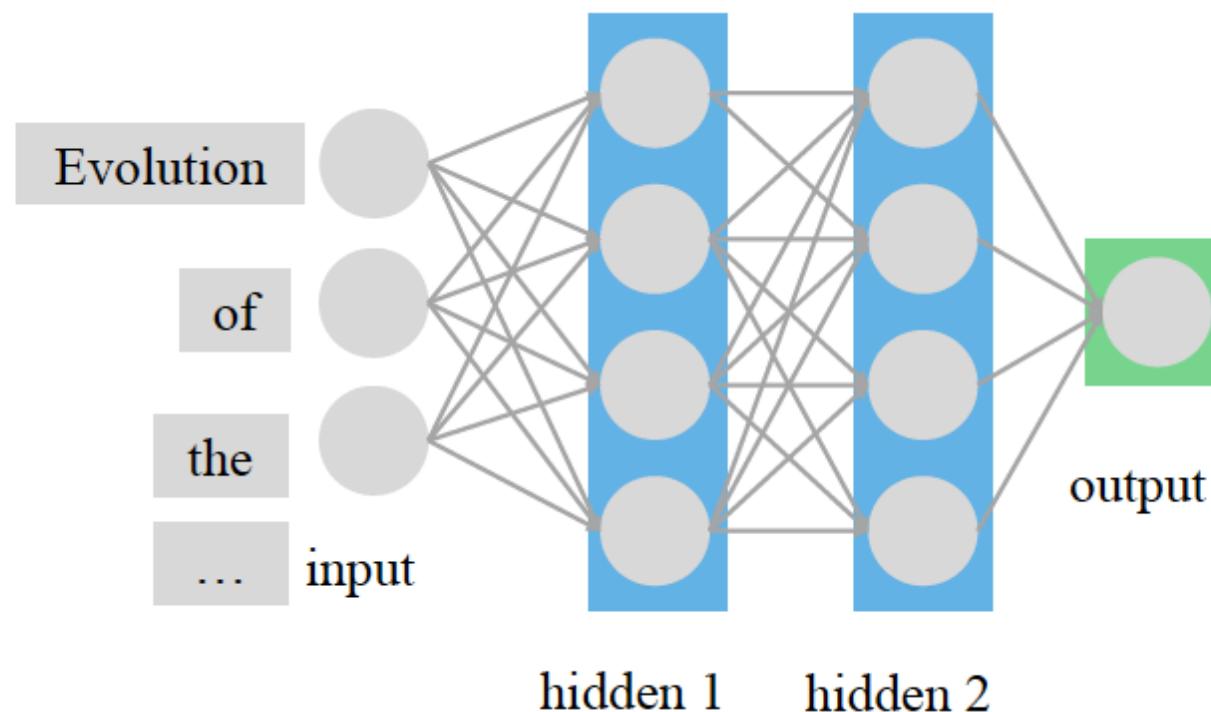
id	T1	T2	T3	T4	T5	.	.	.	n
1	1	1	1	1	1	1	0	0	0
2	1	1	1	0	1	1	1	0	0
3	1	1	1	0	1	0	0	1	1
4	1	1	1	0	1	0	0	1	0

Log reg + BoW



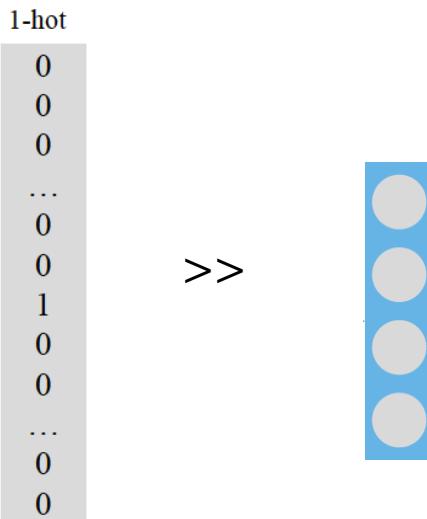
Cuantas variables tendrá este modelo?

NN + BoW

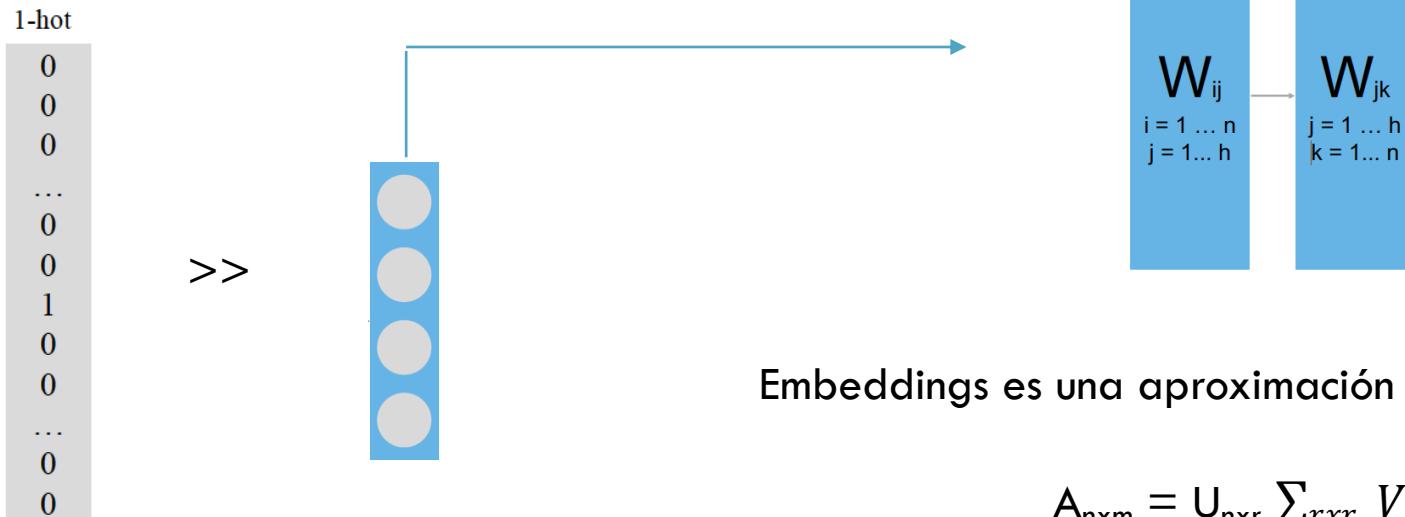


Cuantas variables tendrá este modelo?

Embeddings



Embeddings



Embeddings es una aproximación al teorema de SVD

$$A_{n \times m} = U_{n \times r} \sum_{r \times r} V^T_{r \times m}$$

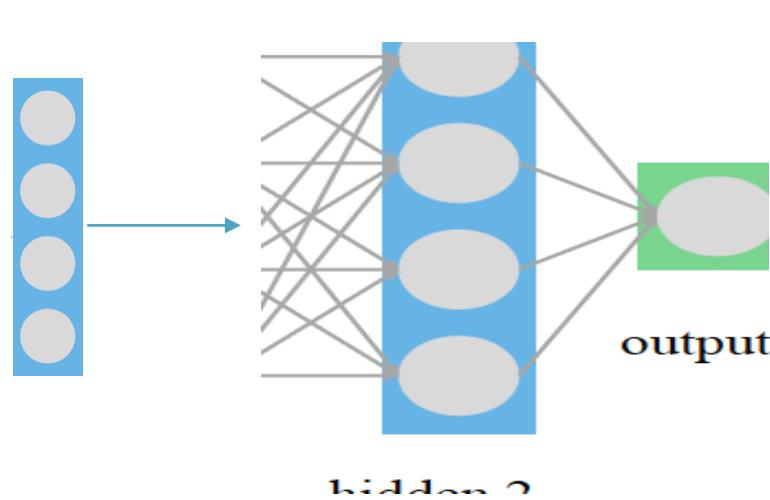
Donde $\sum_{r \times r}$ (matríg diagonal) representa la decomposicion latente de la matrix original $A_{n \times m}$. Y los r factores de la matriz son denominados factores latentes, SVD asegura la existencia de $\sum_{r \times r}$.

Sentiment Analysis

All reviews:

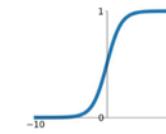


Filtering
Stemming
Lowercase
BoW



Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

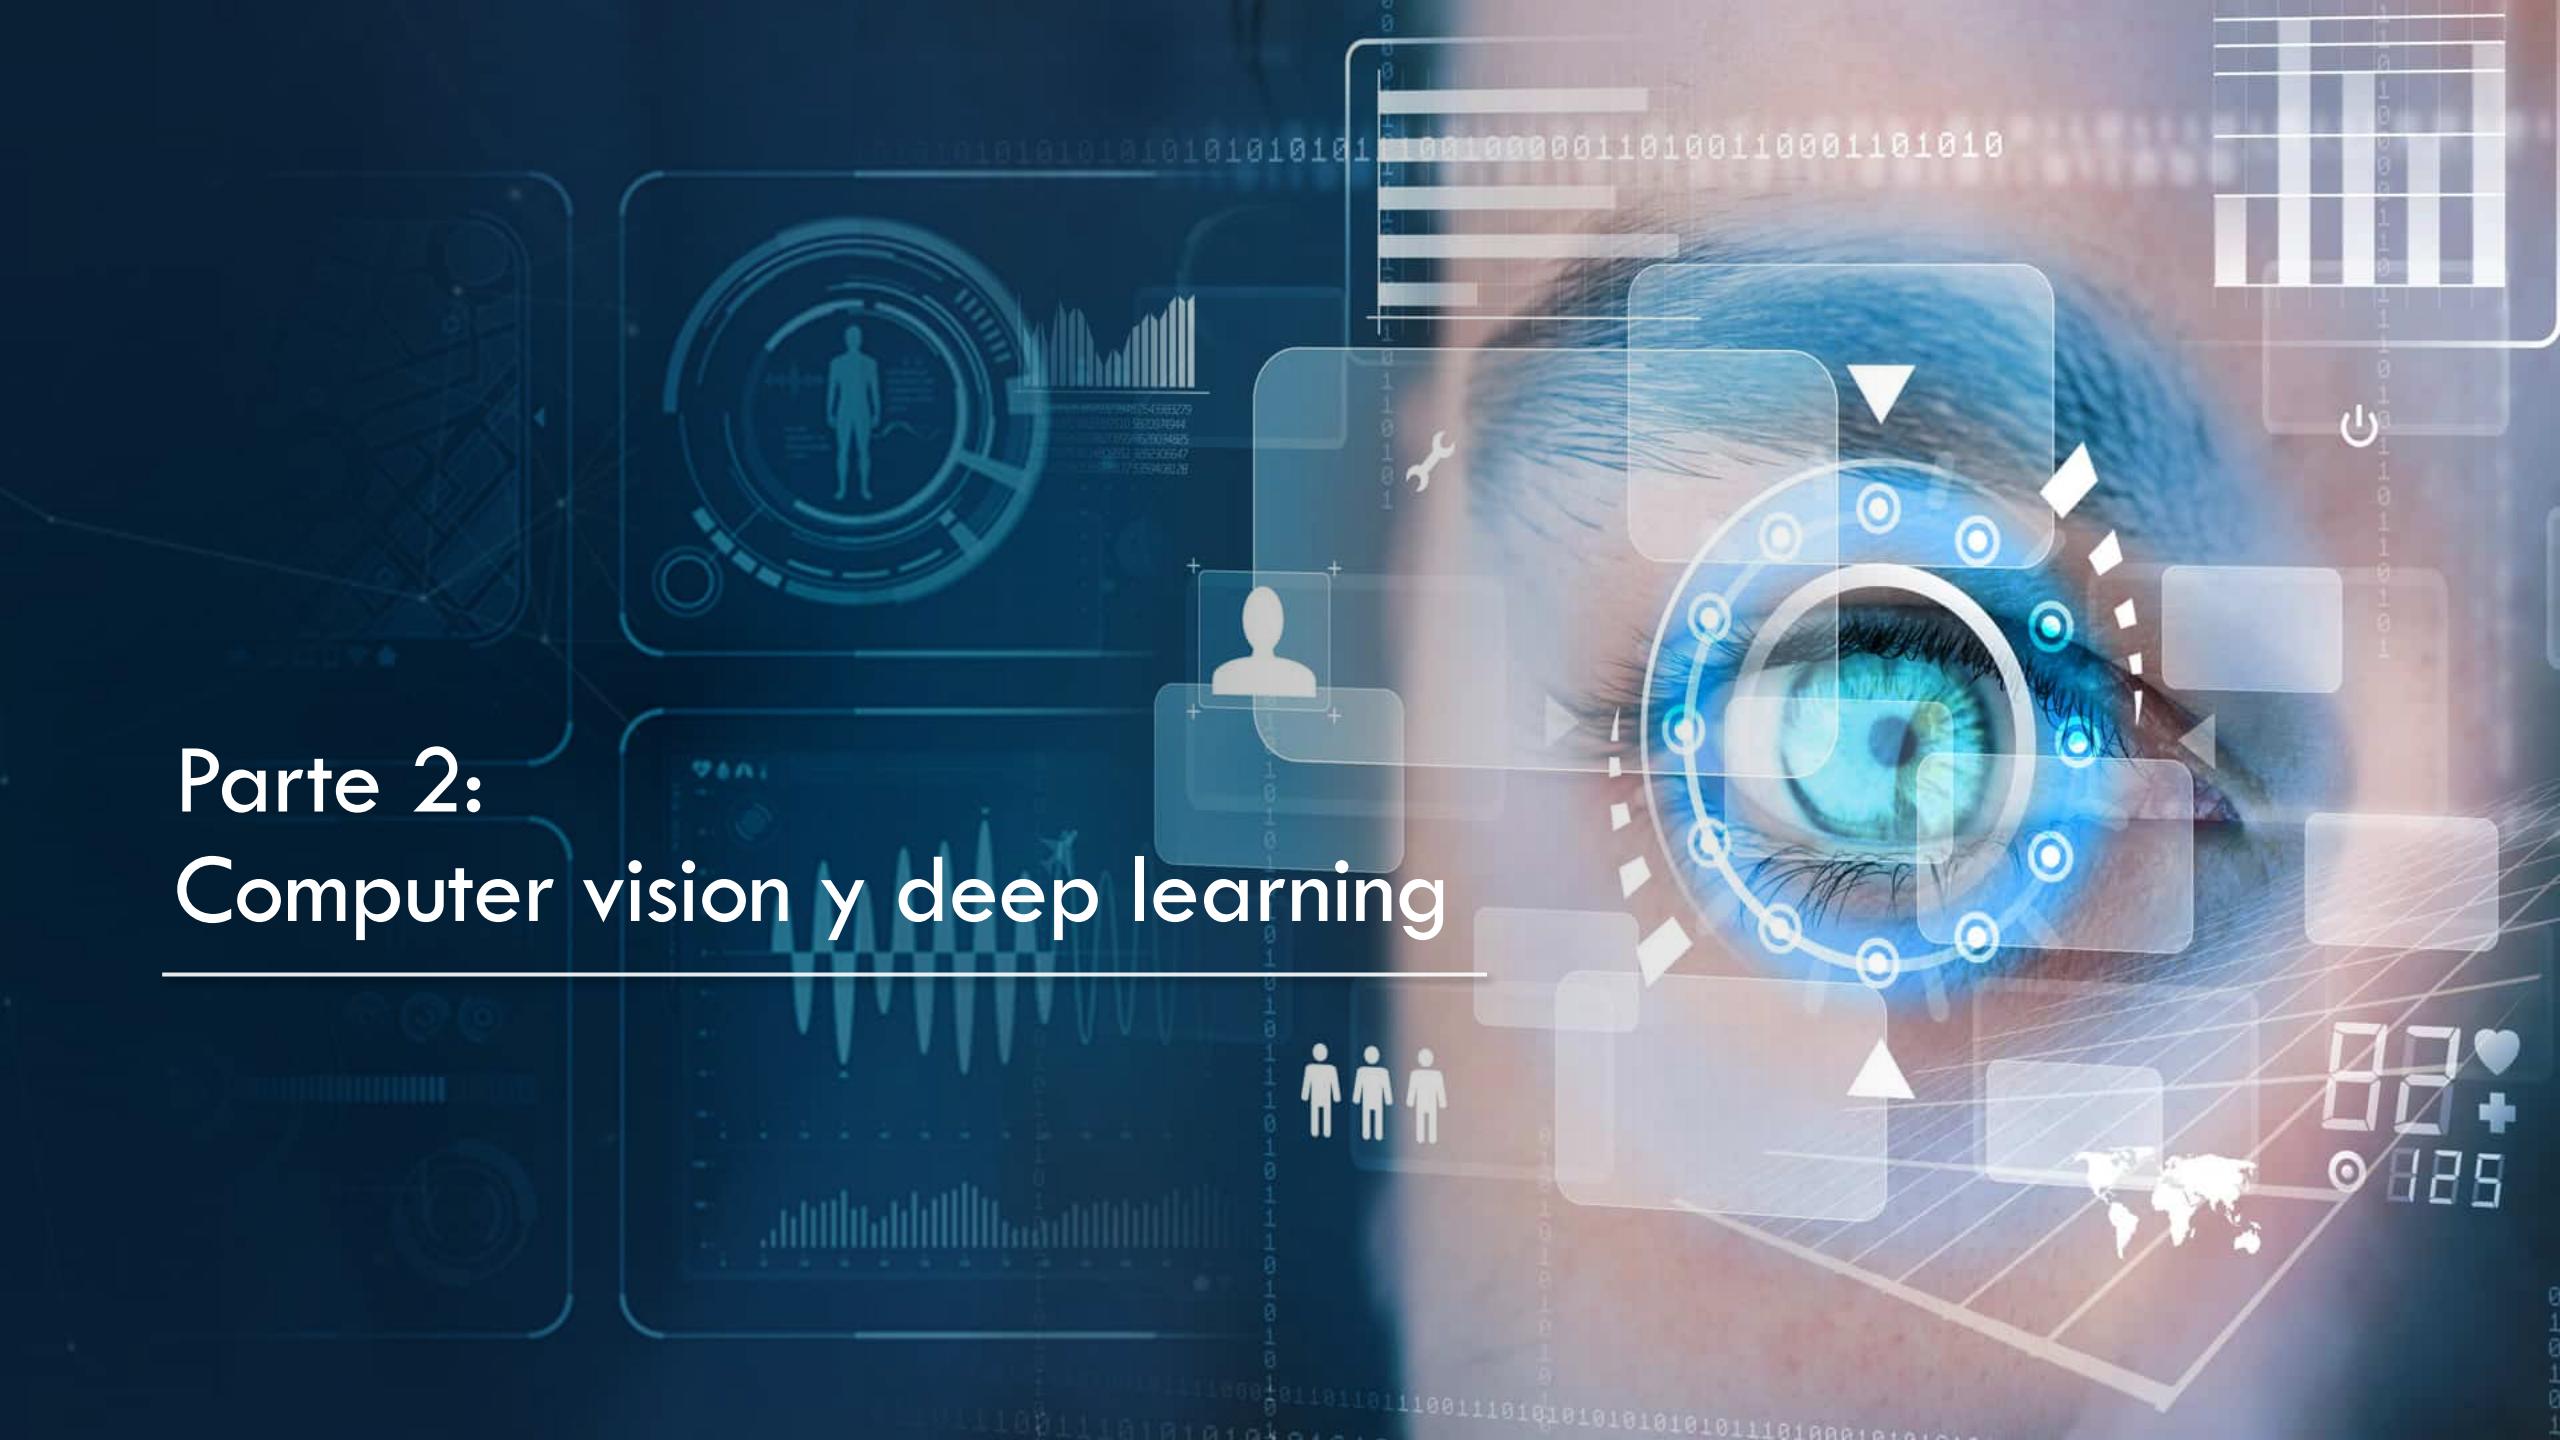


Demo 1

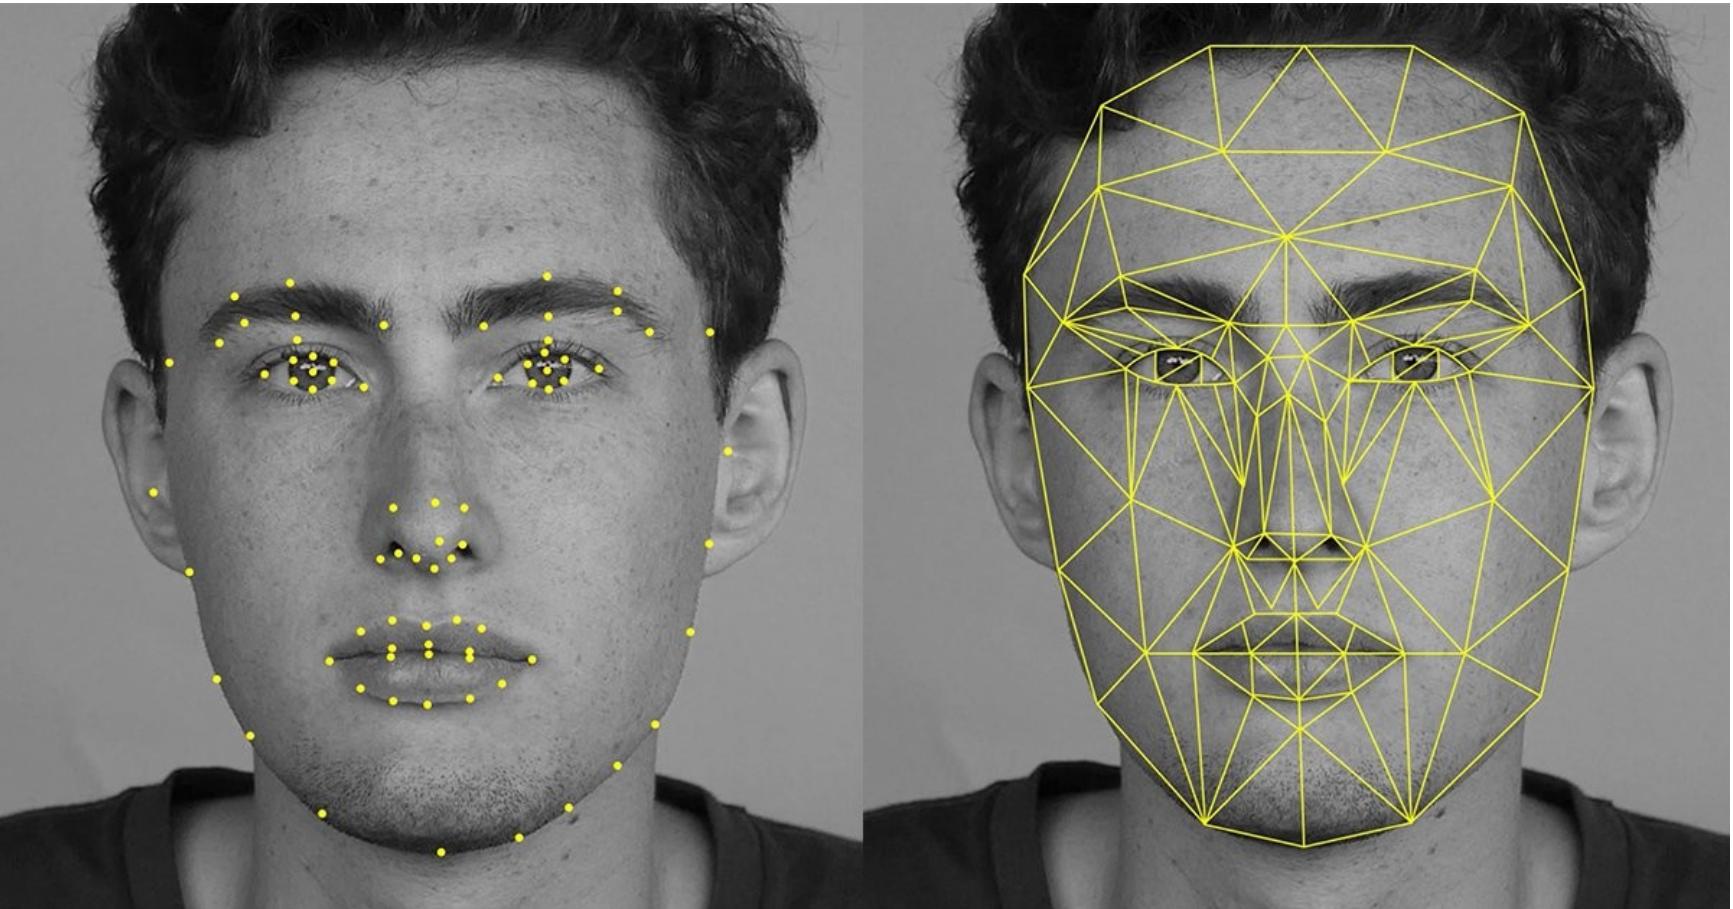
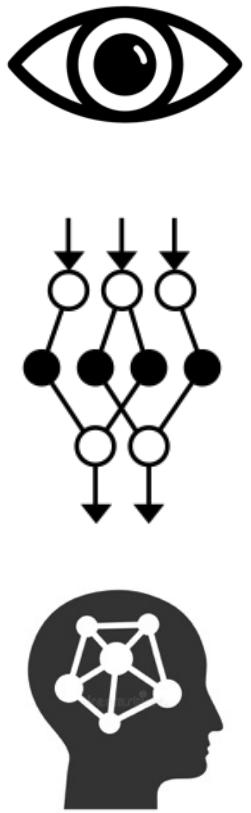
NN In

Tensorflow

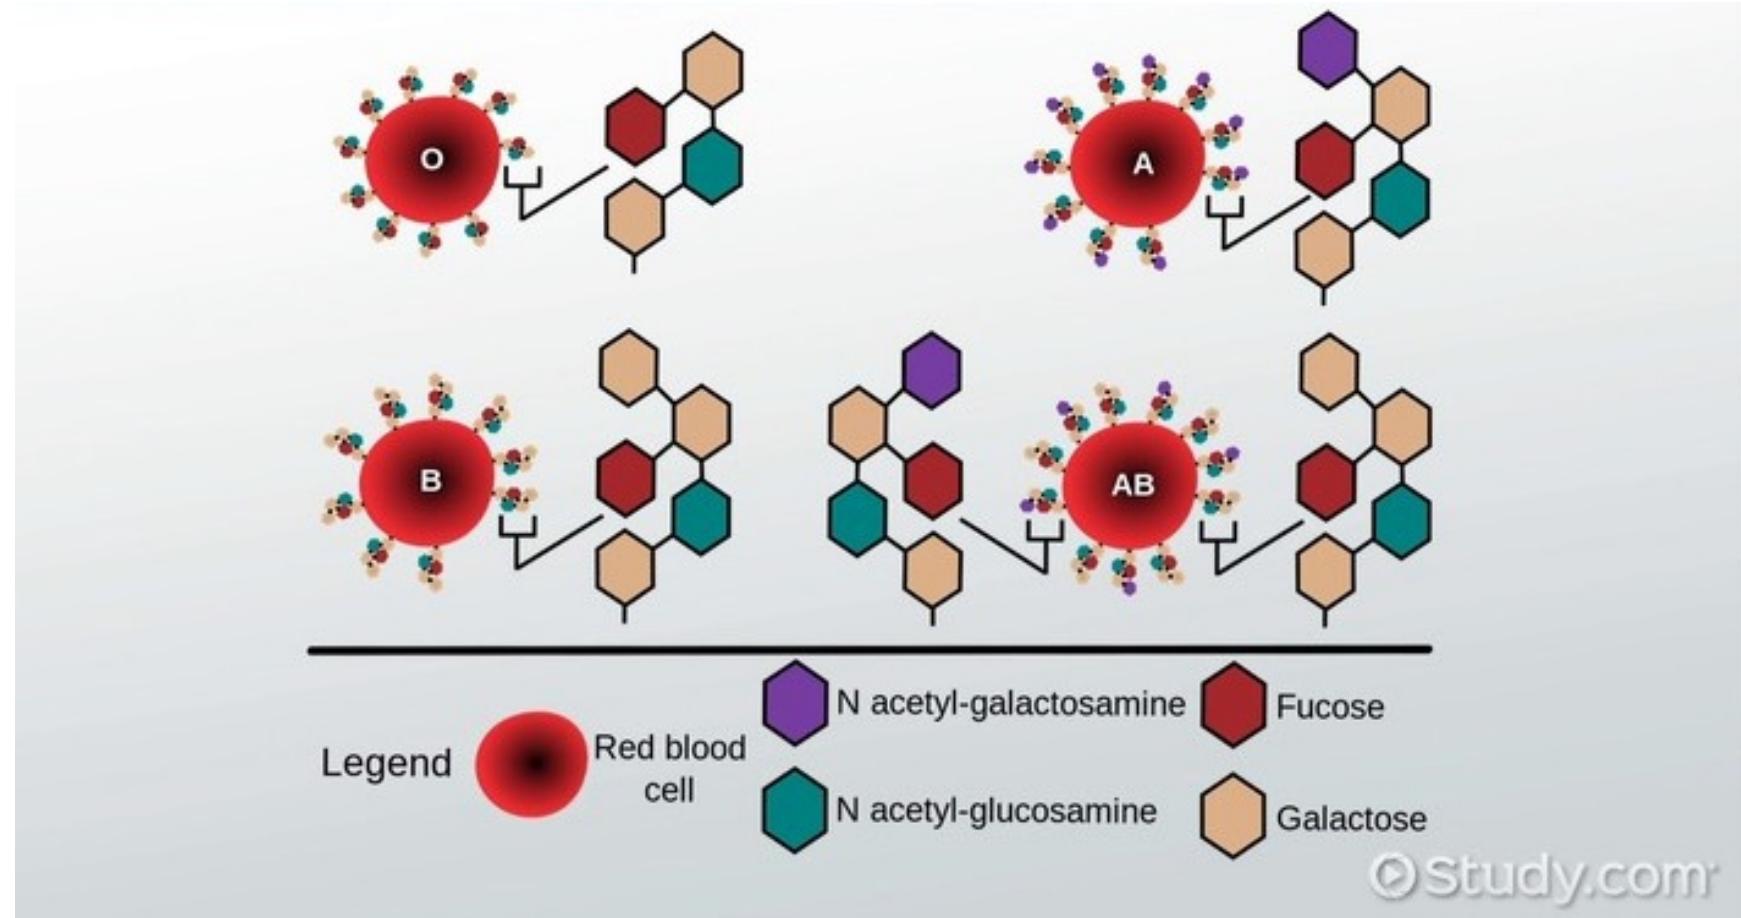
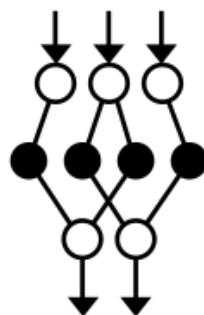
Parte 2: Computer vision y deep learning



Facial detection & recognition



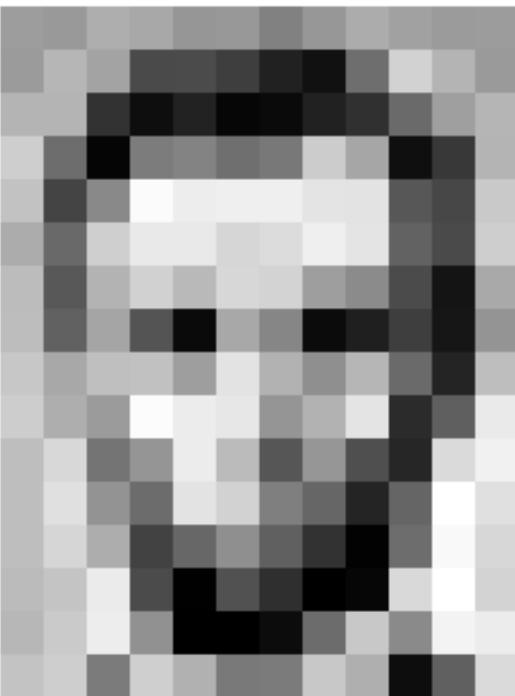
Infection detection & recognition





Cómo ve una computadora?

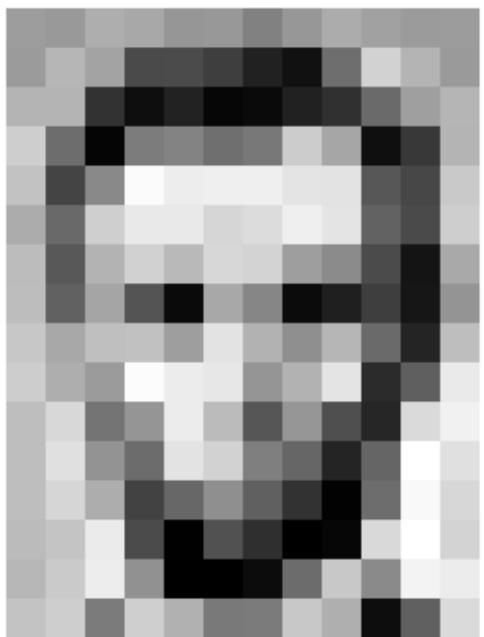
Las imágenes son matrices numéricas



157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	94	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	209	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	94	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	209	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

Clasificación

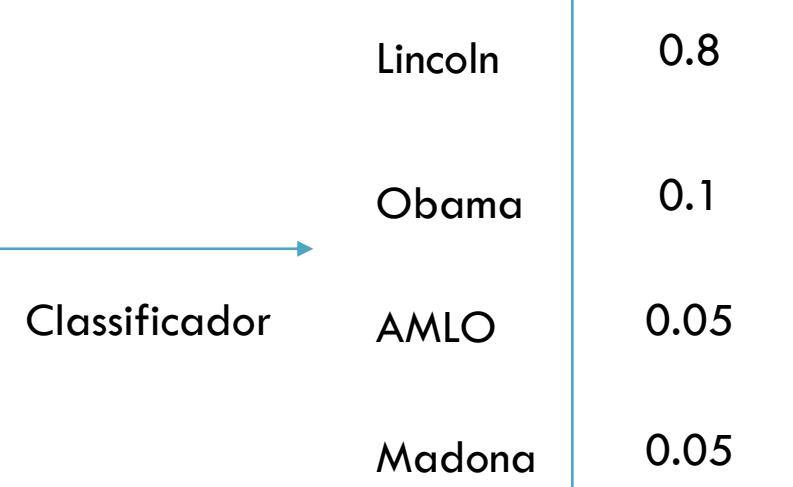


Input

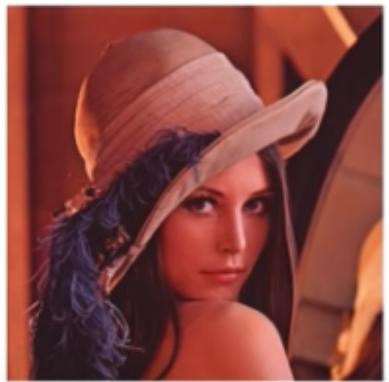


157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	239	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
196	206	123	207	177	121	123	200	175	13	96	218

Representación



Variables



Naríz
Ojos
Boca



Llantas
Placa
Luces



Ventanas
Puerta
Jardín

Variables



Naríz
Ojos
Boca



Llantas
Placa
Luces



Ventanas
Puerta
Jardín

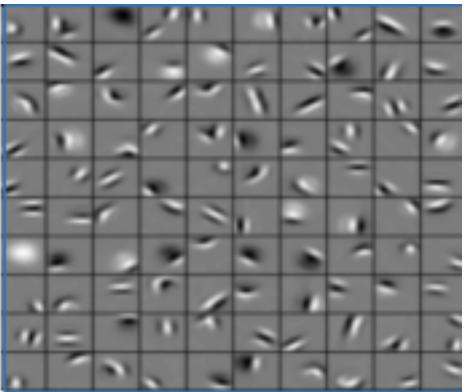
Opcion1: Usar domain Knowledge

- Crear
- Probar
- Seleccionar variables.

Cuál es el problema?

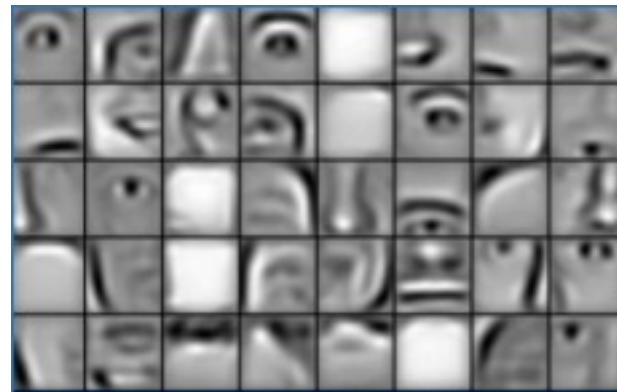
Neural Networks

Low Level



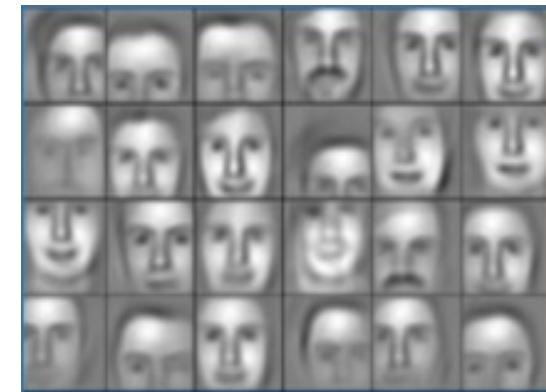
Hidden layer 1

Mid Level



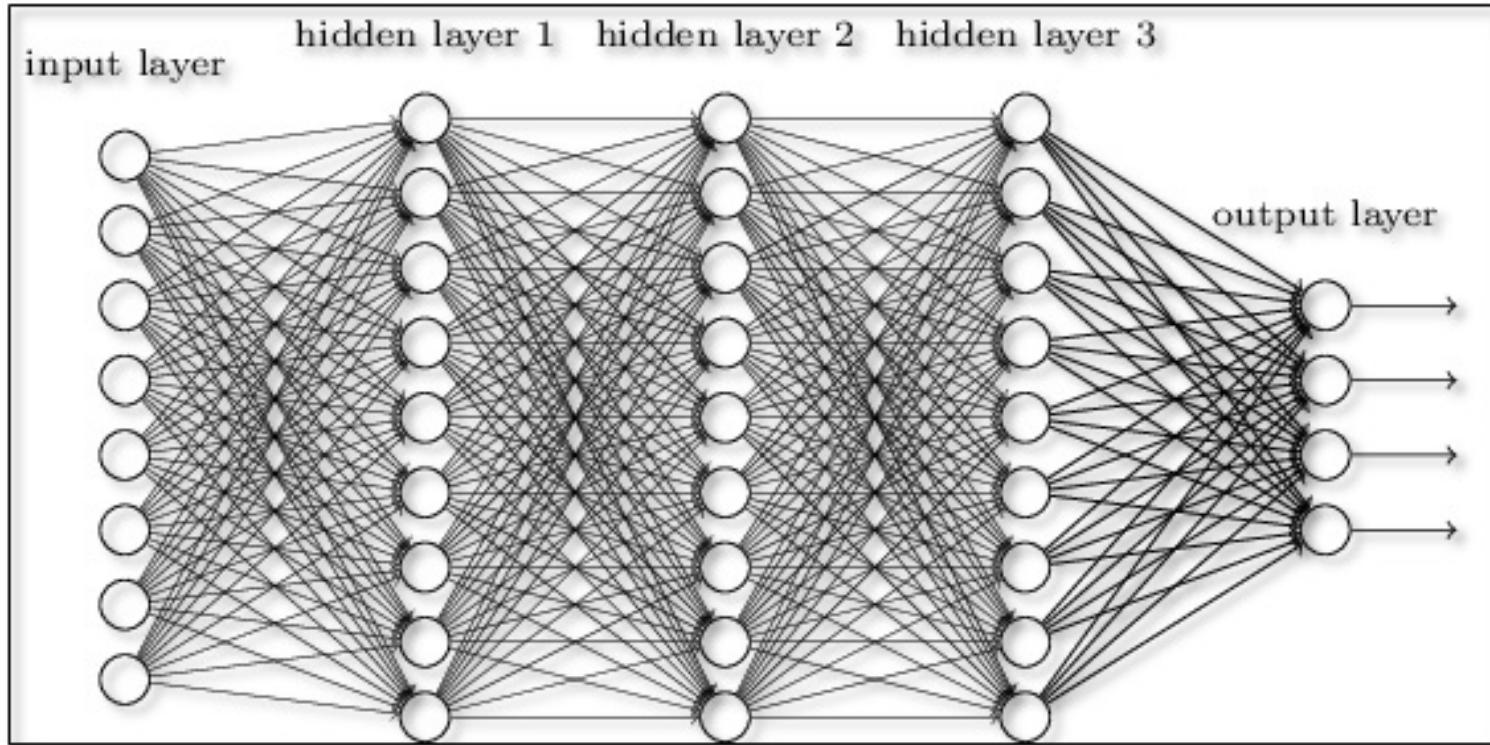
Hidden layer 2

High Level features



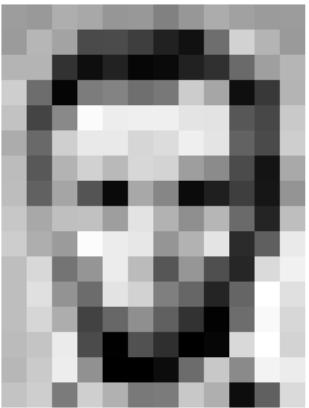
Hidden layer 3

Obteniendo Variables

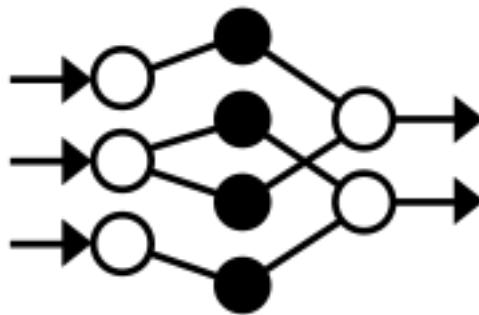


Deep Neural Network

2D Input



[255 x 255]

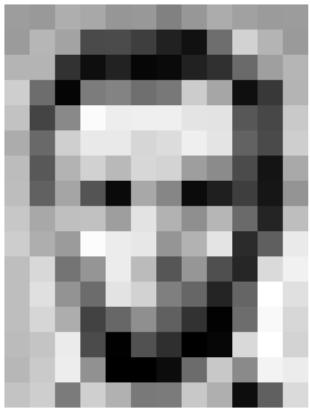


[65,025 N]

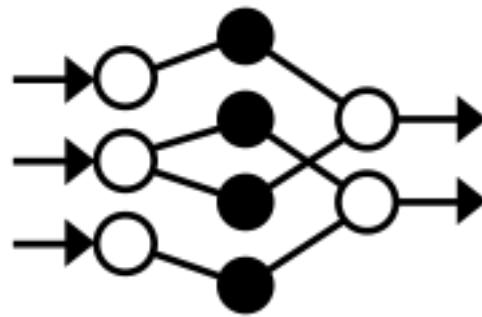
FNN

[10]

2D Input



[255 x 255]



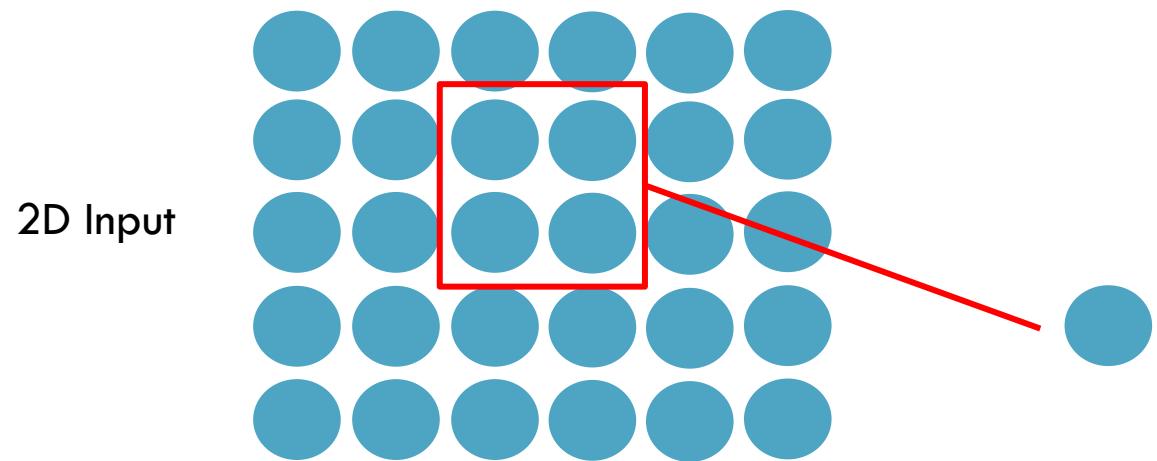
[65,025 N]

FNN

[10]

Qué otras opciones hay?

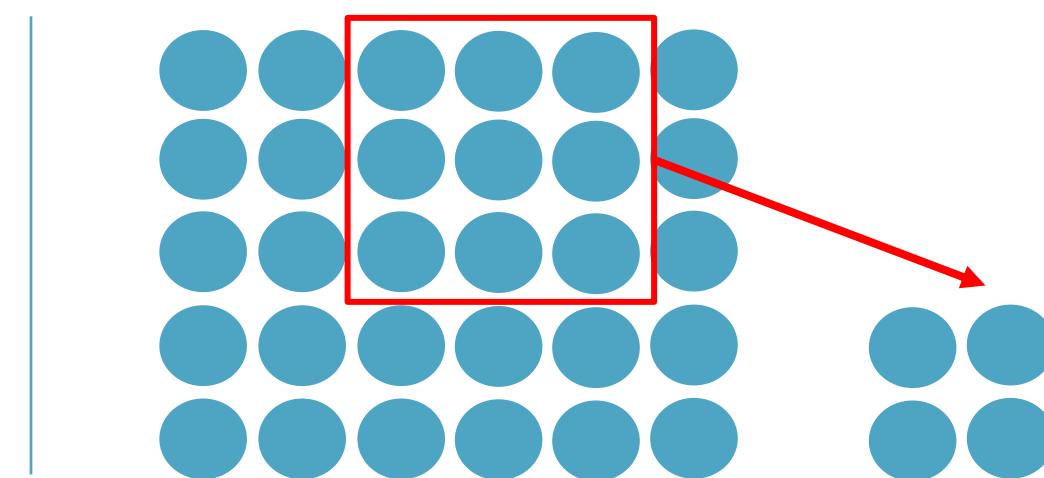
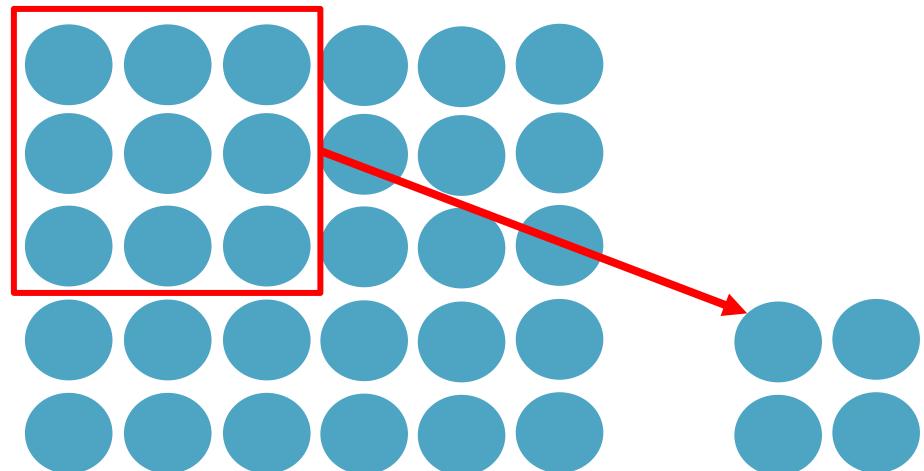
Utilizando la estructura espacial



Idea:

Connectar subconjuntos de datos que afecten un solo nodo de la red.

Utilizando la estructura espacial



Utilizamos un desplazamiento (slide) para obtener features de la imagen original

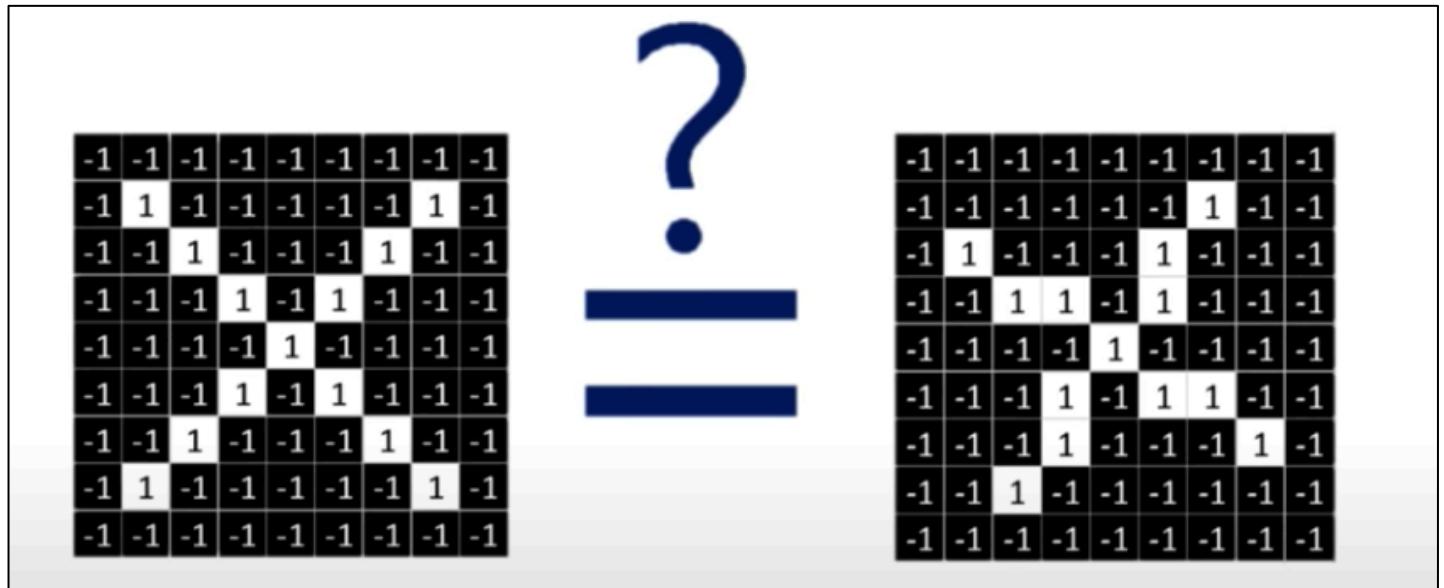
En la practica, ésta operación en Rojo se llama convolución con un filtro de 3x3.

Obteniendo Variables con Convoluciones

X vs X

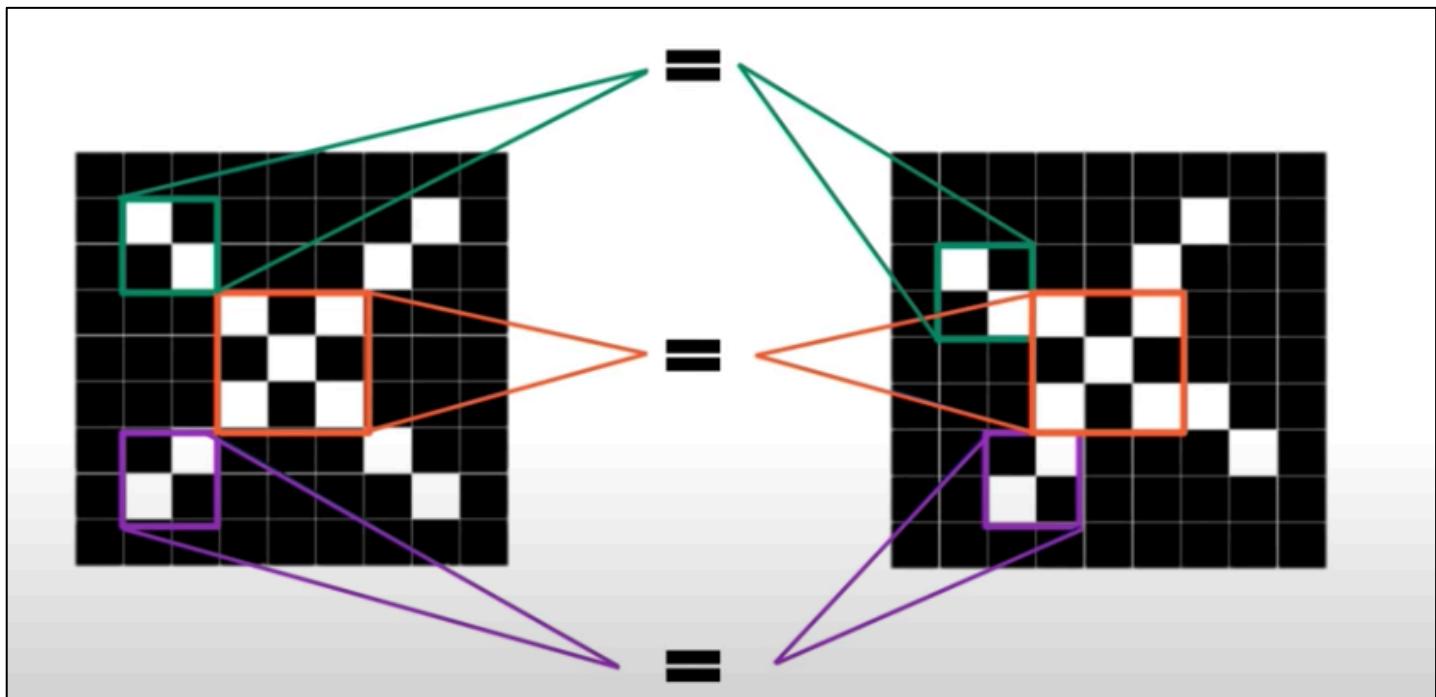
La imagen es representada como una matriz de pixeles...

Queremos clasificar una imagen como X incuso si está distorsionada o invertida,

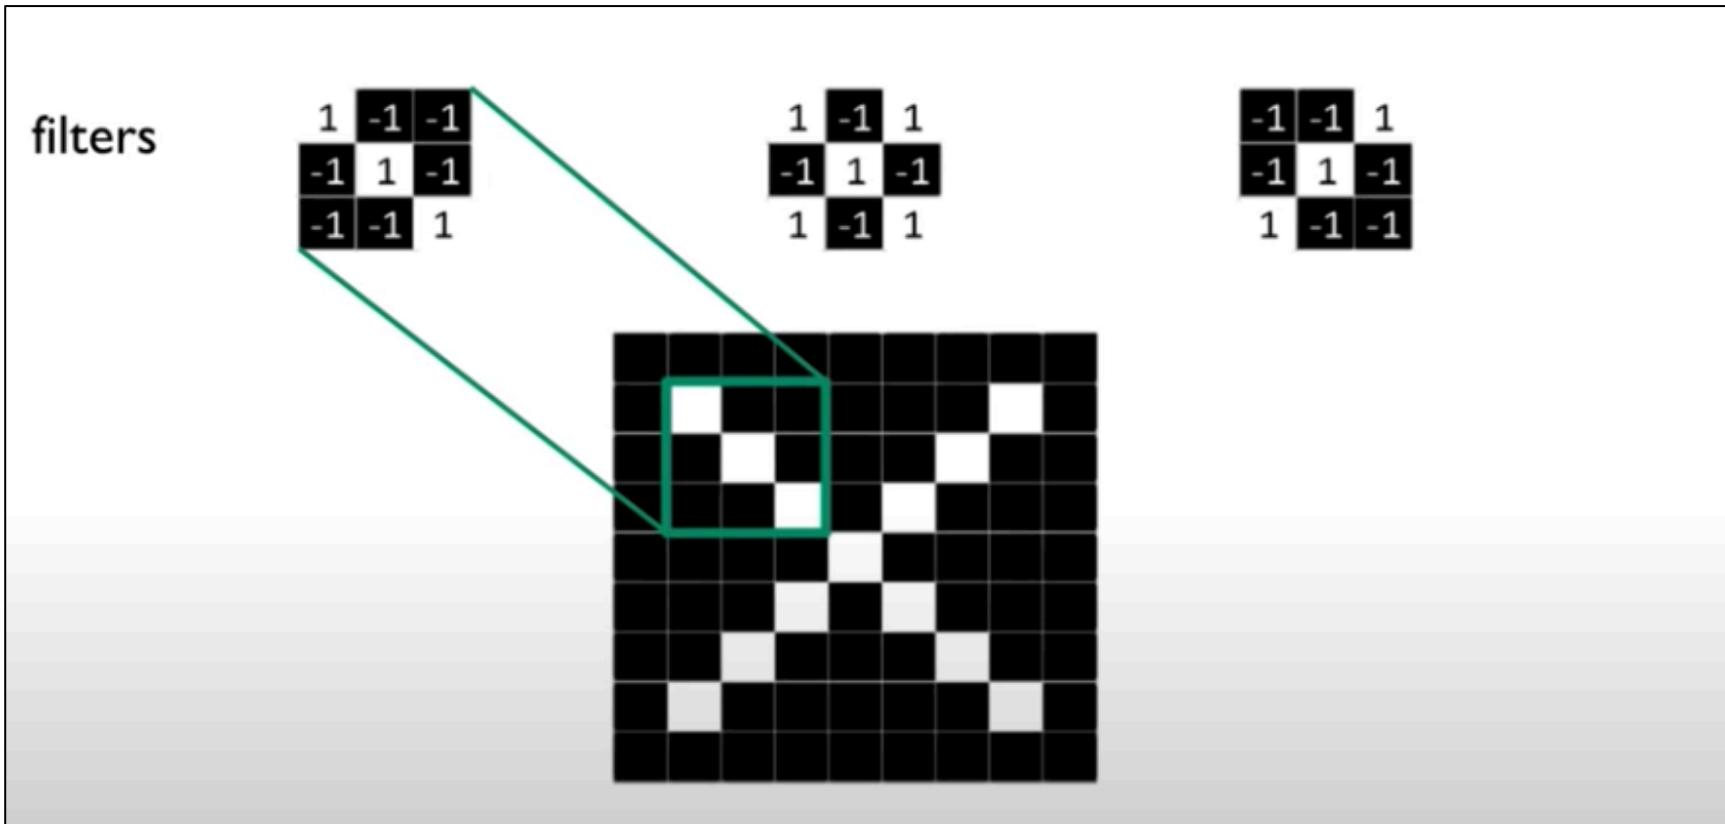


X vs X

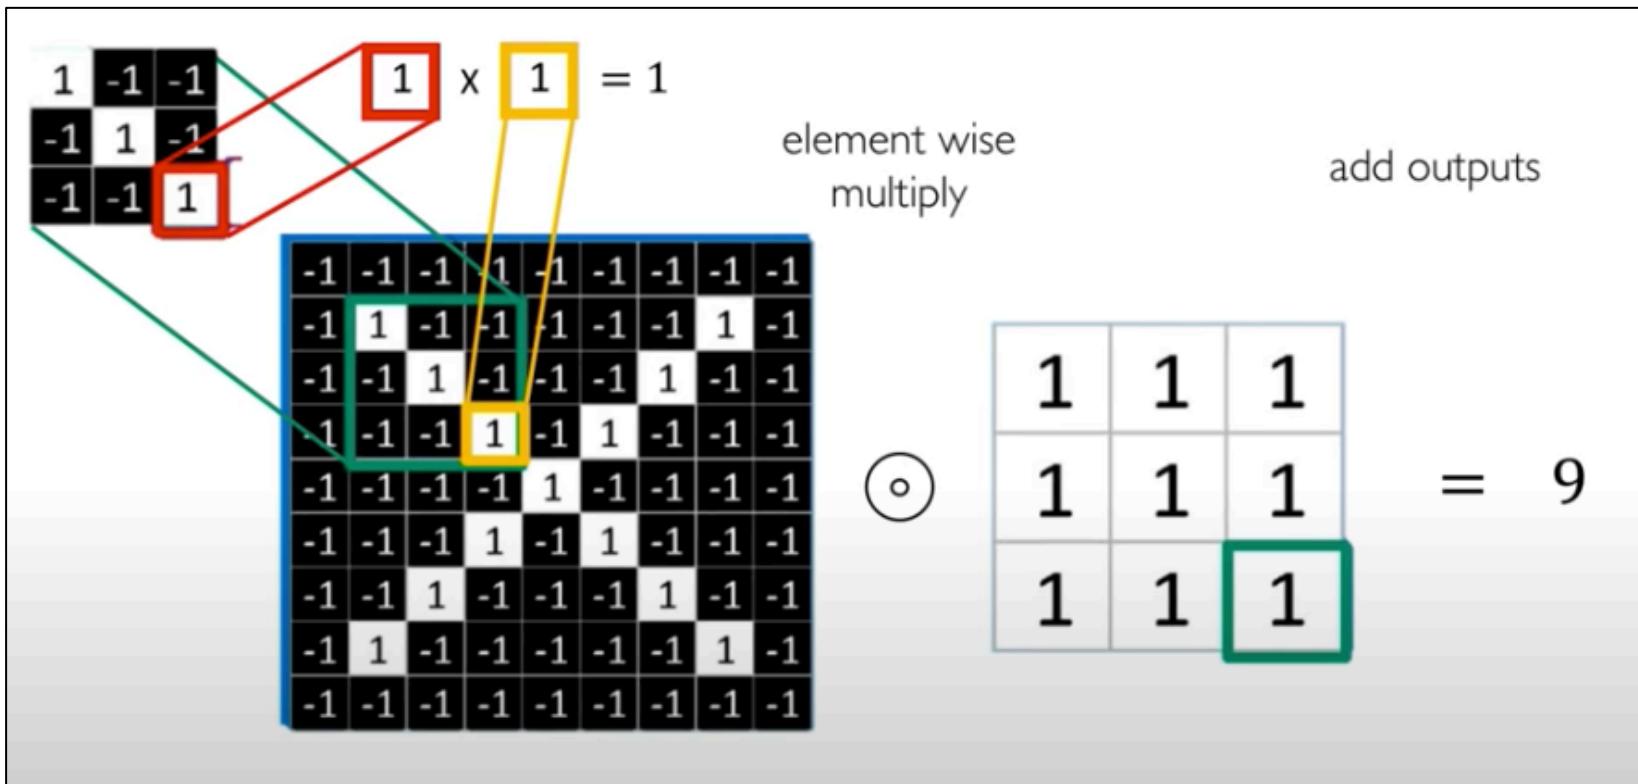
Lo que queremos es obtener variables que tengan similitud considerando la estructura.



Utilizamos filtros para detectar patrones

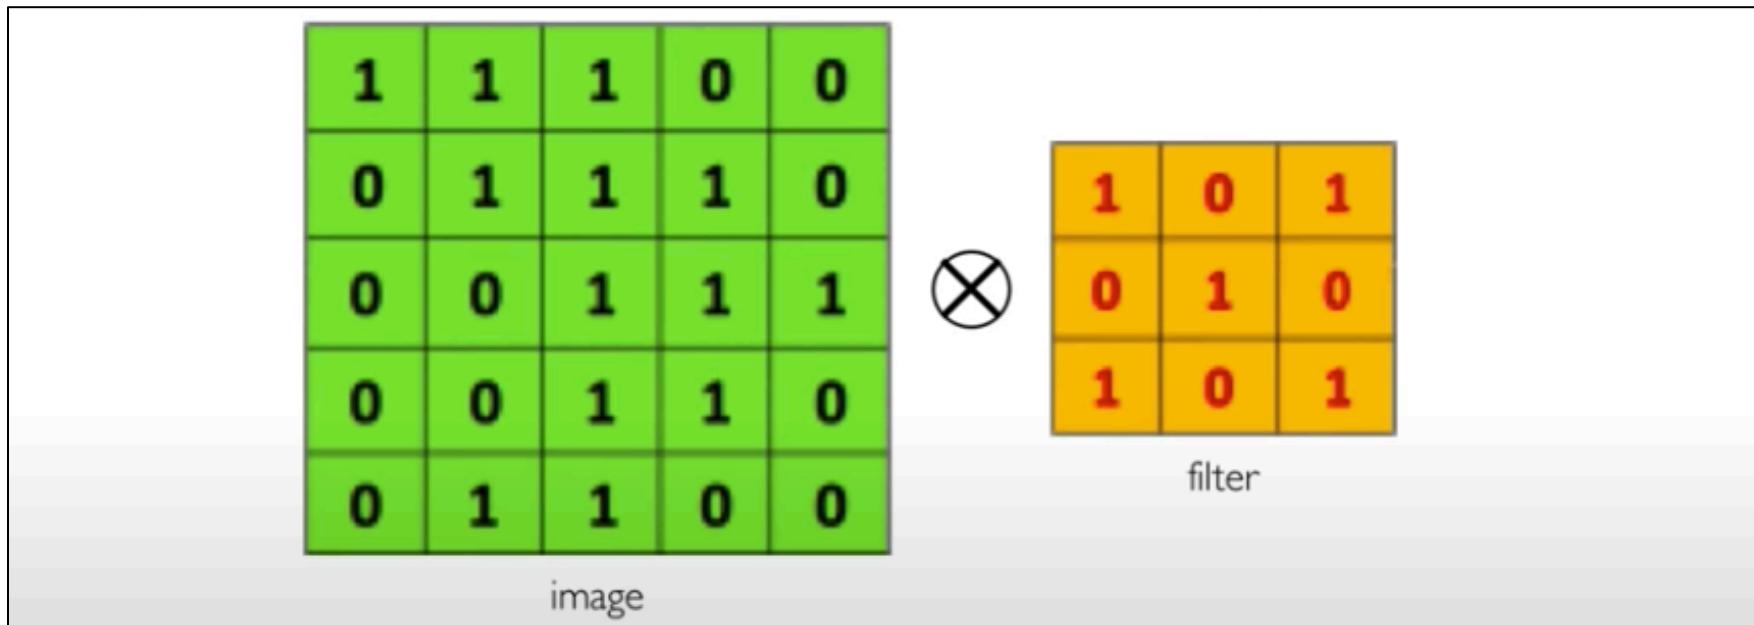


Utilizamos filtros para detectar patrones



Utilizamos filtros para detectar patrones

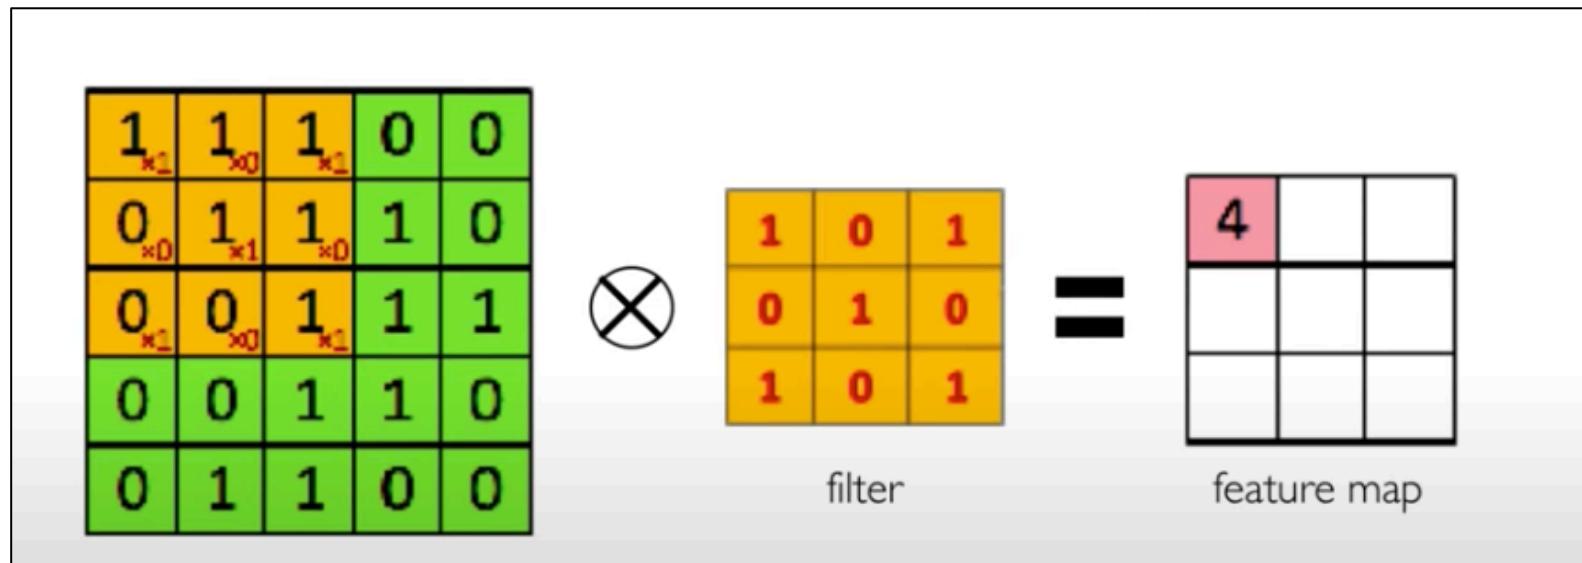
Vamos a calcular la convolución de una imagen de 5x5 con un filtro de 3x3



Vamos a deslizar el filtro de 3x3 sobre la imagen, multiplicar cada elemento y sumar.

Utilizamos filtros para detectar patrones

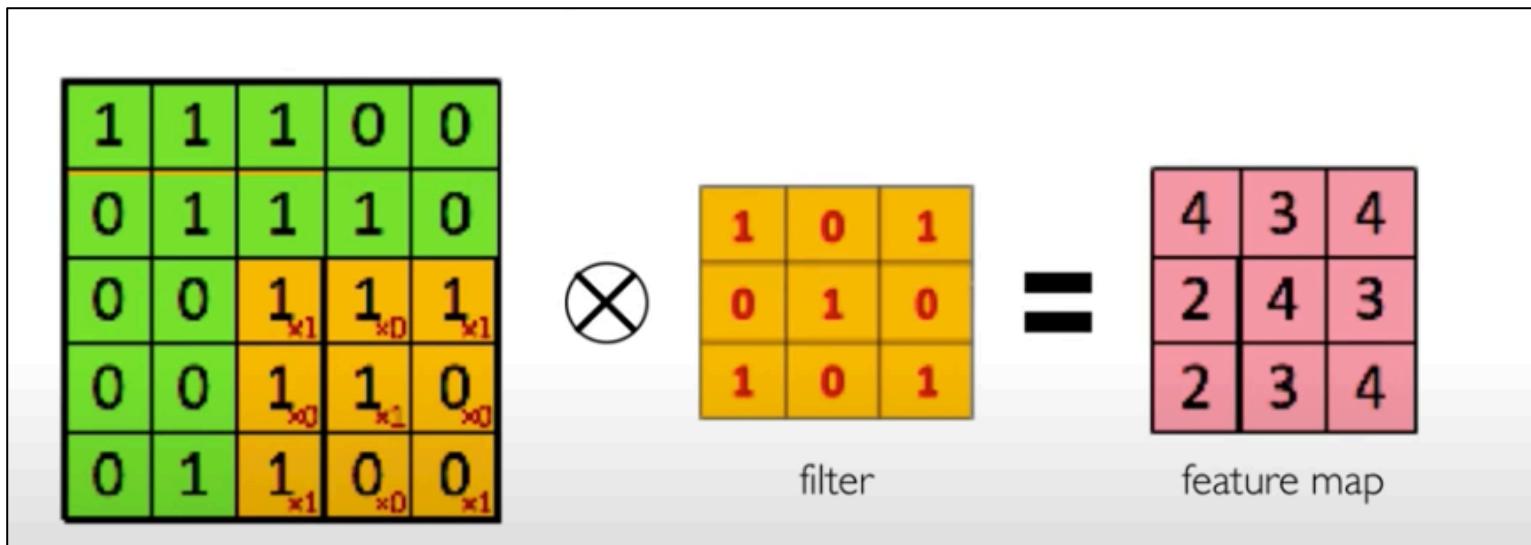
Vamos a calcular la convolución de una imagen de 5x5 con un filtro de 3x3



Vamos a deslizar el filtro de 3x3 sobre la imagen, multiplicar cada elemento y sumar.

Utilizamos filtros para detectar patrones

Vamos a calcular la convolución de una imagen de 5x5 con un filtro de 3x3



El mapeo de variables (feature map) nos indica en donde hay datos relevantes en la imagen original

Utilizar filtros diferentes



Original



Sharpen



Edge Detect

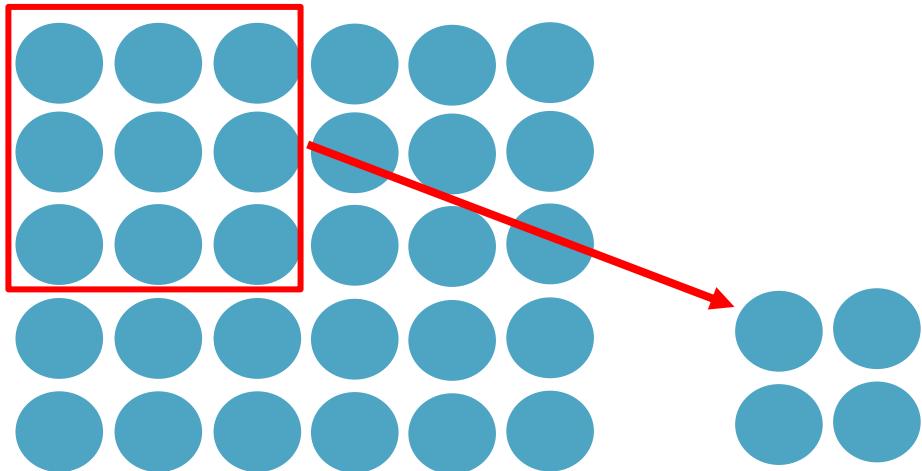


"Strong" Edge
Detect

Resumen:

1. Aplicar filtros nos ayuda a extraer variables locales
2. Podemos utilizar diferentes filtros para obtener multiples variables.
3. Las convoluciones "comparten" información de la imagen original

Convolución

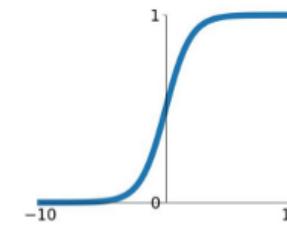


Convolución → $\sum_i \sum_j w_{ij} + x_{i+p, j+q} + b$

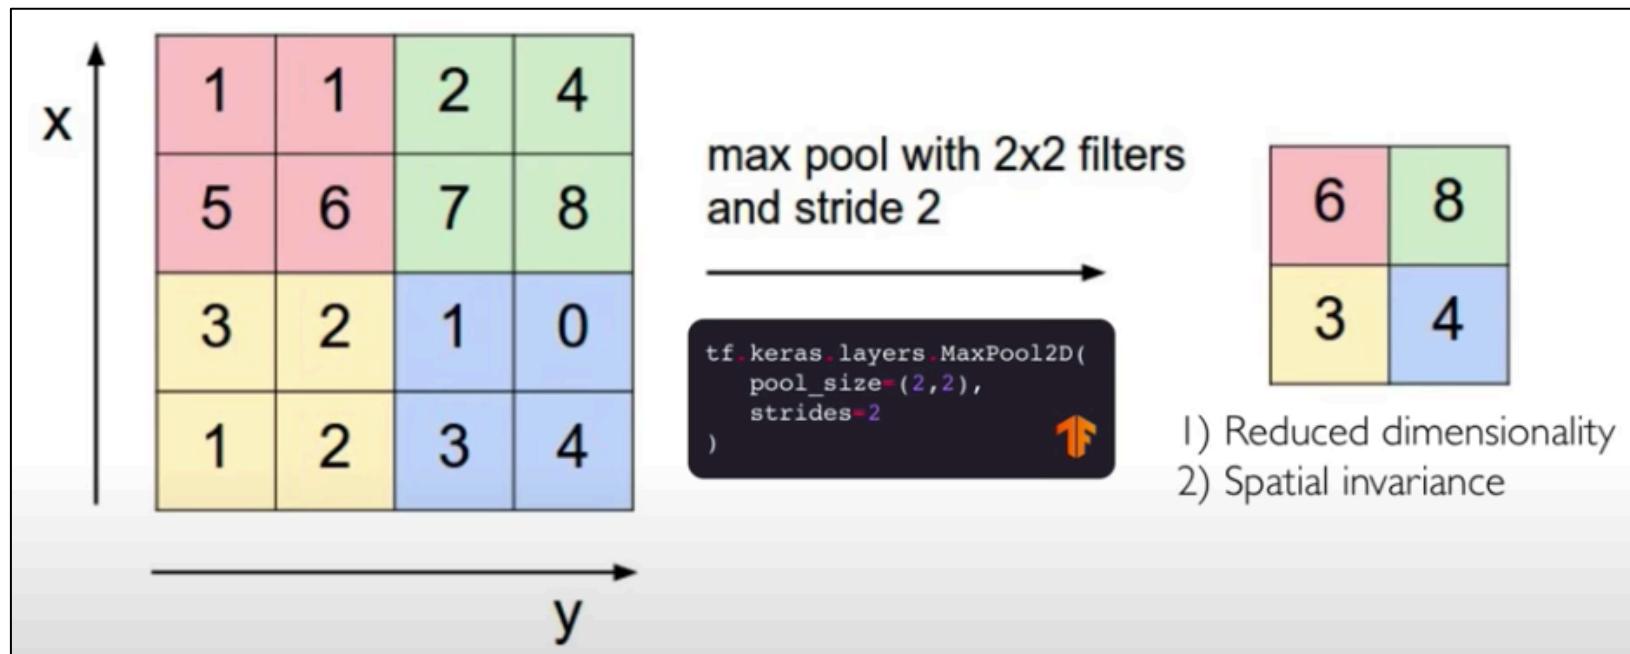
Para neuronas p, q en hidden layer

Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



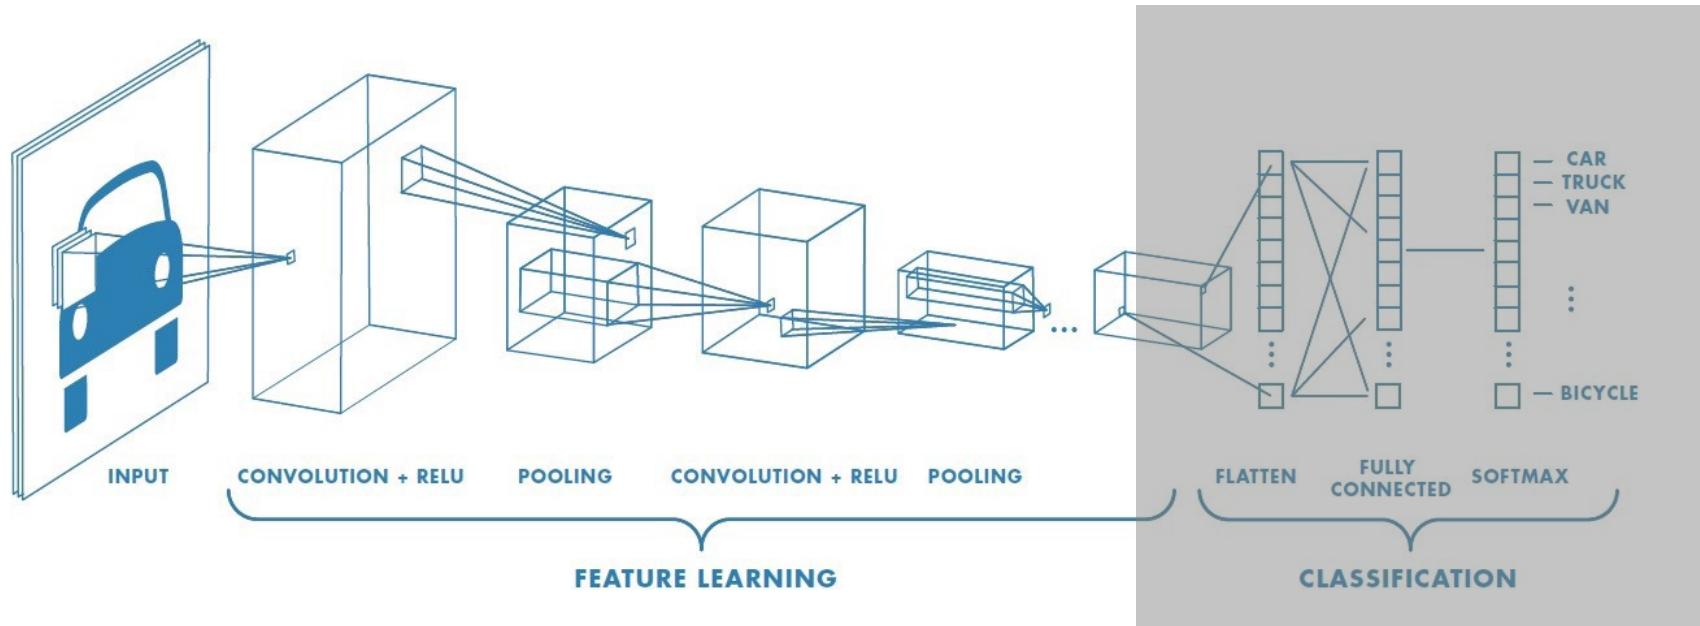
Maxpooling



Cuál es otra forma de reducir la dimensión después de una convolución?

Convolutional Neural Networks (CNNs)

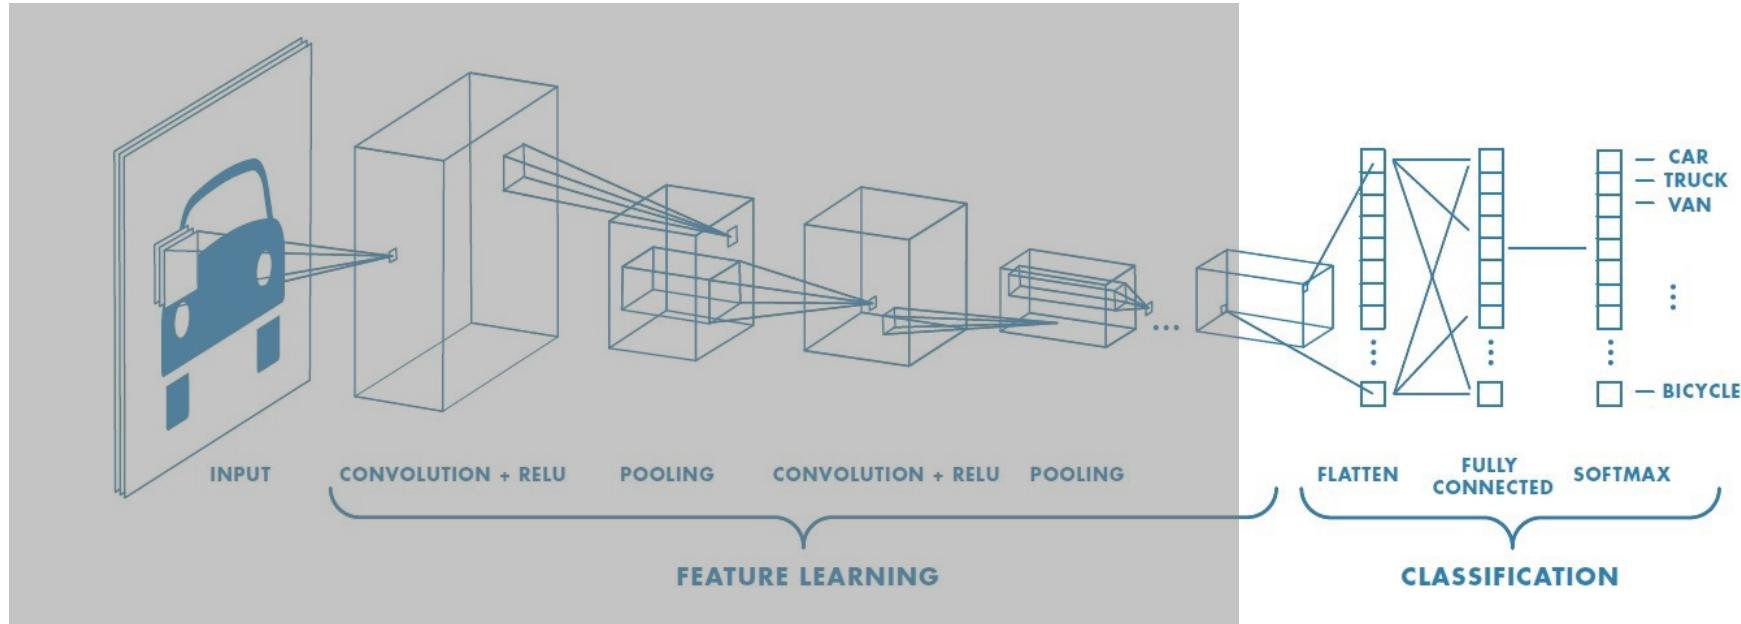
CNN



1. Convoluciones aplican filtros para general feature maps.
2. Aplicar función de activación no linear como ReLU.
3. Subsampling via maxpooling para reducir la dimención de los features maps.

```
tf.keras.layers.Conv2D  
tf.keras.activations.*  
tf.keras.layers.MaxPool2D
```

CNN



4. Fully connected layers sobre las variables que obtuvimos de las convucciones.
5. Función de activación sigmoidal para obtener una probabilidad.

$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

CNN In Tensorflow



Hand bag

Shopping Bags

Shopping Bag

Shopping Bags

Context classification

Object recognition for
Self-driving cars.

—

—

—

—

—

—

—

—

ATION FLOW

LANE LINES

LANE LINES

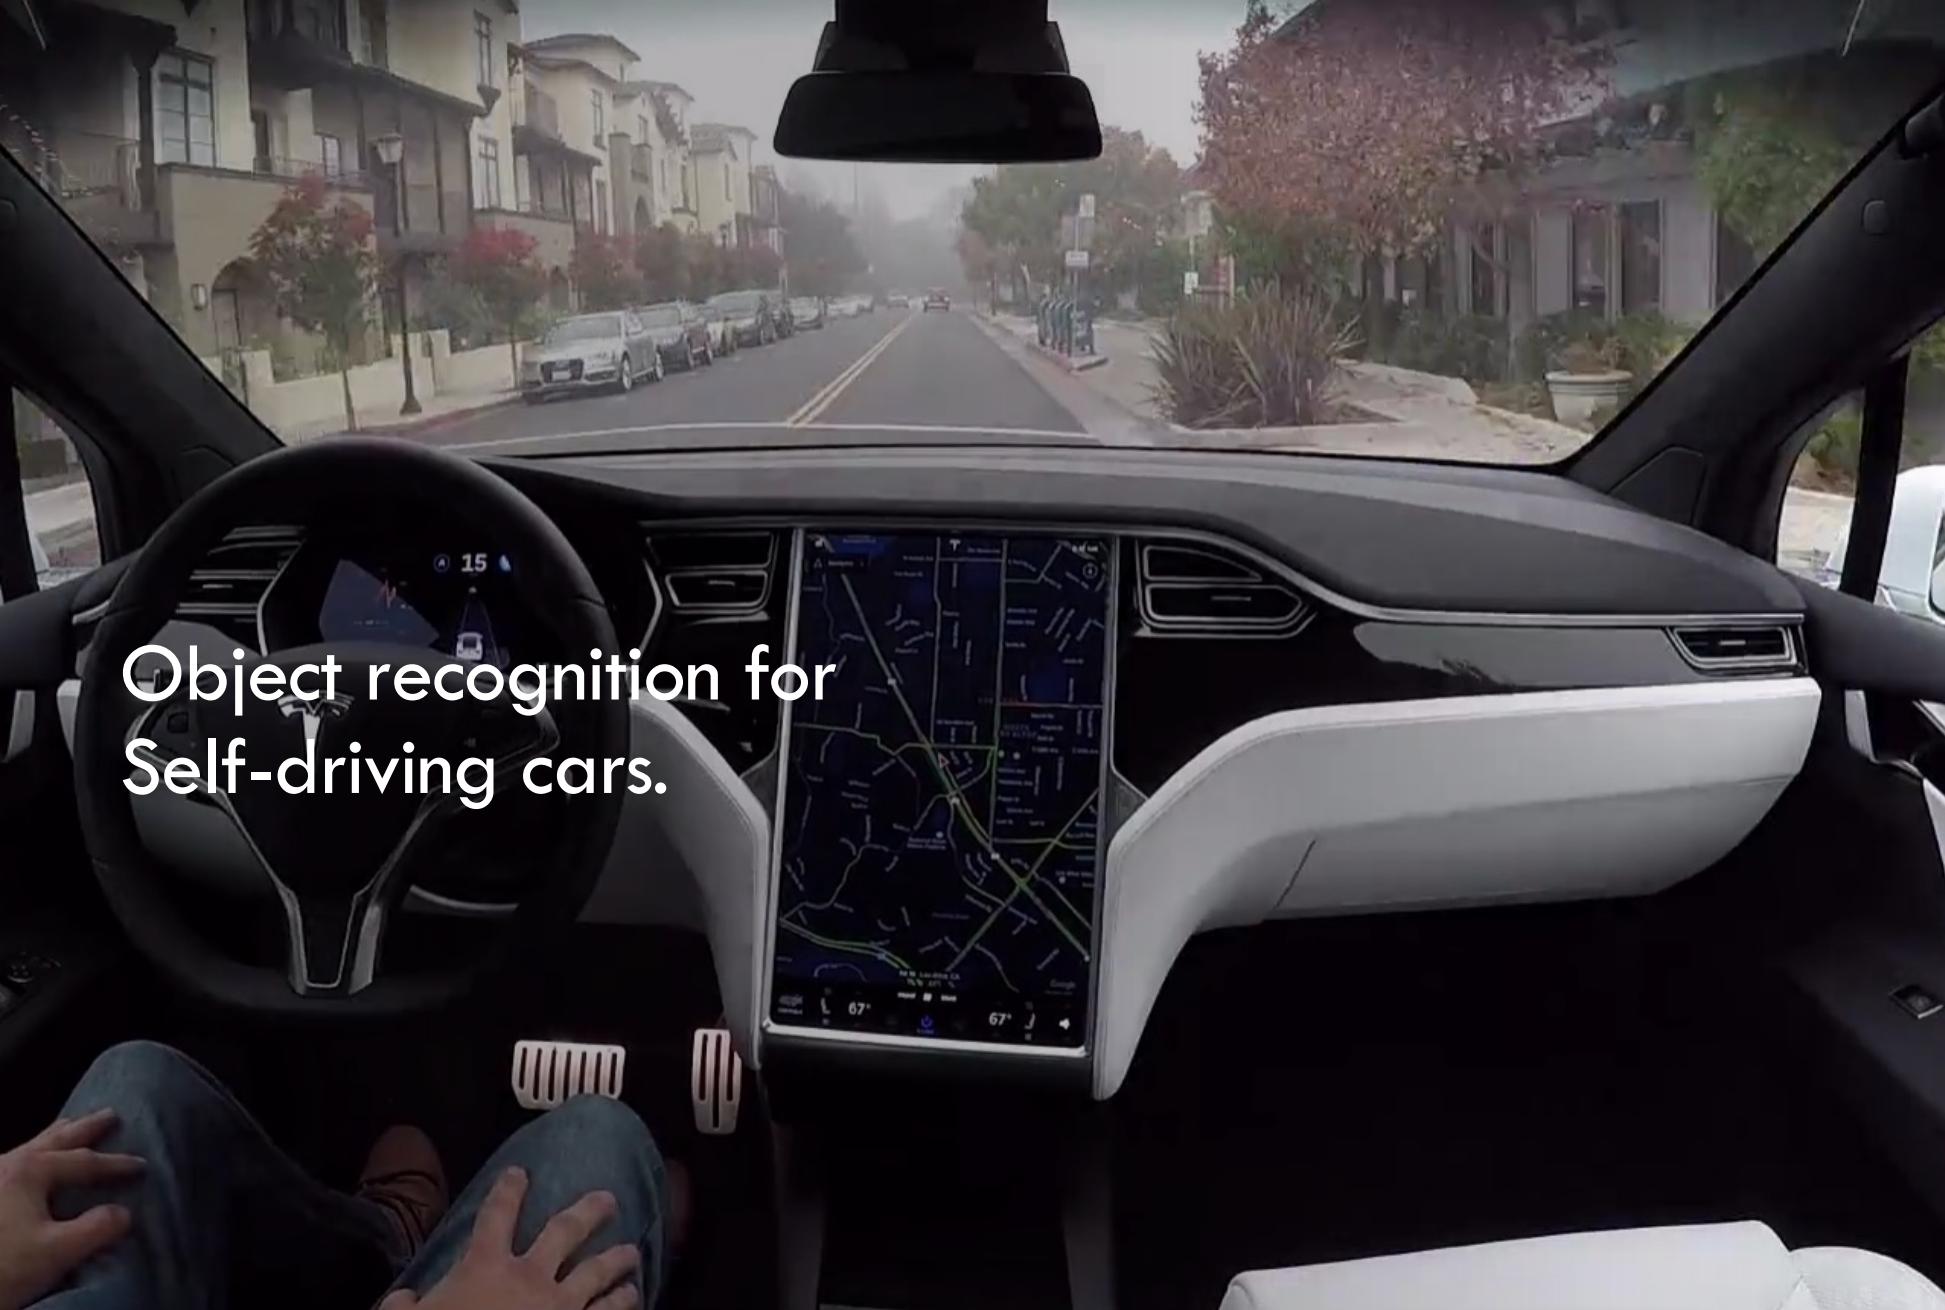
ROAD FLOW

IN-PATH OBJECTS

ROAD LIGHTS

OBJECTS

ROAD SIGNS



RIGHT REARWARD VEHICLE CAM

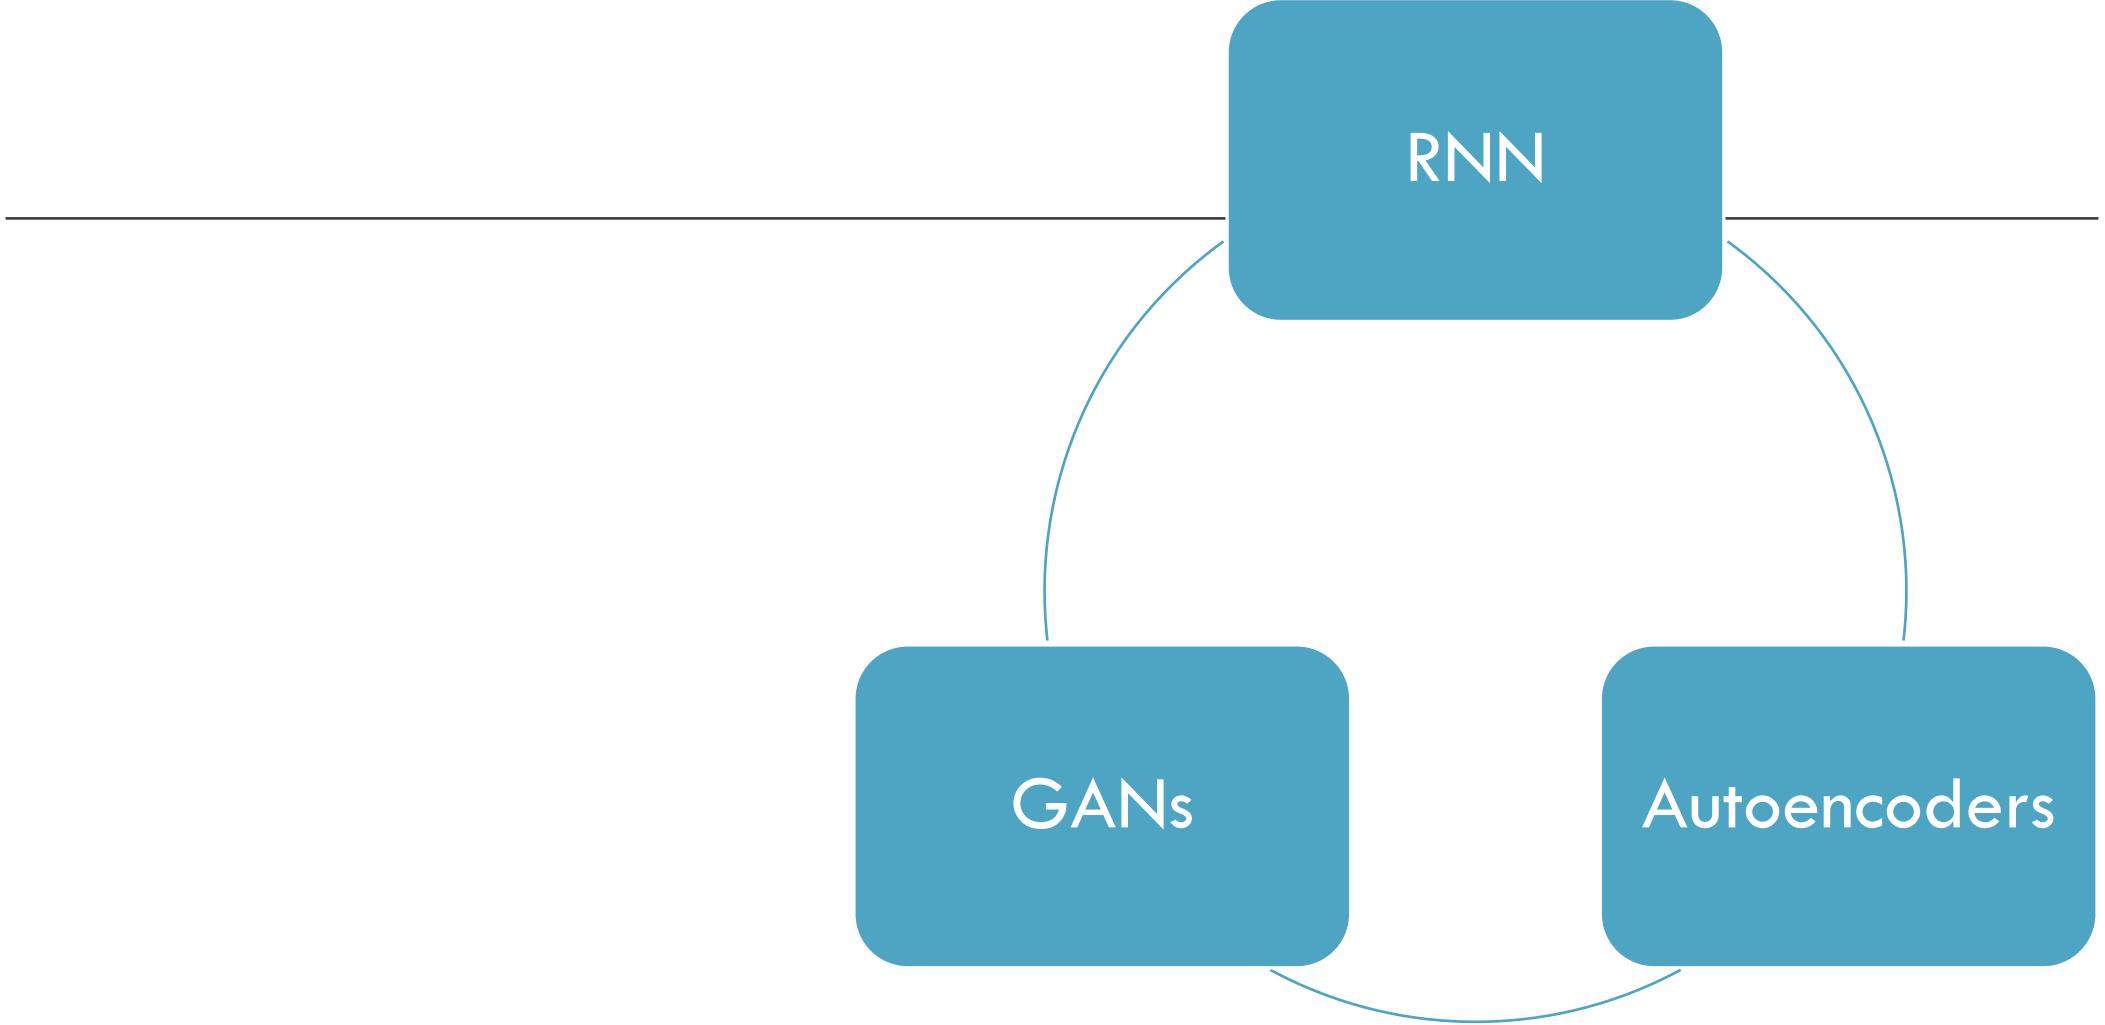
The background features a stylized brain shape composed of a glowing blue circuit board pattern. Numerous small, bright blue lights (nodes) are scattered across the circuit lines, creating a sense of neural activity or data flow. The overall effect is futuristic and represents the theme of advanced architecture.

Arquitecturas avanzadas

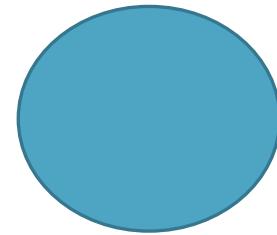
Arquitecturas avanzadas



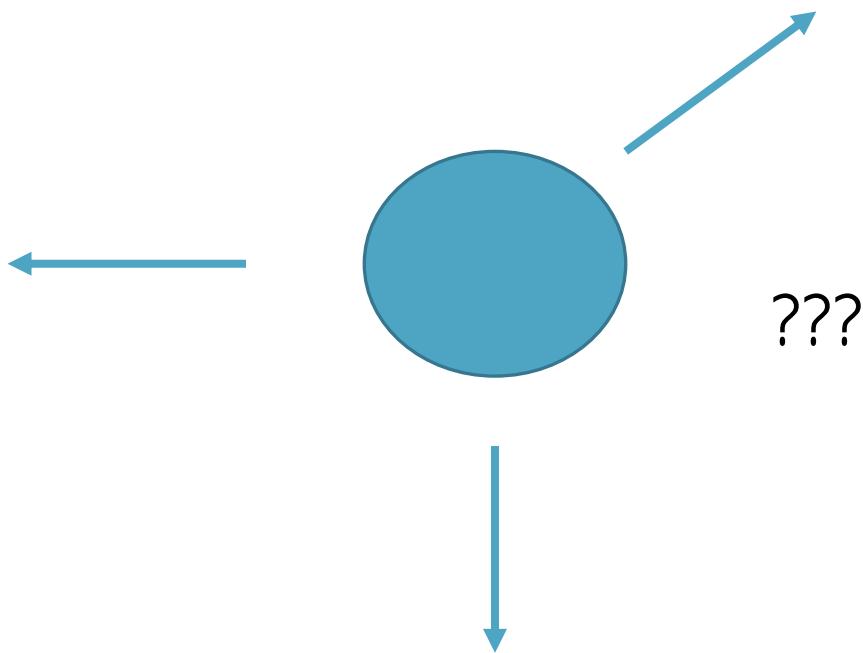
Ahora si se viene lo chido



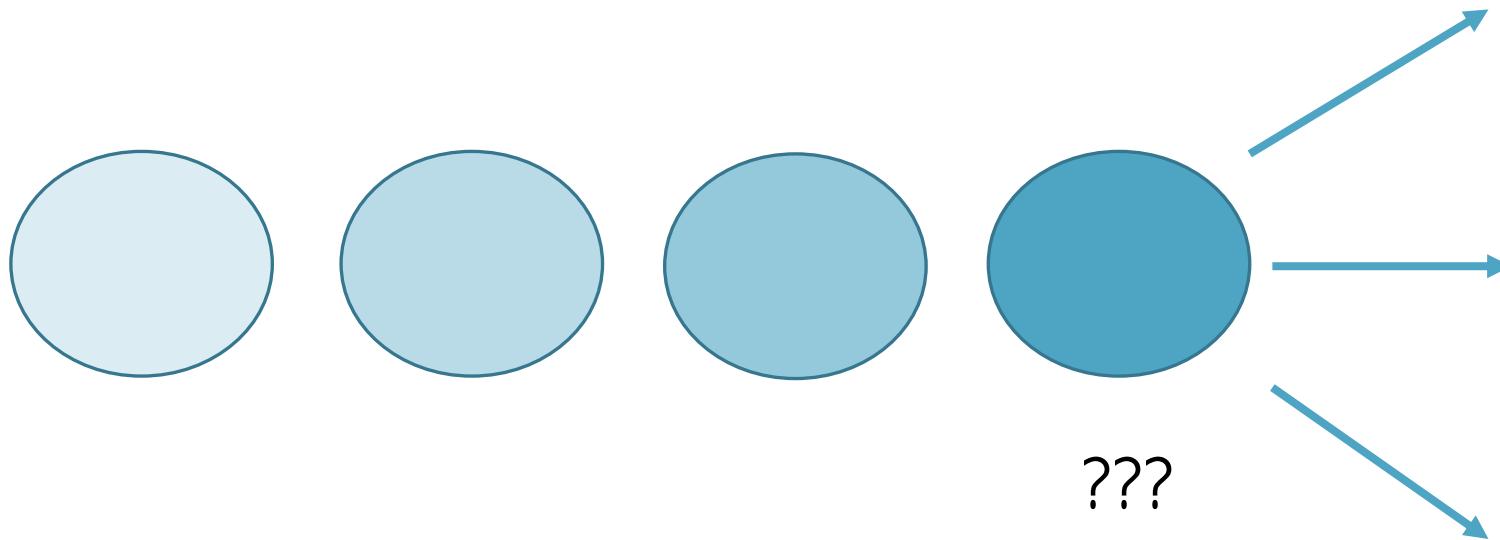
Supongamos que tenemos la imagen de una pelota,
podemos predecir a donde se va a mover?



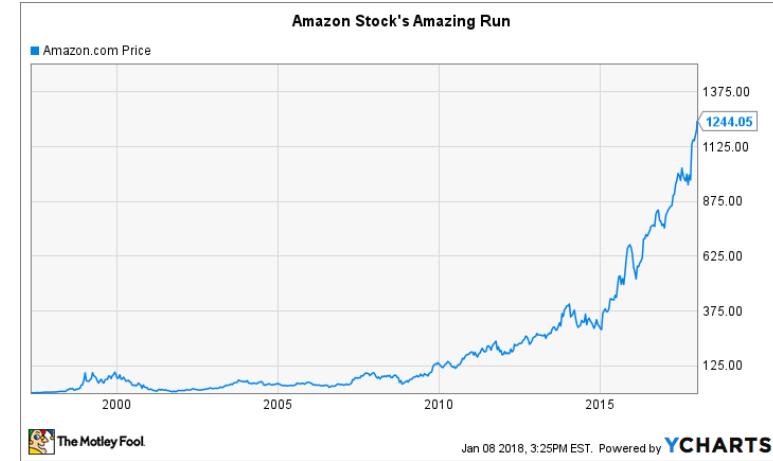
Supongamos que tenemos la imagen de una pelota,
podemos predecir a donde se va a mover?



Supongamos que tenemos la imagen de una pelota,
podemos predecir a donde se va a mover?



En todos lados
hay secuencias
de datos



Problema:

Dado un texto,
predecir la
siguiente palabra



Predecir la siguiente palabra

This morning I took my dog for a walk

Predecir la siguiente palabra

This morning I took my dog for a walk

Datos

Predicción

Predecir la siguiente palabra

This morning I took my dog for a walk

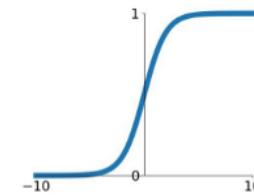
Datos

Predicción

Opción 1: Regresión Logistica

Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



Predecir la siguiente palabra

This morning I took my dog for a walk

Datos

Predicción

Opción 1: Regresión Logistica

Opción 2 : Modelos probabilisticos como Naïve Bayes o LDA

$$\begin{aligned} P(\text{text}) &= P(x_0, \dots, x_n) = \\ &= P(x_0)P(x_1|x_0)P(x_2|x_0, x_1)\dots P(x_n|\dots) \end{aligned}$$

Predecir la siguiente palabra

This morning I took my dog for a walk

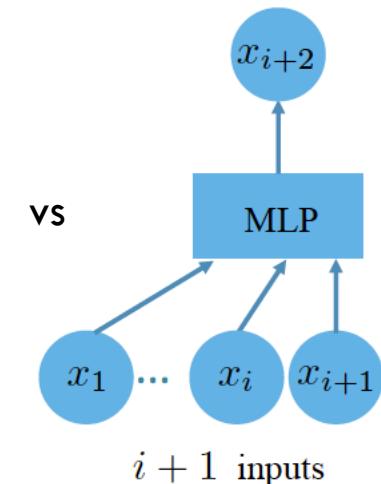
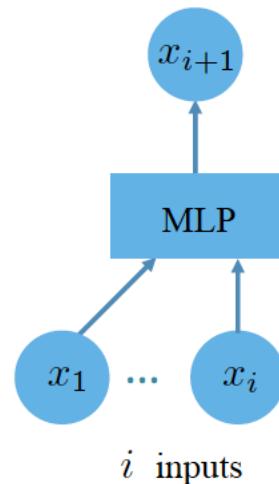
Opción 1: Regresión Logistica

Opción 2 : Modelos probabilisticos como Naïve Bayes o LDA

Opción 3: Feedforward NN

Datos

Predicción





Entonces, qué necesitamos?

Un modelo (o layer) que pueda:

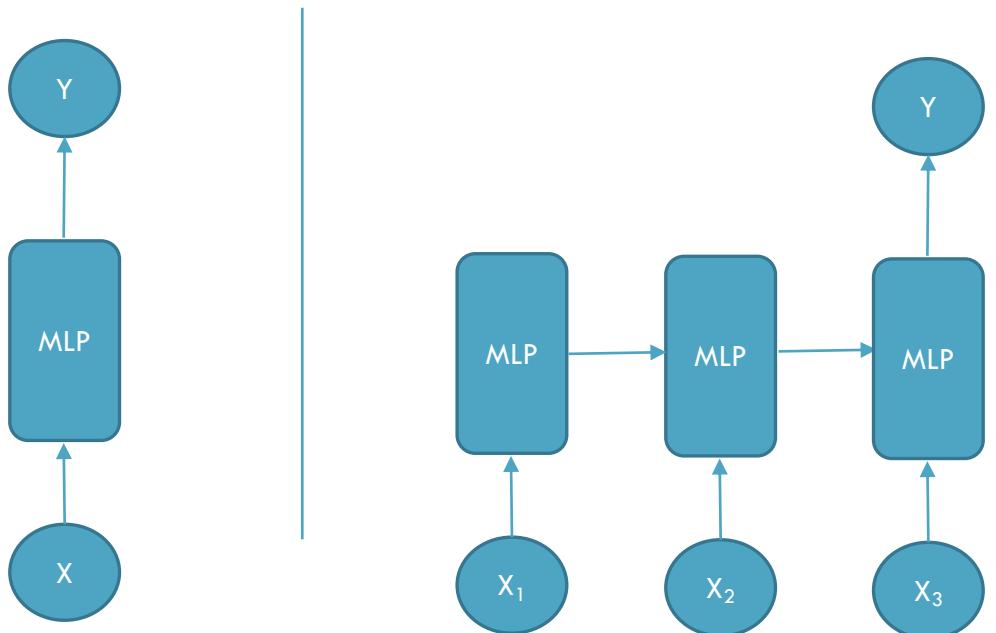
- Aceptar longitud variable como input
- Capturar dependencias en datos independientemente de su posición
- Compartir parámetros a travez de la secuencia.

Recurrent Neural Network

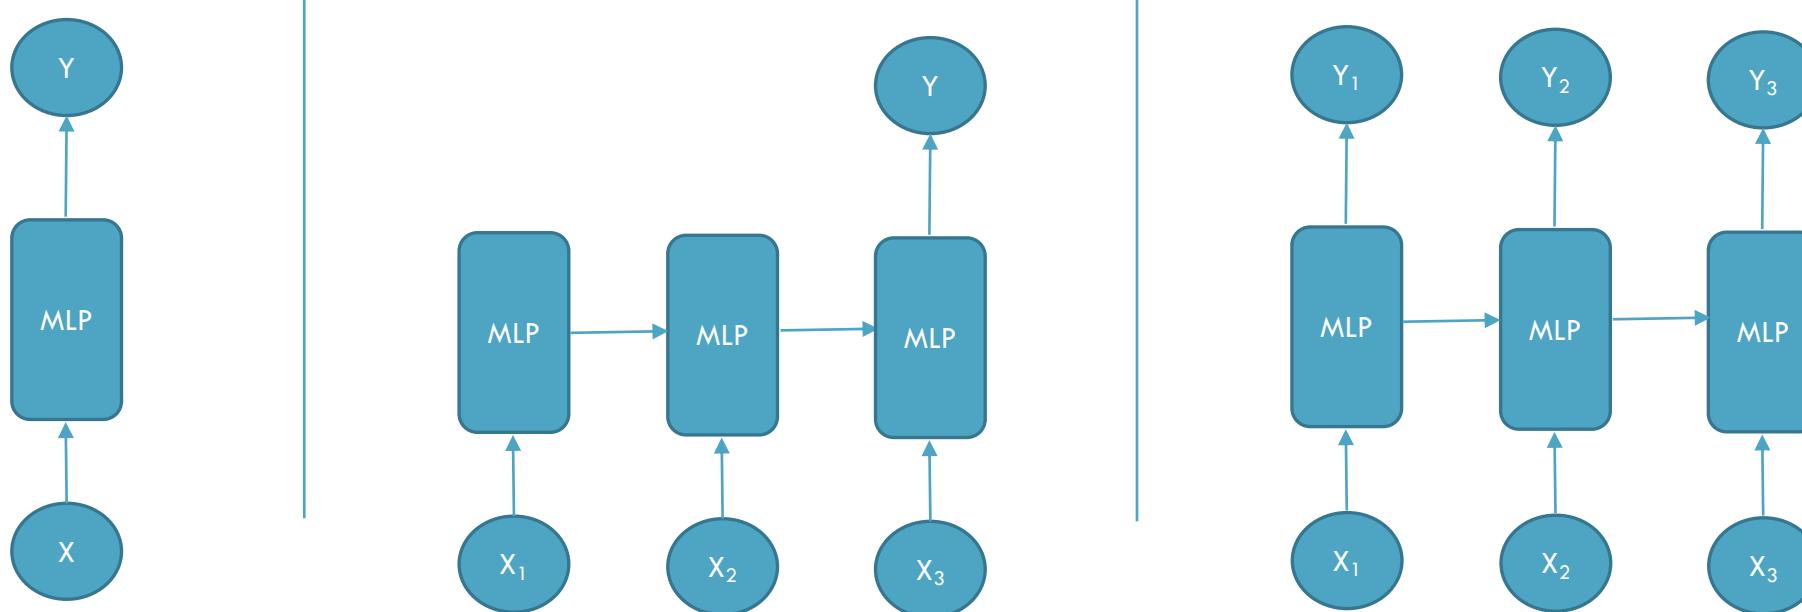
Recurrent Layer



Recurrent Layer

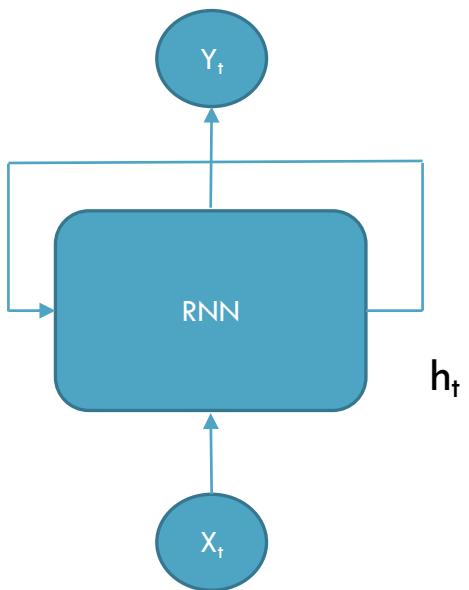


Recurrent Layer

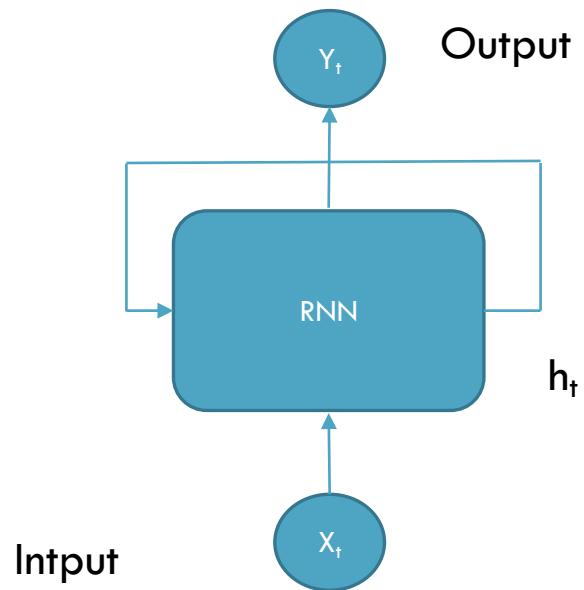


Y muchas otras
arquitecturas más...

Recurrent Layer



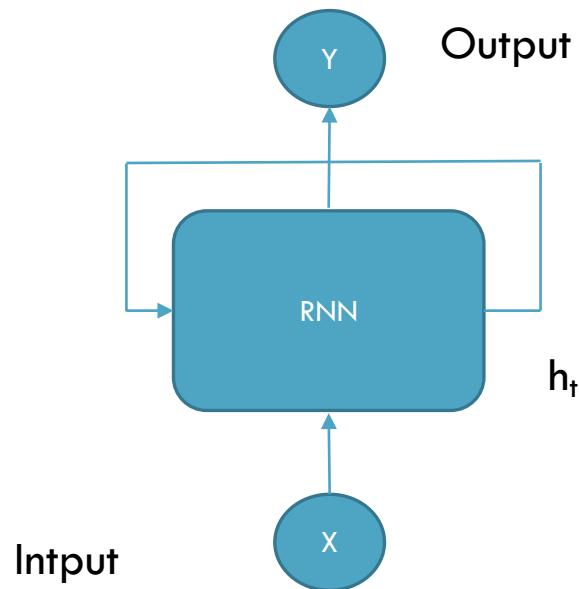
Recurrent Layer



RNN aplican una relación secuencial en cada paso para procesar la secuencia

$$h_t = f_W(h_{t-1}, x_t)$$

Recurrent Layer



RNN aplican una relación secuencial en cada paso para procesar la secuencia

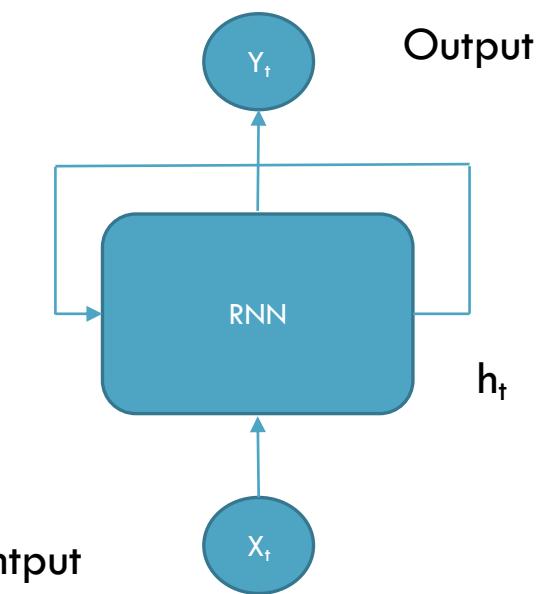
$$h_t = f_W(h_{t-1}, x_t)$$

Diagram illustrating the computation of the hidden state h_t in an RNN:

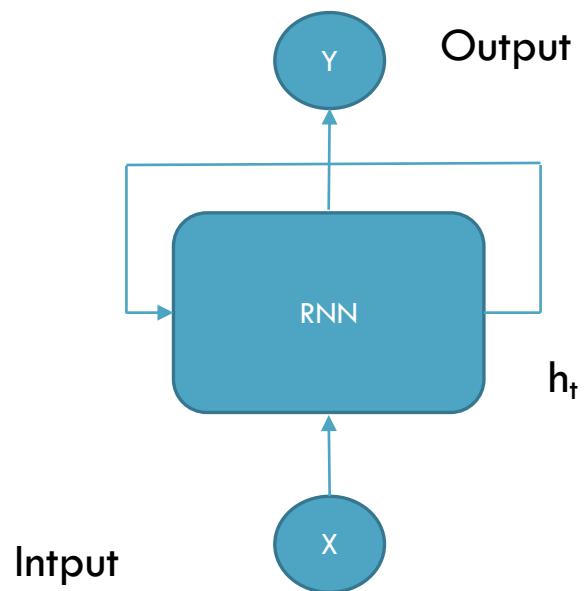
- Cell state** (green line): The previous cell state h_{t-1} .
- Matriz de pesos W** (green line): The weight matrix f_W .
- Cell state anterior** (red line): The previous cell state h_{t-1} .
- Nueva observación** (blue line): The new observation x_t .

Recurrent NN intuición

```
my_rnn()  
  
hidden_state = [0 , 0 , 0 , 0]  
  
sentence = ["I" , "love" , "recurrent" , "neural" ]  
  
For word in sentence:  
    prediction, hidden_state = my_rnn ( word hidden_state )  
  
Next_word = prediction
```



Entrenamiento

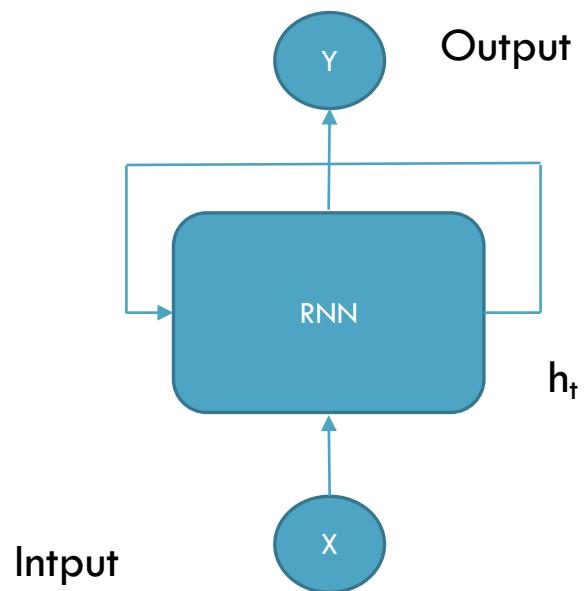


Como se entrena una RNN?

$$h_t = \tanh(W_{hh}^T h_{t-1}, W_{xh}^T x_t)$$

$$Y_t = W_{hy}^T h_t$$

Entrenamiento

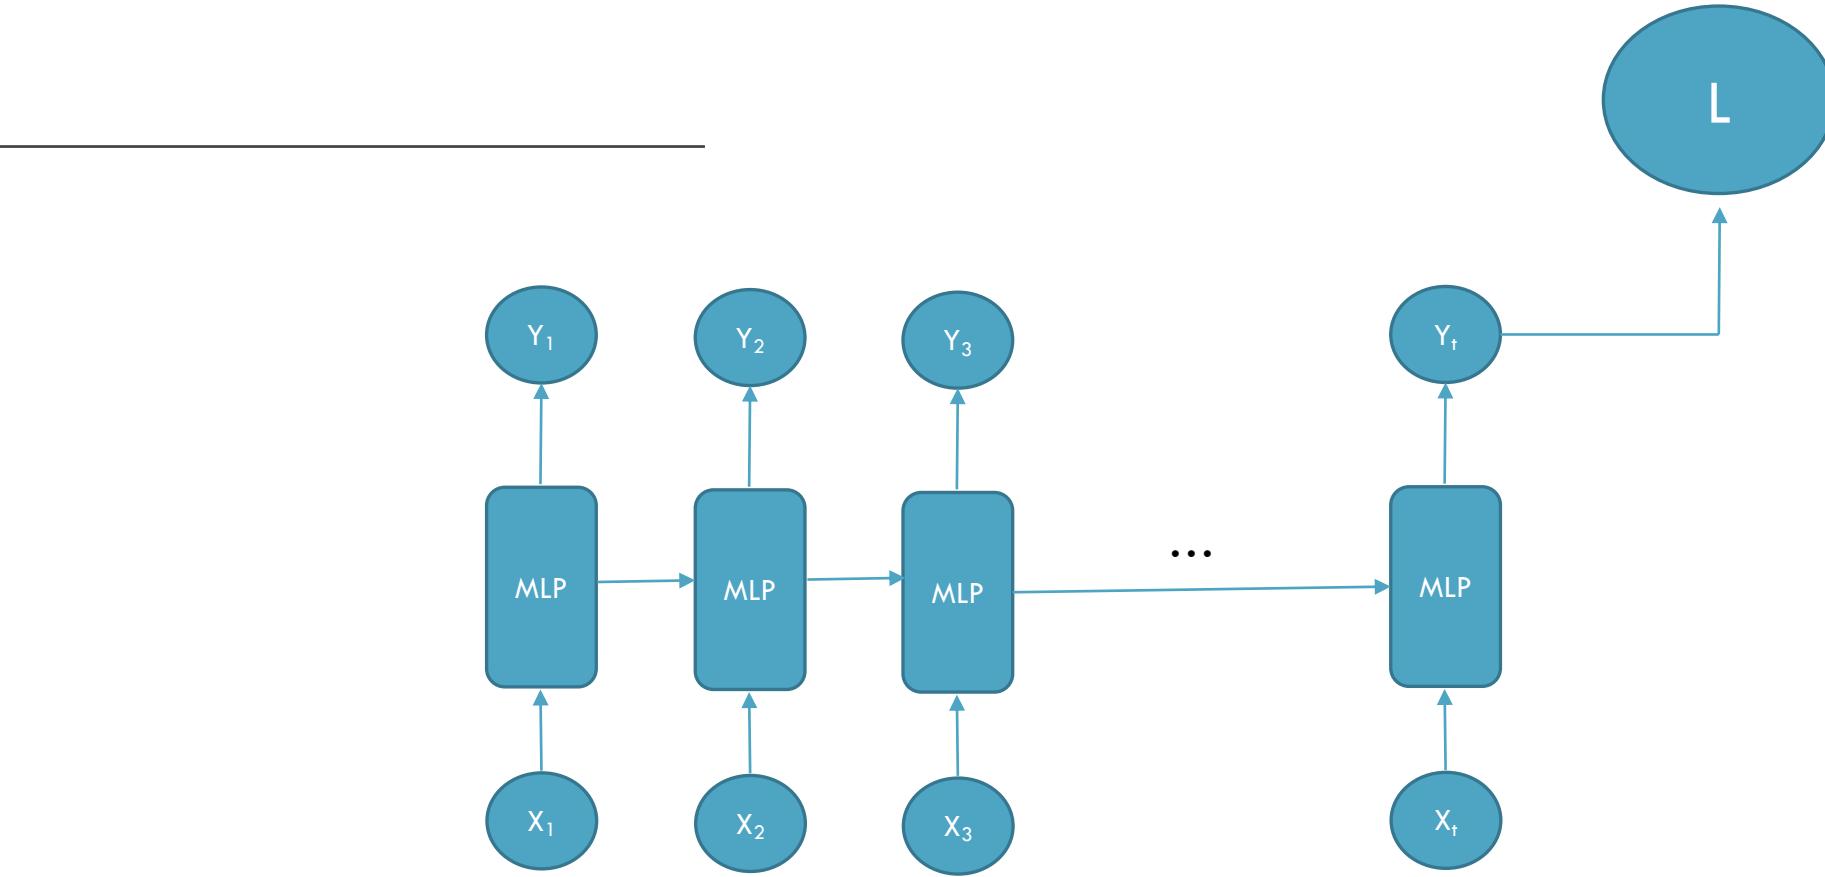


Como se entrena una RNN?

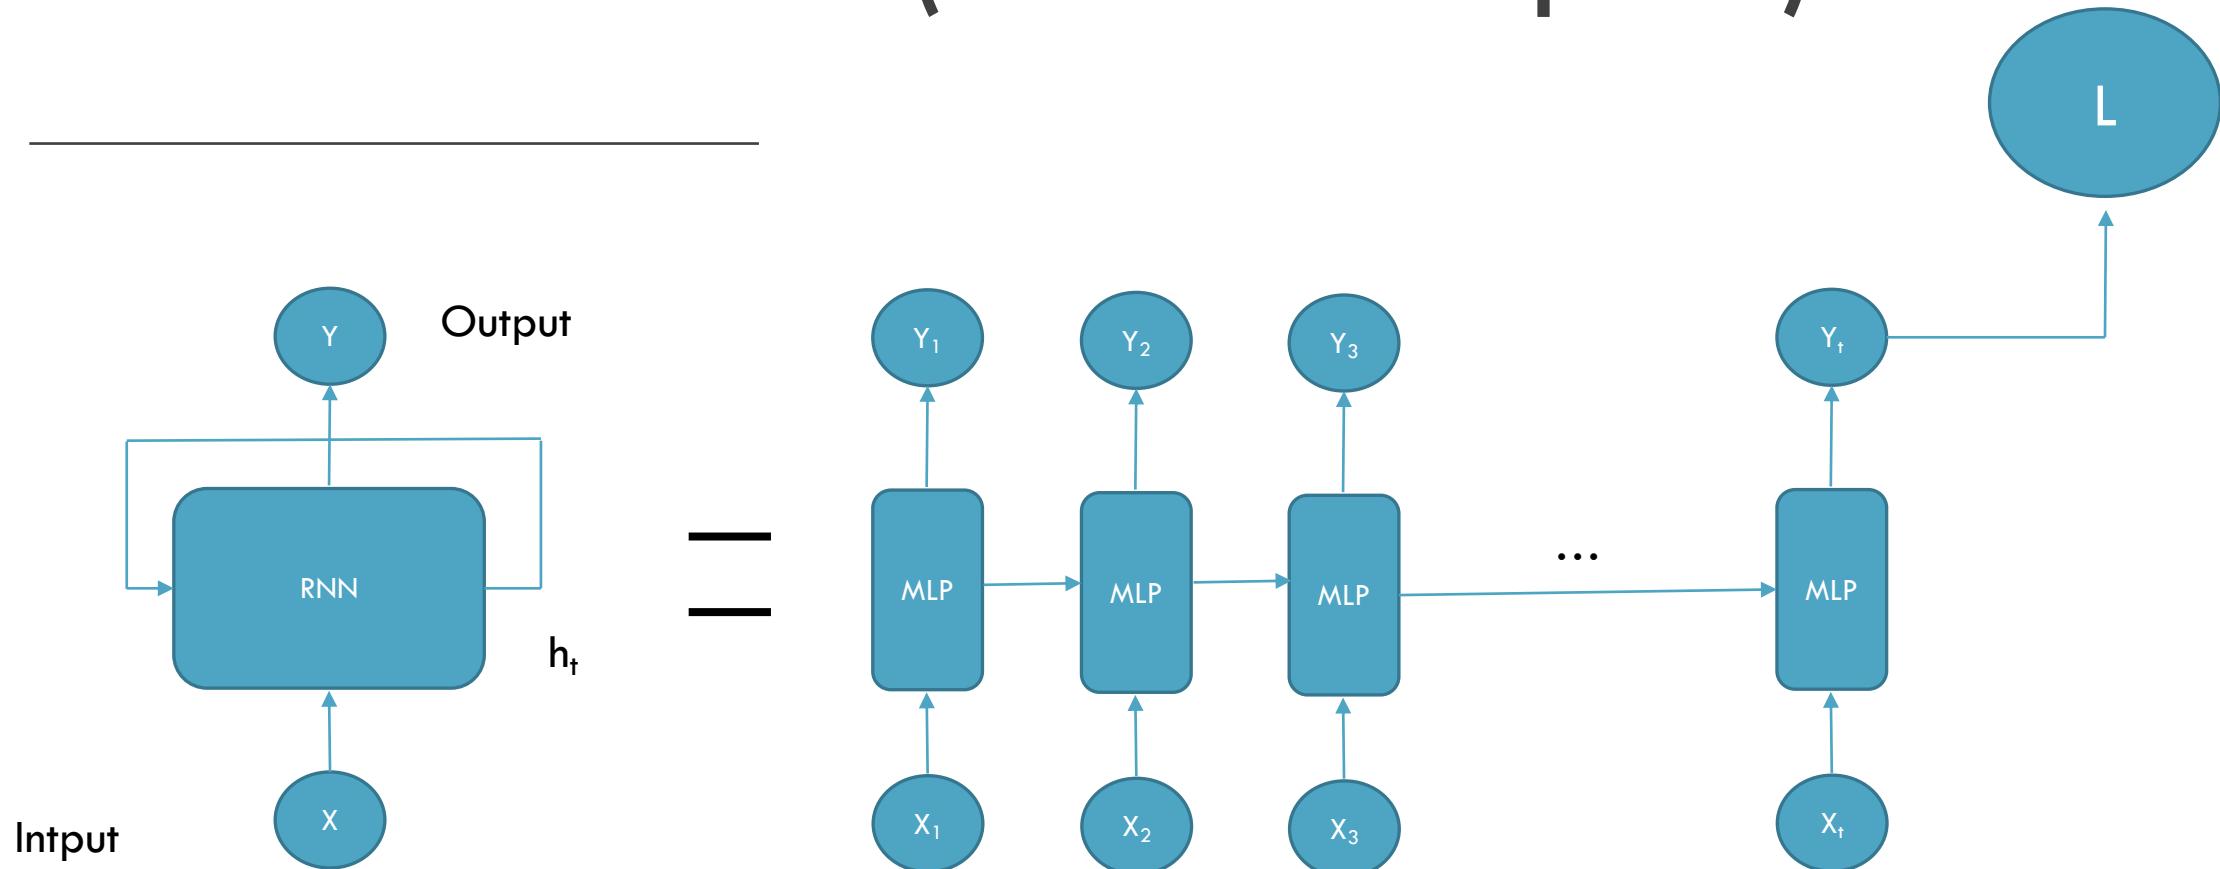
$$h_t = \tanh(W_{hh}^T h_{t-1}, W_{xh}^T x_t)$$

Backpropagation
Through Time

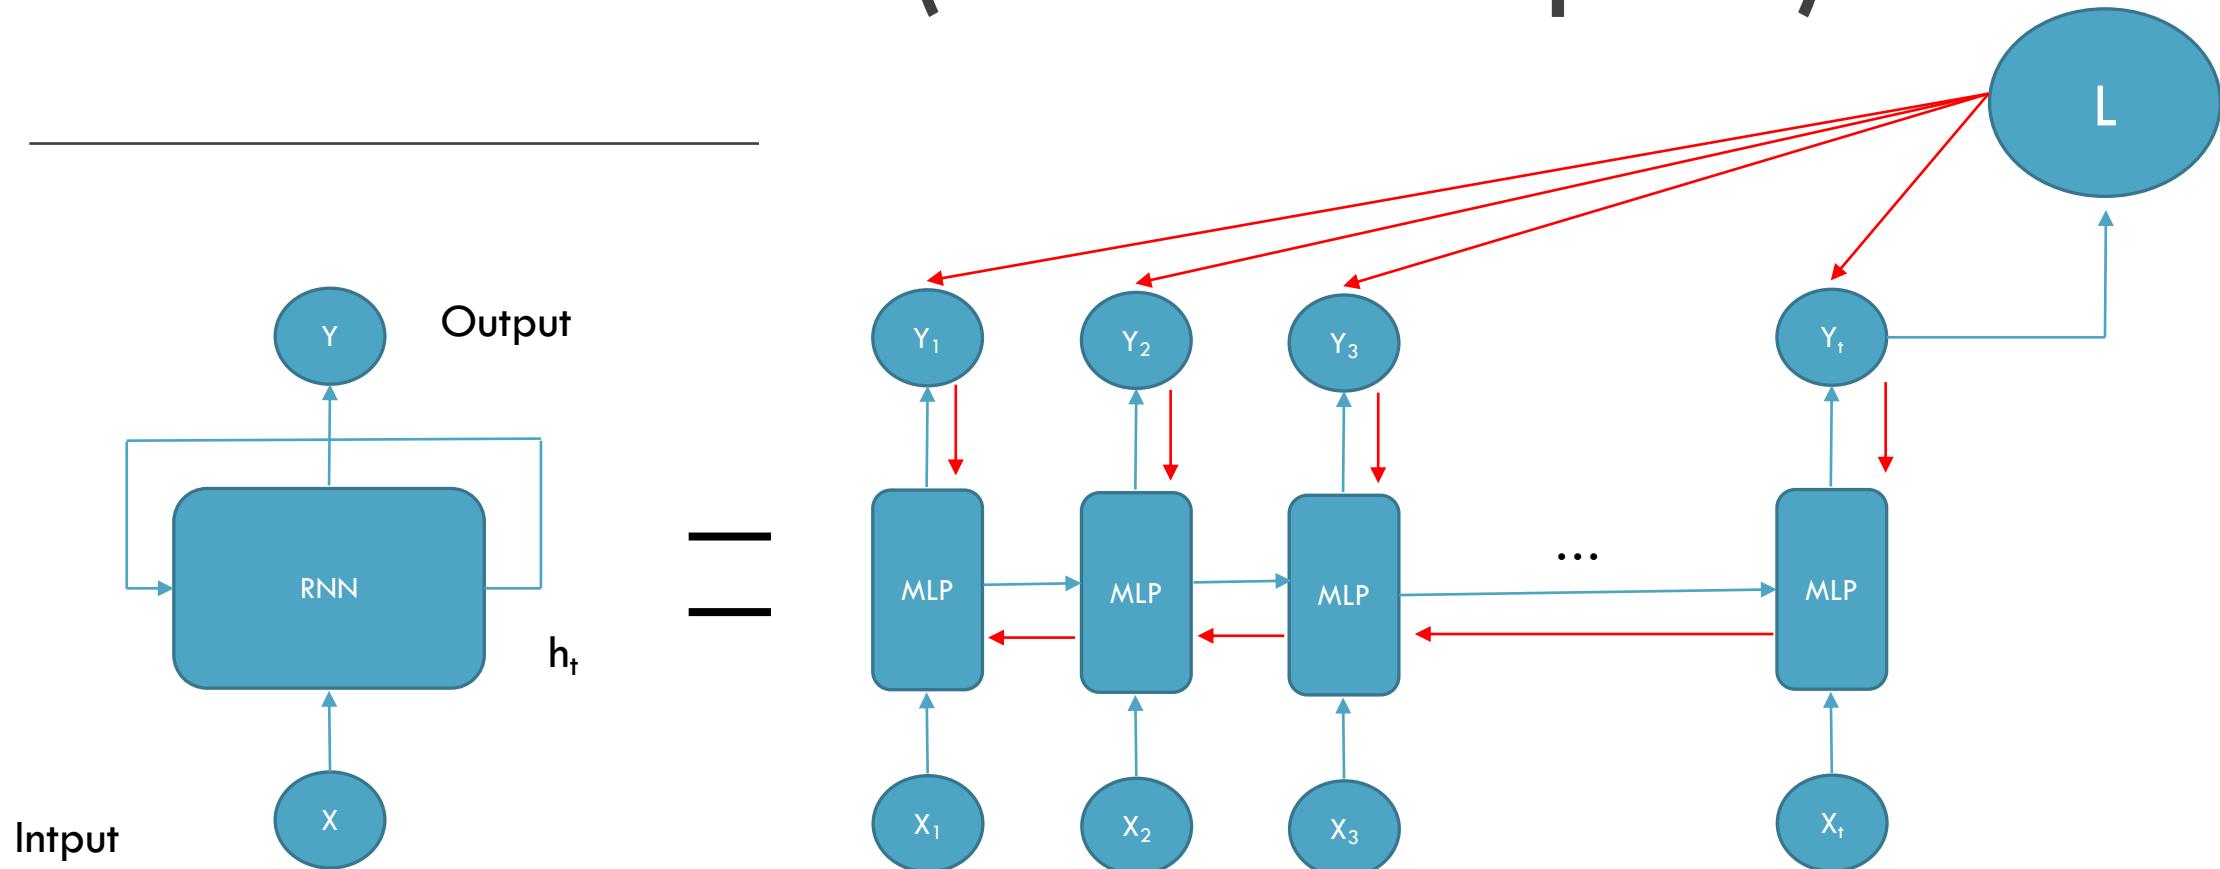
$$Y_t = W_{hy}^T h_t$$



BBTT intuición (Forward pass)



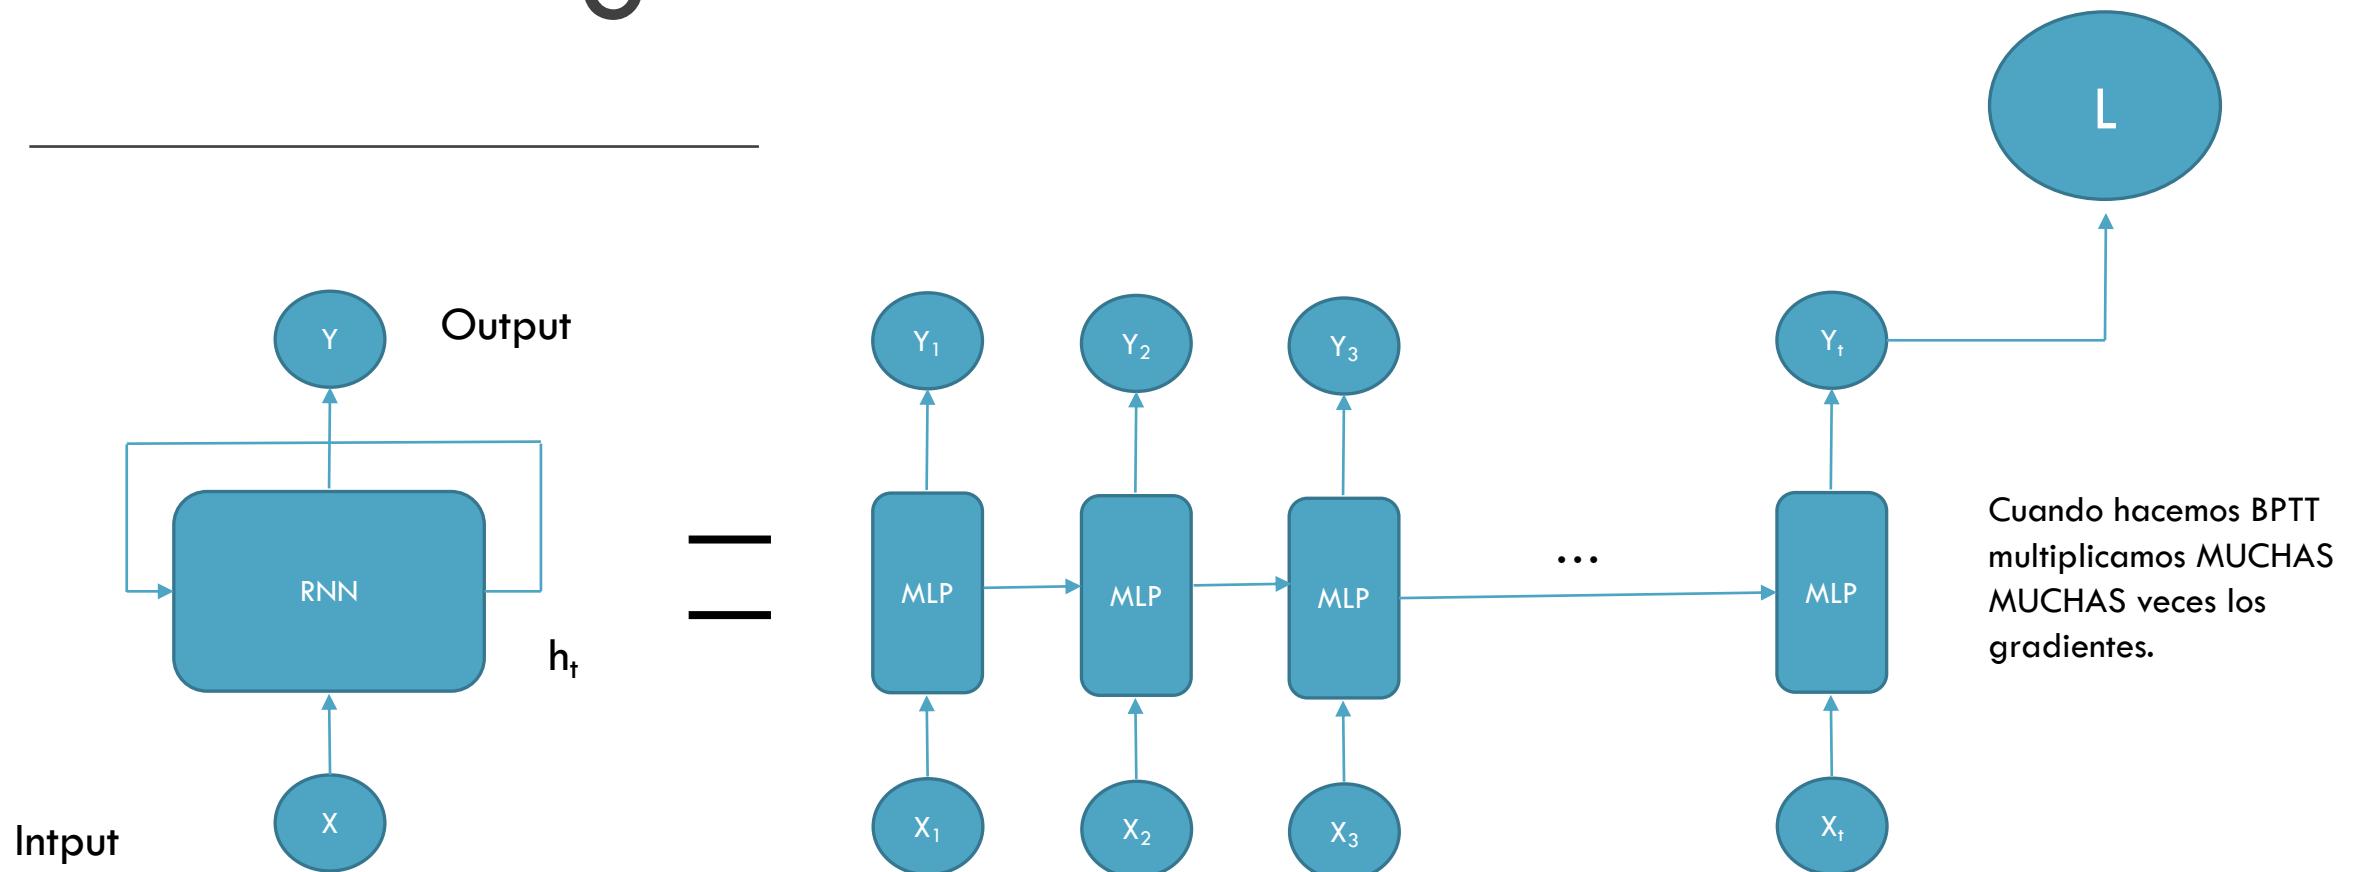
BBTT intuición (Backward pass)





Questions

Vanishing Gradient



Qué pasa con los gradientes?

$$\lim_{n \rightarrow \infty} \prod a$$

Si $\alpha < 0$

Si $\alpha > 0$

Qué pasa con los gradientes?

$$\lim_{n \rightarrow \infty} \prod a$$

Si $a < 0$

Si $a > 0$

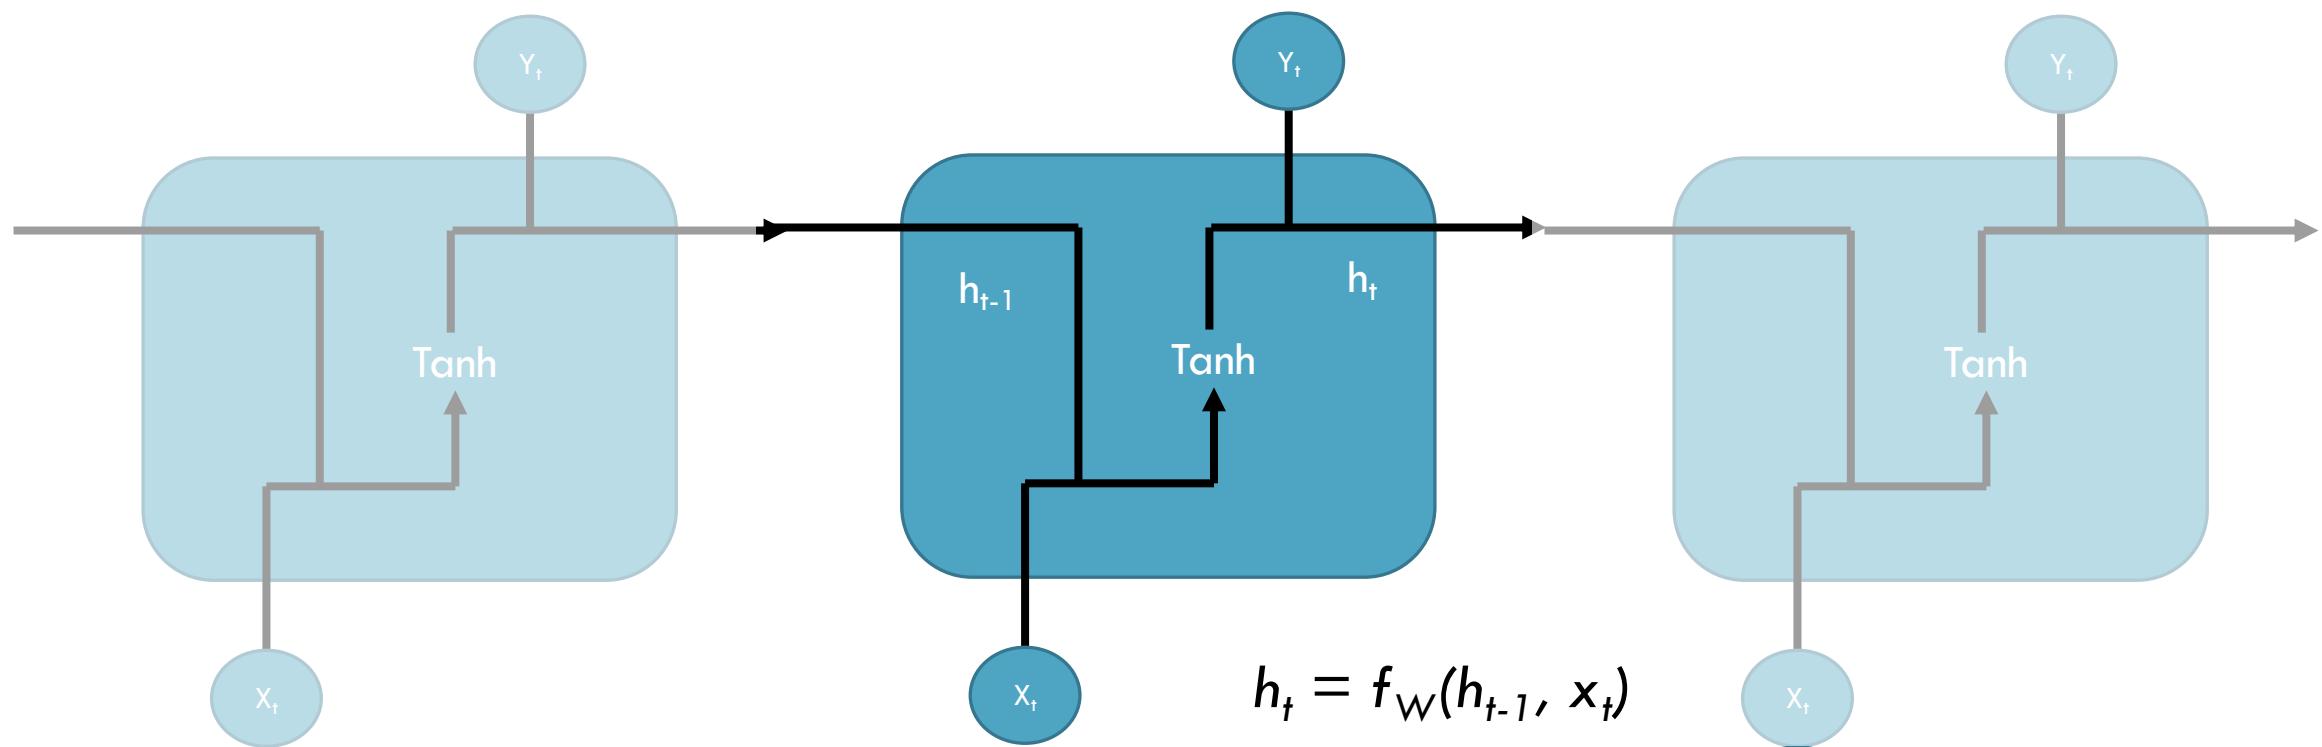
Entonces, necesitamos que los gradientes sean igual (o lo mas similares posibles) a 1

Dummy solution: utilizar ReLU como función de activación, por qué?

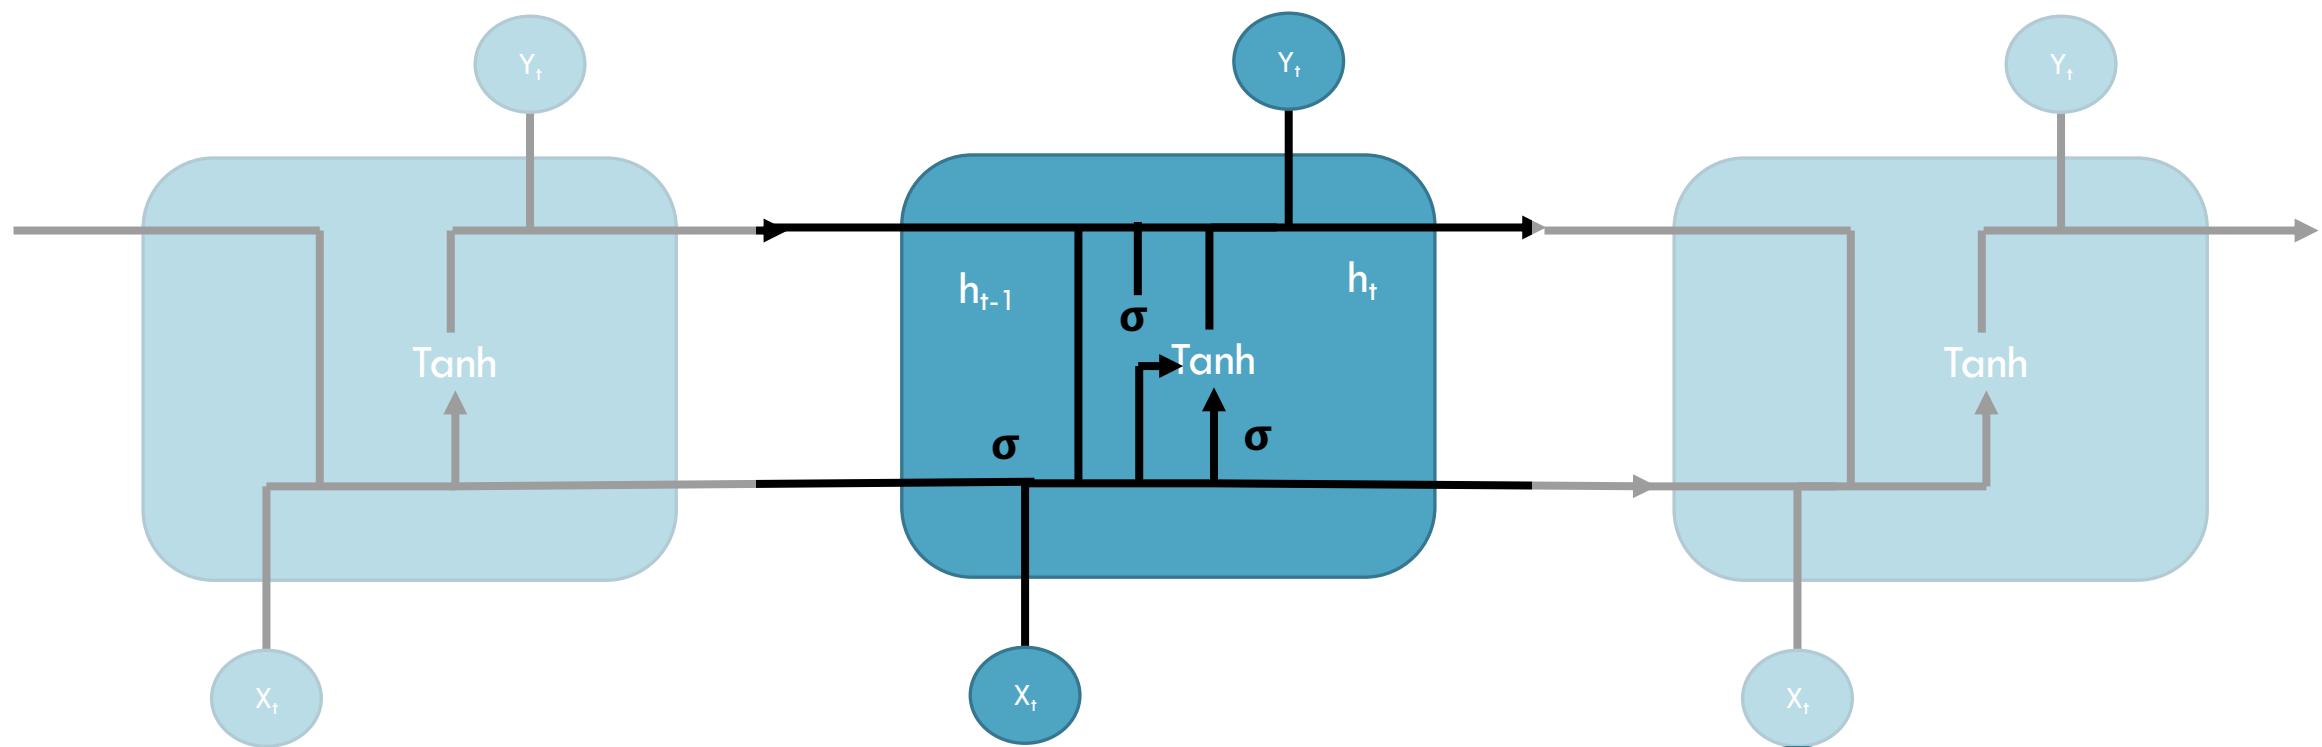
LSTM (Long-Short Term Memory)



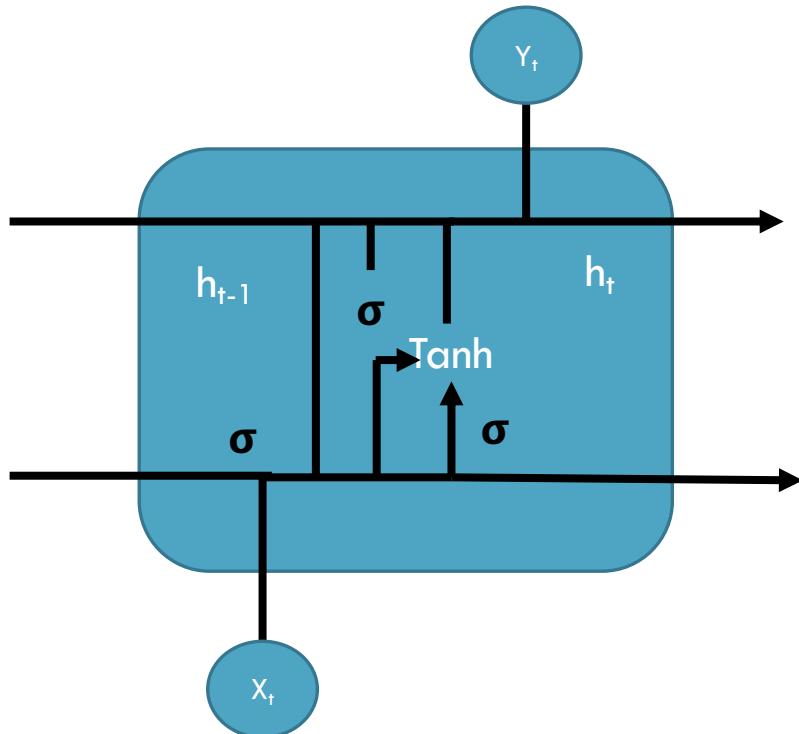
LSTM (Long-Short Term Memory)



LSTM (Long-Short Term Memory)



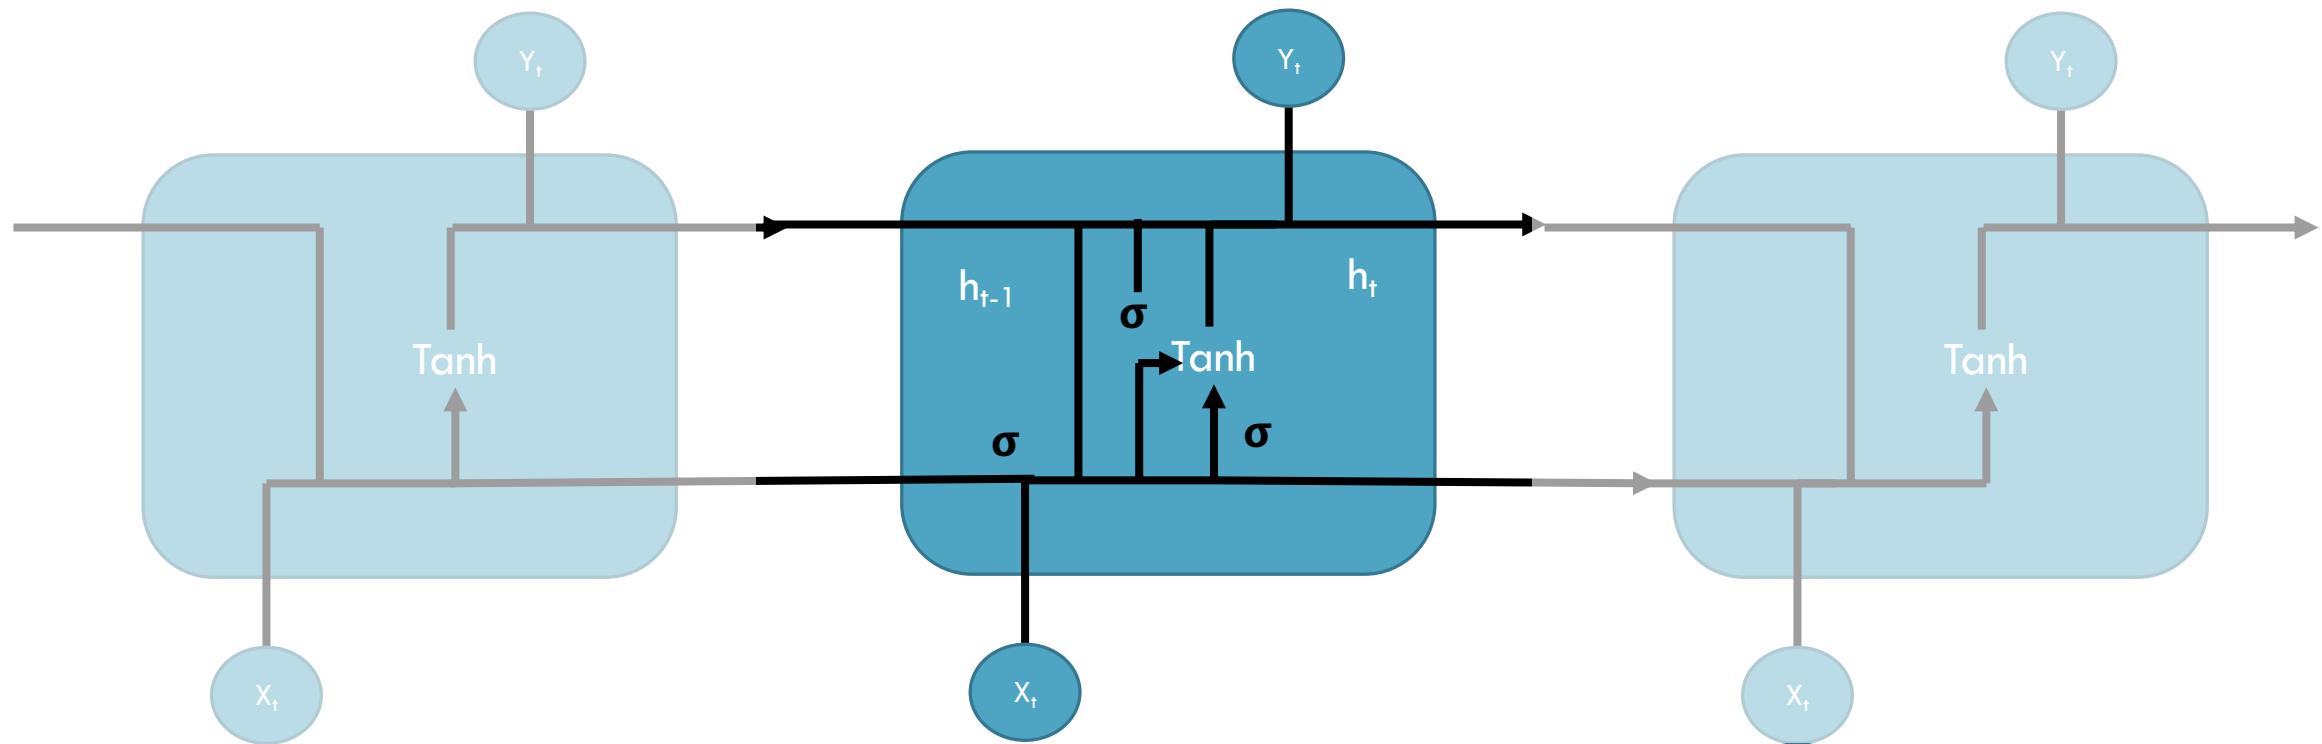
LSTM (Long-Short Term Memory)



LSTM cumplirá 4 diferentes propósitos con σ :

1. Olvidar (forget)
2. Guardar (store)
3. Actualizar (Update)
4. Predecir (Output)

LSTM (Long-Short Term Memory)



LSTM (RNN) Resumen

1. Mantiene las celdas separadas en cada iteración
2. Utiliza 4 puertas para controlar la información relevante

Olvida la información irrelevante.

Guarda lo la información importante

Actualiza los pesos de la matriz W

Predice el output final

RNN Aplicaciones

Choping music generated with LSTM:

<https://www.youtube.com/watch?v=j60J1cGINX4>

RNN Aplicaciones

Choping music generated with LSTM:

<https://www.youtube.com/watch?v=j60J1cGINX4>

Sentiment Classification

Python Notebook

RNN Aplicaciones

Choping music generated with LSTM:

<https://www.youtube.com/watch?v=j60J1cGINX4>

Sentiment Classification

Python Notebook

Traducción de textos

	eng	hin
0	give your application an accessibility workout	अपने अनुप्रयोग को पहुंचनीयता व्यायाम का लाभ दें
1	accerciser accessibility explorer	एक्सेसिलिसर पहुंचनीयता अन्वेषक
2	the default plugin layout for the bottom panel	निचले पटल के लिए डिफॉल्ट प्लगइन खाका
3	the default plugin layout for the top panel	ऊपरी पटल के लिए डिफॉल्ट प्लगइन खाका
4	a list of plugins that are disabled by default	उन प्लगइनों की सूची जिनमें डिफॉल्ट रूप से निष...
5	highlight duration	अवधि को हाइलाइट रक्कें
6	the duration of the highlight box when select...	पहुंचनीय आसंधि नोड को चुनते समय हाइलाइट बक्से ...
7	highlight border color	सीमांत बोर्डर के रंग को हाइलाइट करें
8	the color and opacity of the highlight border	हाइलाइट किए गए सीमांत का रंग और अपारदर्शिता।
9	highlight fill color	भराई के रंग को हाइलाइट करें
10	the color and opacity of the highlight fill	हाइलाइट किया गया भराई का रंग और पारदर्शिता।



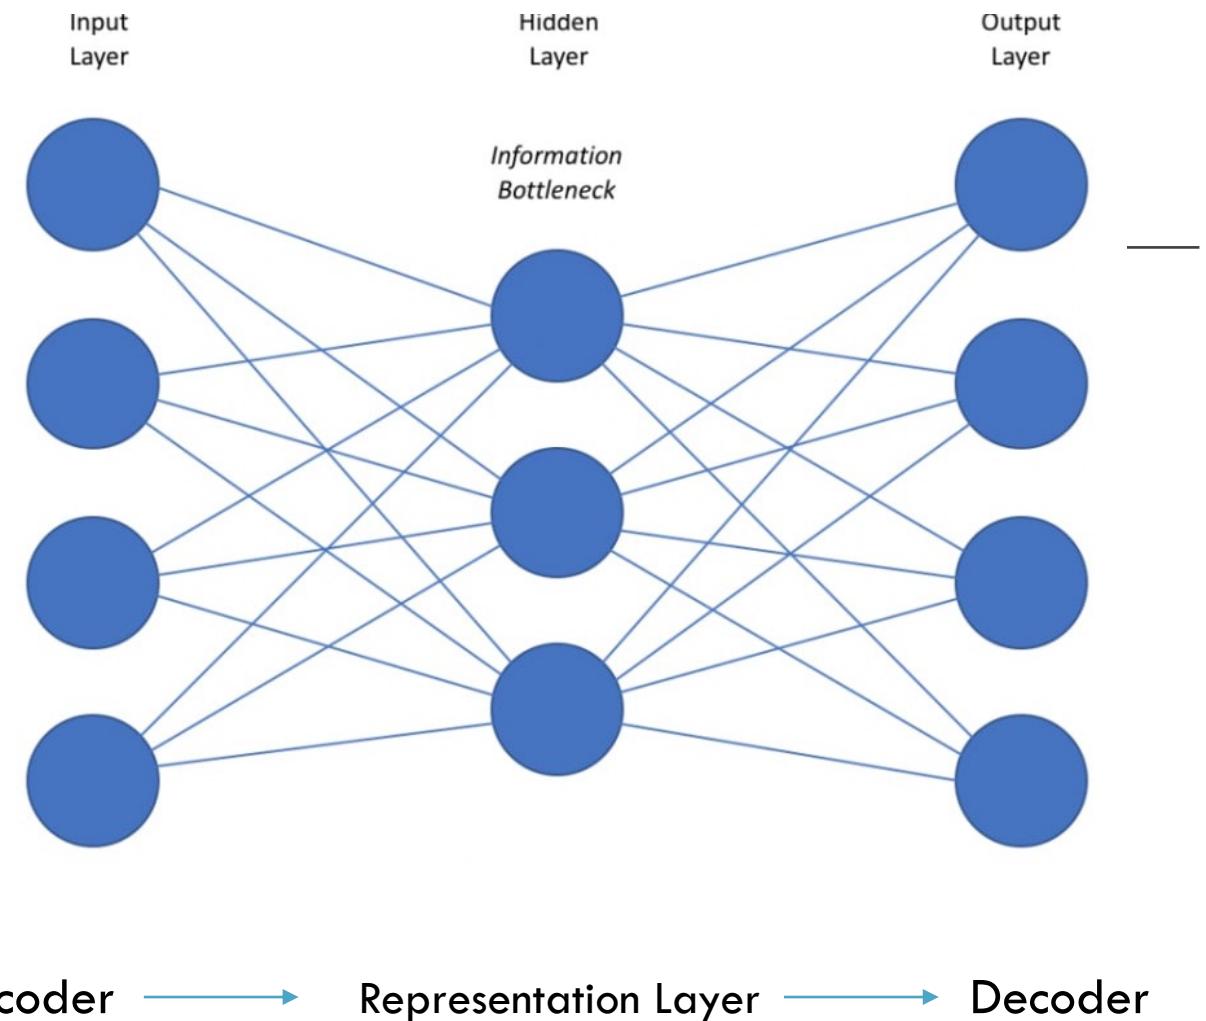
Questions

Autoencoder (Unsupervised DL)

Idea:

Tomar el input, pasarlo por una estructura de red neuronal y despues reconstruirlo

$$\text{Decoder}(\text{Encoder}(x)) = X$$

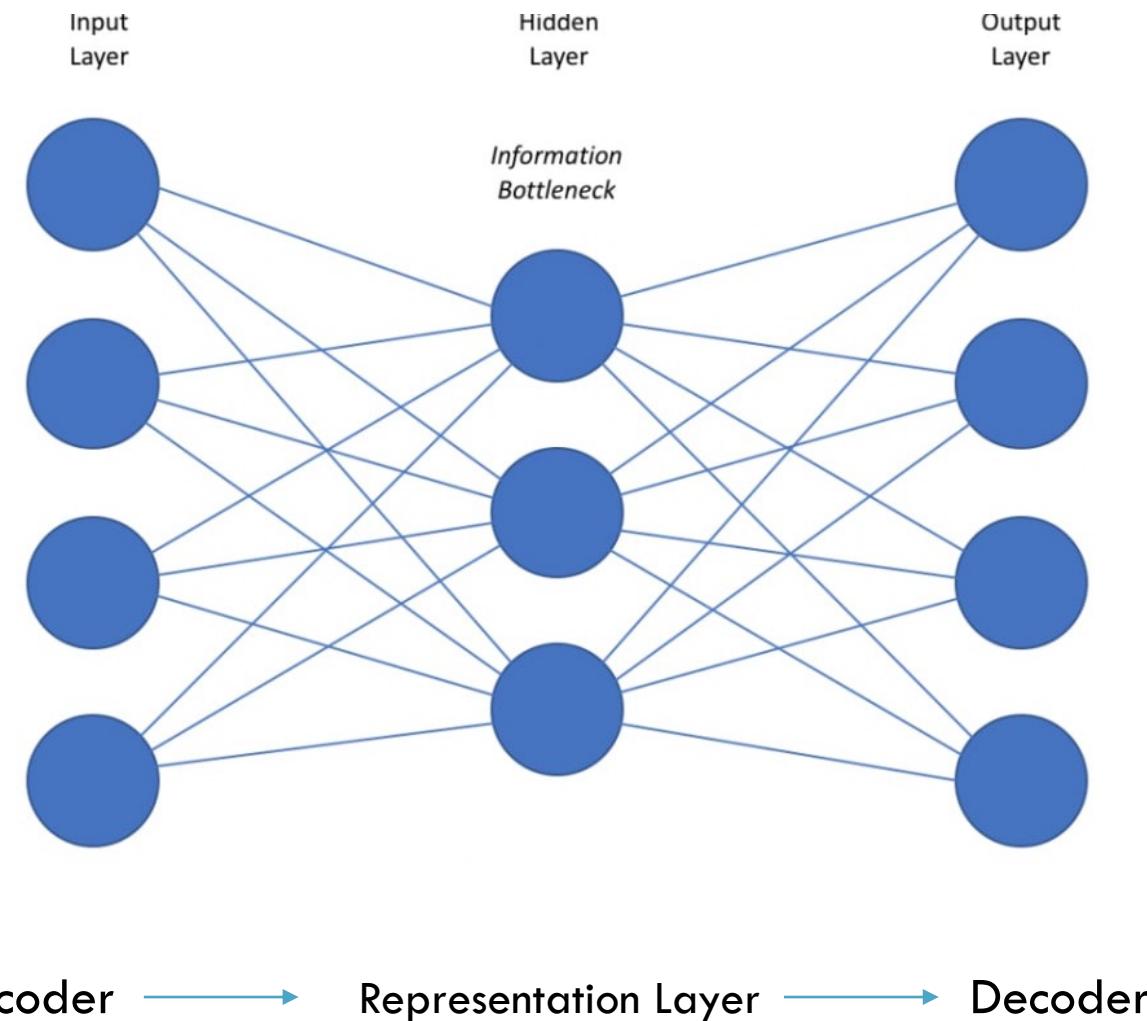


Autoencoder (Unsupervised DL)

Idea:

Tomar el input, pasarlo por una estructura de red neuronal y despues reconstruirlo

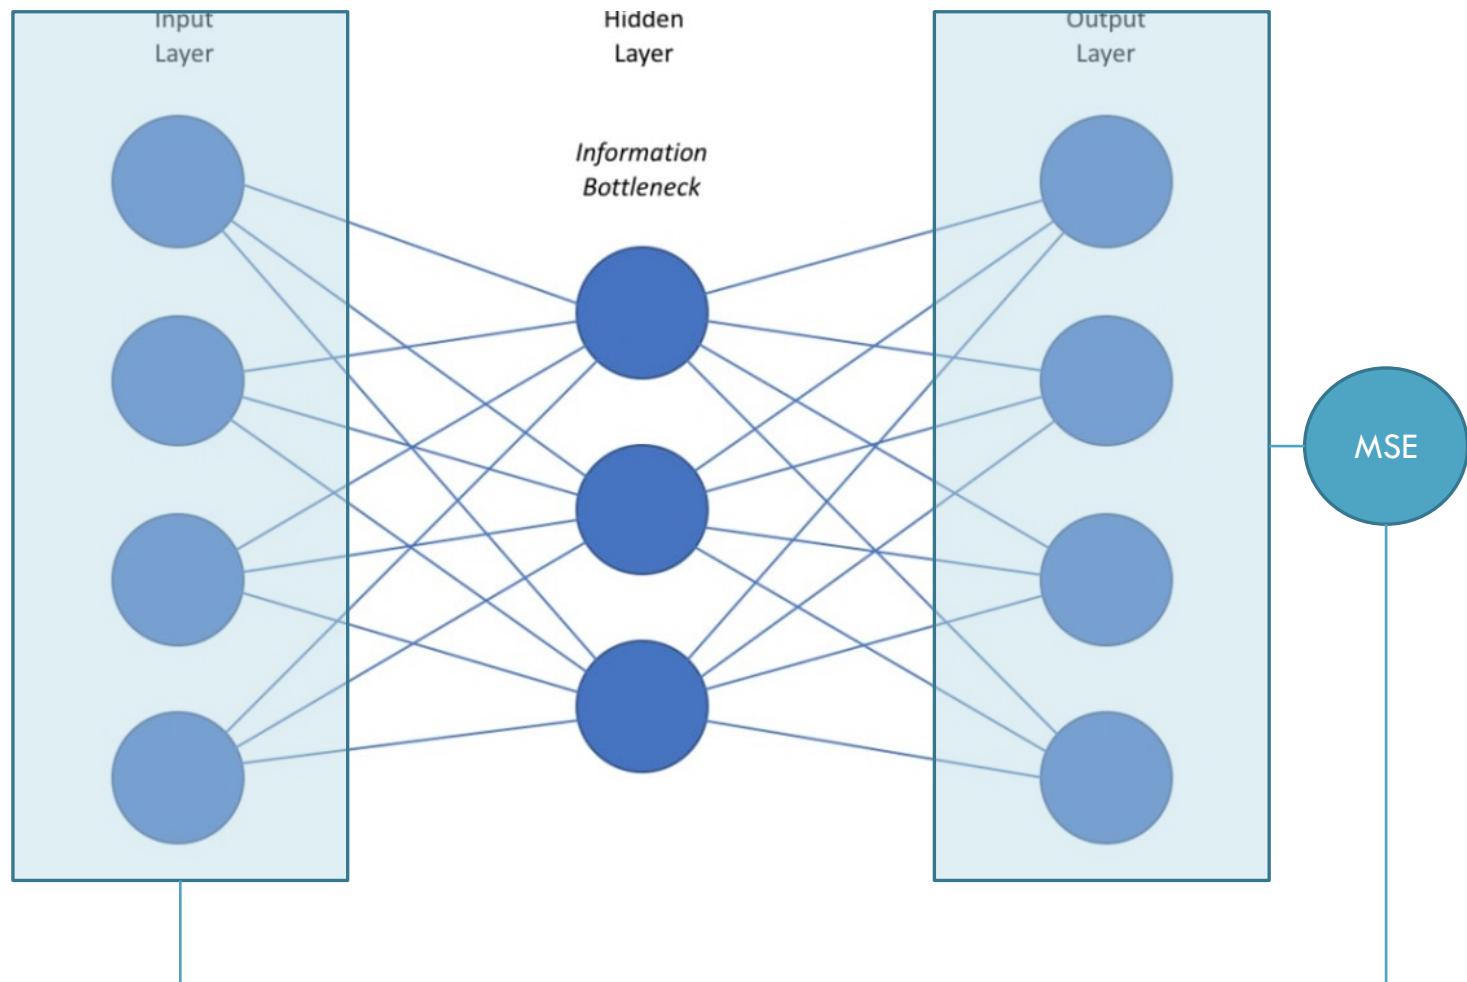
$$\text{Decoder}(\text{Encoder}(x)) = X$$



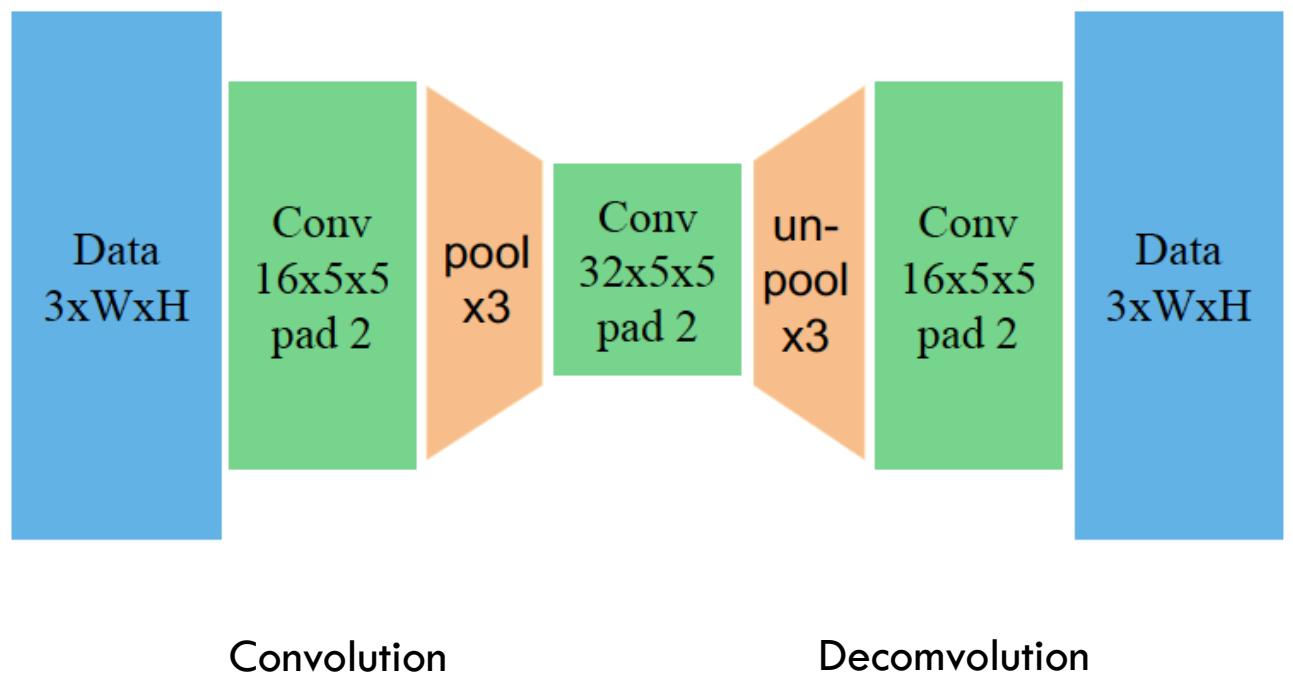
Cómo se entrena?

Backpropagation!

Minimizando el MSE entre el input y el output.

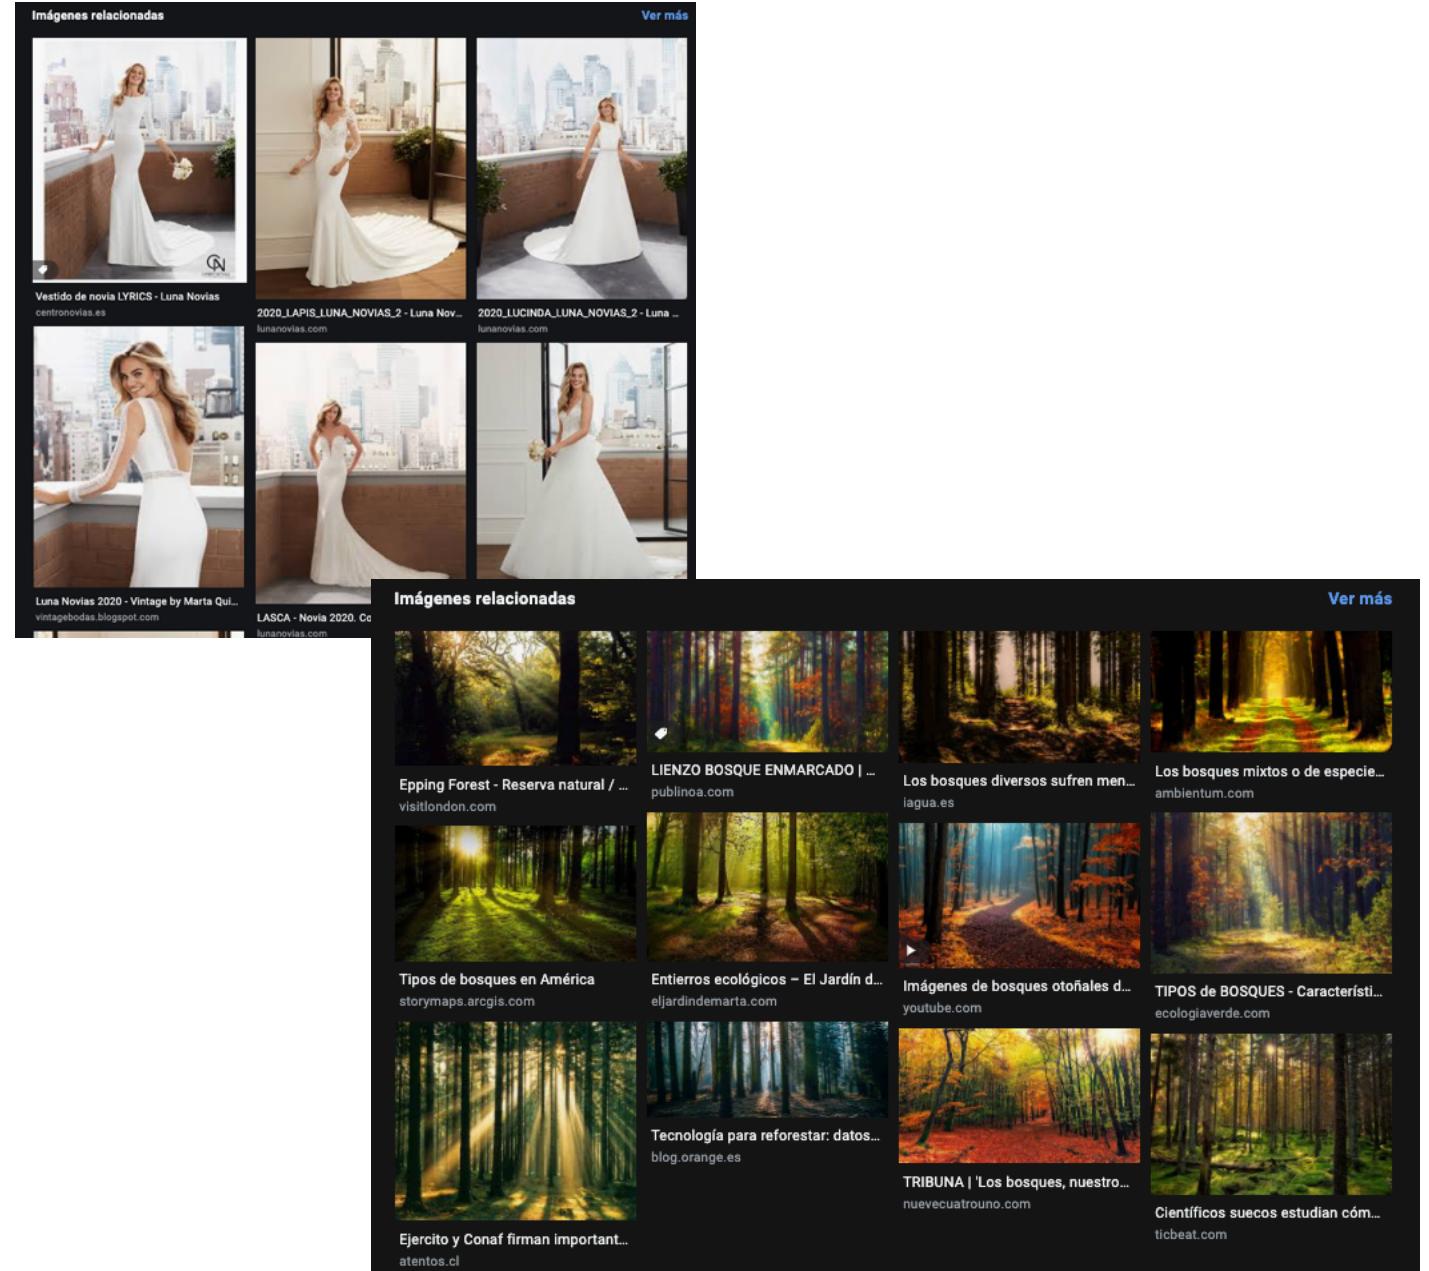


Qué pasa si el input es una imagen?



Autoencoder Aplicaciones

Image Clustering



Autoencoder Aplicaciones

Style mixing



Tipos de Autoencoders

Podemos utilizar diferentes pre-entrenados AE

Sparse Autoencoder

$$L = \|X - Dec(Enc(X))\| + \sum_i |Enc_i(X)|$$

variational Autoencoder

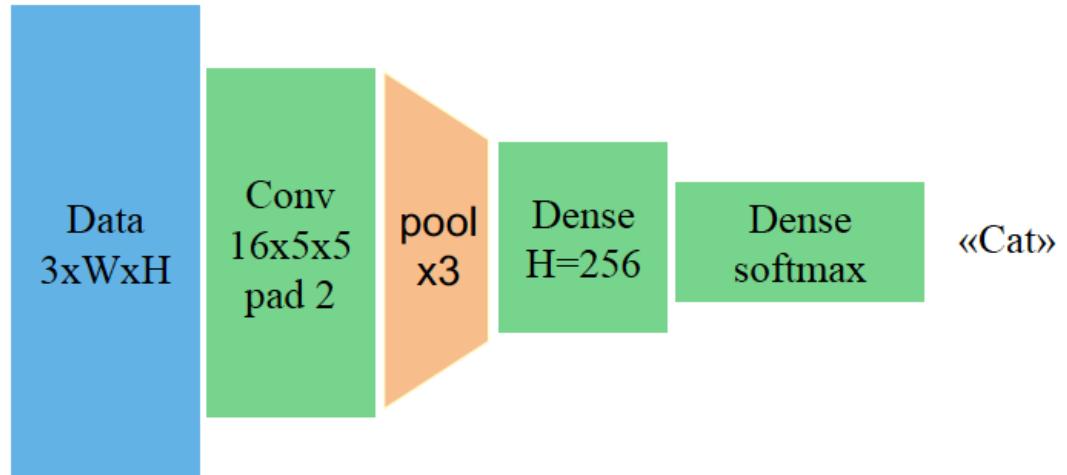
$$L = \|X - Enc(Dec(Noize(X)))\|$$

Expanding Autoencoder

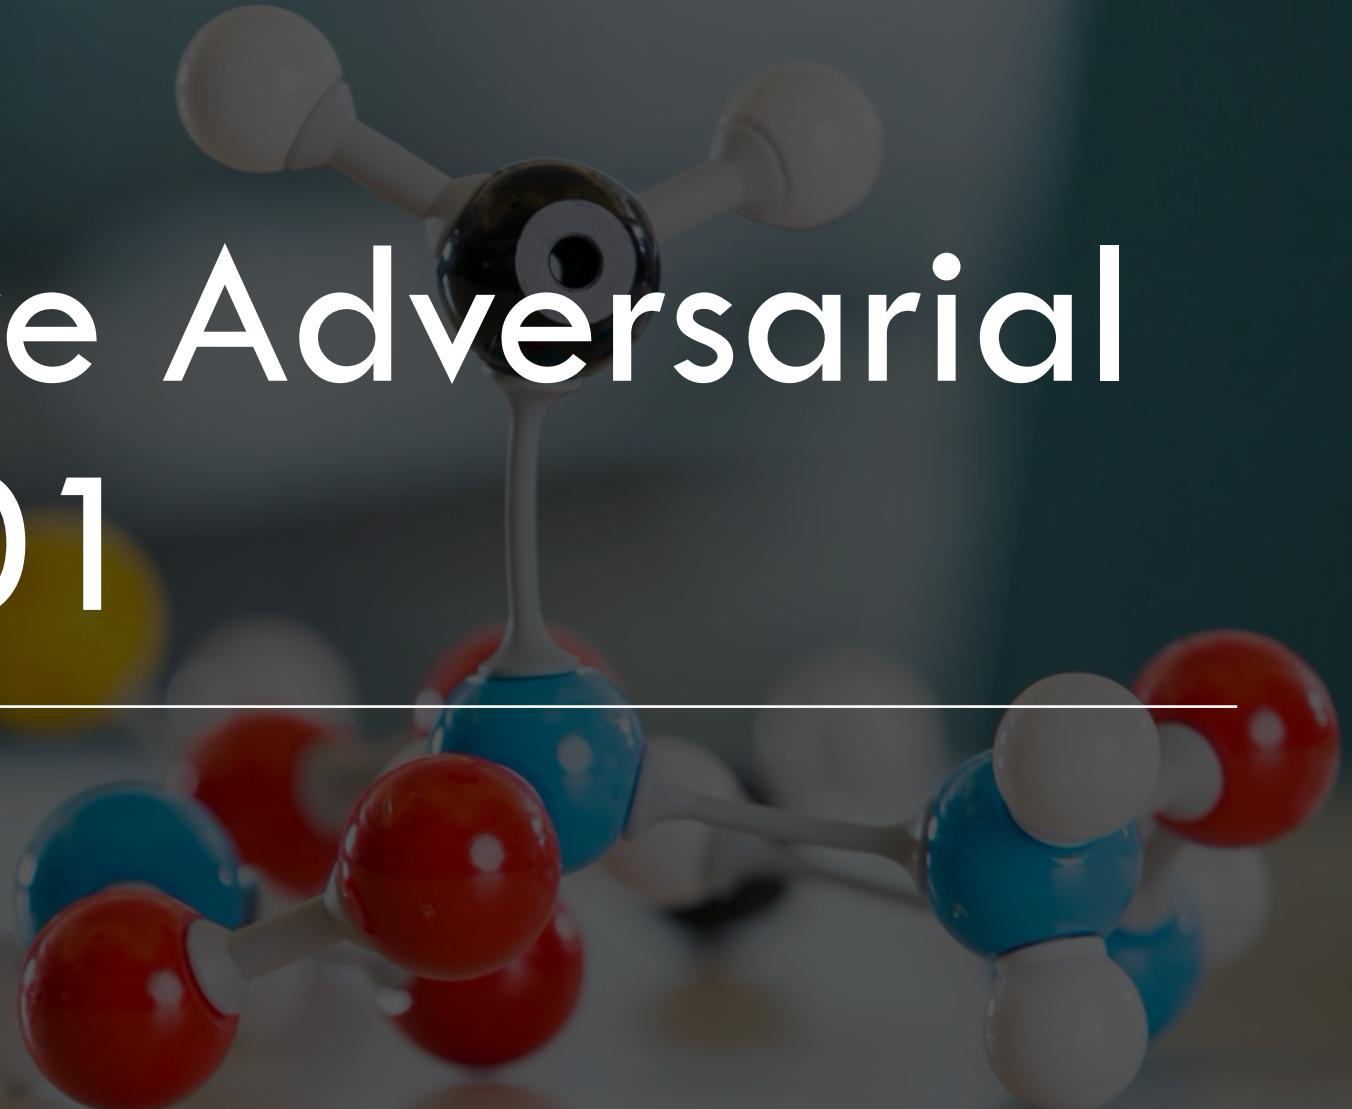
$$L = \|X - Dec(Enc(X))\|$$

Denoizing Autoencoder

$$L = \|X - Enc(Dec(Noize(X)))\|$$



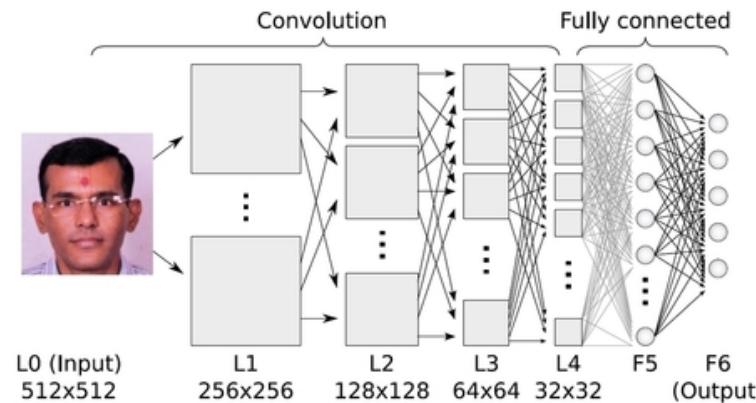
Generative Adversarial models 101



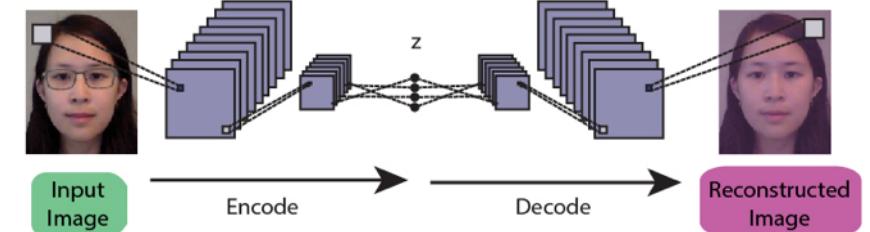
Objetivo:

Queremos generar nuevas imágenes “Falsas” que parezcan reales, vamos a dejarlo en rostros de personas.

CNN puede clasificar
imágenes de rostros



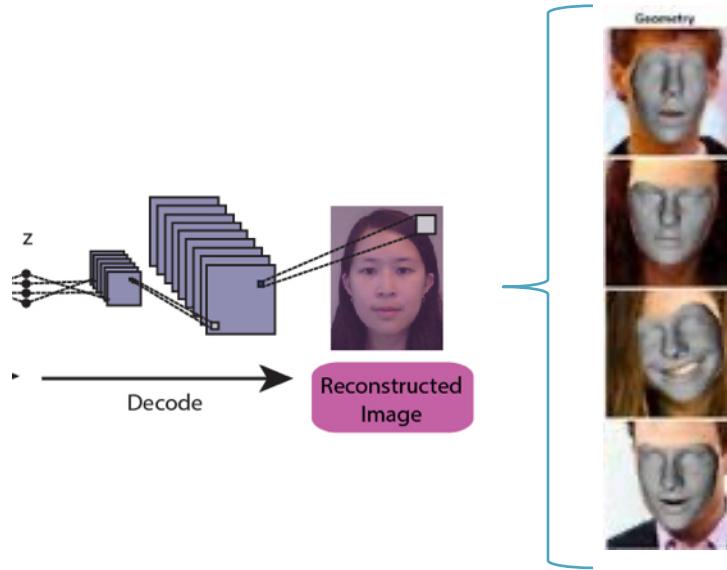
Autoencoder puede descomponer y reconstruir
imágenes de rostros



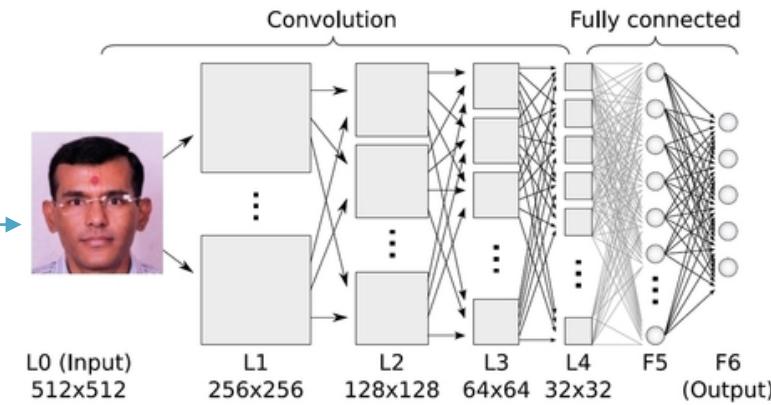
Objetivo:

Queremos generar nuevas imágenes “Falsas” que parezcan reales, vamos a dejarlo en rostros de personas.

Muestreo
aleatorio de
n objetos



CNN puede clasificar
imágenes de rostros

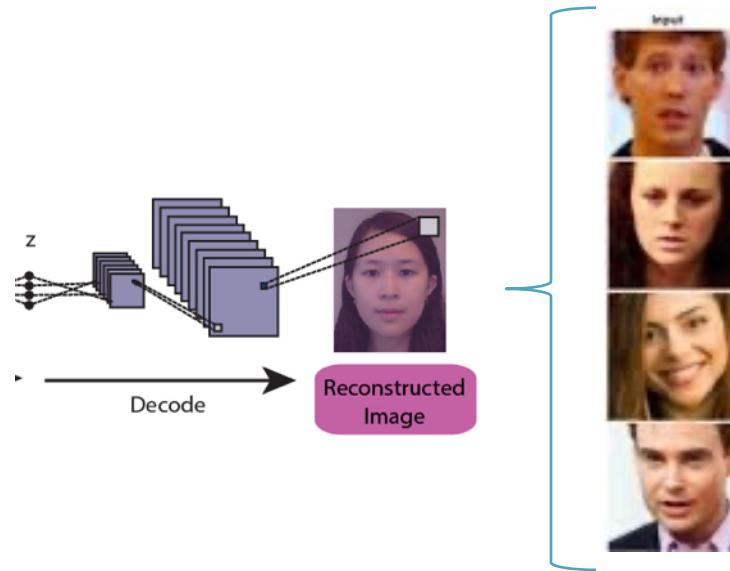


MSE

Objetivo:

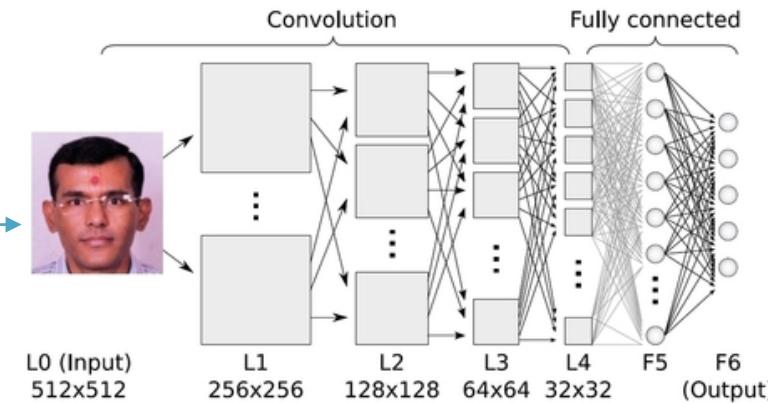
Queremos generar nuevas imágenes “Falsas” que parezcan reales, vamos a dejarlo en rostros de personas.

Muestreo aleatorio de n objetos



Una vez que este proceso converge, nuestro sampler se convierte en un generador de rostros capaz de engañar a una red neuronal

CNN puede clasificar imágenes de rostros





Questions

Que vimos hoy?

Calentamiento

- SGD
- MLP y Backpropagation
- FNN
- NLP y sentiment Analysis

Computer vision

- Convoluciones
- Feature extraction
- Image classification

Arquitecturas avanzadas

- RNN
- Autoencoders
- GANs

La punta del iceberg



Otros paradígmasis

Semi-supervised Learning

Transfer Learning

Reinforcement Learning

Recursos Adicionales



Dive into Deep Learning

An interactive deep learning book with code, math, and discussions

Provides NumPy/MXNet, PyTorch, and TensorFlow implementations

<https://d2l.ai/>

FREE COURSE

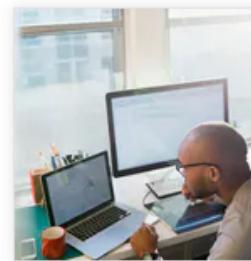
Intro to TensorFlow for Deep Learning

by  TensorFlow

This course is a practical approach to deep learning for software developers

[START FREE COURSE](#)

coursera



Machine Learning

University of Washington

Especialización

 4.6 (13,927) | 380K estudiantes | **PLUS**

 Intermediate

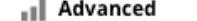


Advanced Machine Learning

National Research University Higher School of Economics

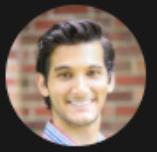
Especialización

 4.5 (3,482) | 290K estudiantes | **PLUS**

 Advanced

<https://www.udacity.com/course/intro-to-tensorflow-for-deep-learning--ud187>

<https://www.coursera.org/specializations/machine-learning>
<https://www.coursera.org/specializations/aml>



Alexander Amini
75,900 suscriptores



Abhishek Thakur
32,000 suscriptores



TensorFlow ✅
290,000 suscriptores



DeepMind
251,000 suscriptores



stanfordonline
201,000 suscriptores

Thank you

